

Nasaq at AraGenre 2026: Definition-Guided Approaches for Hierarchical Arabic Genre Classification

Shahad Altamimi¹, Samar Ahmed¹, Morooj Alrajhi¹, Yahya Assiri¹, Abdullah I. Alharbi¹

¹King Salman Global Academy for Arabic Language (KSAA),
Riyadh, Saudi Arabia

{sealtamimi, sahmad, malrajhi, Yasiri, aialharbi}@ksaa.gov.sa

Abstract

We describe **SynAnchor**, Nasaq’s definition-guided system for AraGenre 2026. The task requires hierarchical Arabic genre classification under disjoint specific-genre inventories across training, development, and test, making generalization to unseen genres a central challenge. SynAnchor addresses this setting by generating Arabic anchor texts from genre definitions using GPT-4o and embedding them with BGE-M3 for similarity-based retrieval. High-confidence cases are classified directly, while ambiguous inputs are selectively adjudicated by GPT-5 among the three highest-ranked candidate genres. On the official test set of 27,972 texts spanning 74 unseen specific genres, SynAnchor achieves a Hierarchical Macro-F1 of 0.6006 on Codabench, compared with 0.2472 for the organizers’ official baseline (approximately 2.4×). Approximately 25% of test instances are resolved using retrieval alone, without LLM adjudication. The results suggest that combining definition-guided Arabic anchor synthesis with selective LLM reasoning can support hierarchical genre classification under changing label inventories.

1 Introduction

Arabic genre classification is hard for two reasons: the language itself varies enormously, from Modern Standard Arabic to Classical Arabic to a range of regional dialects (Habash, 2010), and genre-labelled Arabic data at fine granularity is simply scarce. A model that learns to recognize genres by memorizing patterns from a fixed set of training labels runs into trouble as soon as the label set changes at deployment time (Yin et al., 2019). That is exactly the situation this shared task creates.

The AraGenre 2026 shared task (El-Haj et al., 2026) requires hierarchical Arabic genre classification: each text receives a broad genre (one of six) and a fine-grained specific genre. Texts span Modern Standard Arabic, Classical Arabic, and

several dialects. The official metric is hierarchical macro-F1 (H-F1). Specific genre inventories are disjoint across training (7 genres), development (6), and test (74), while the six broad categories recur across splits; systems must therefore generalize from English genre definitions rather than memorizing labels.

This disjointness is not incidental to the task design but its central difficulty: a classifier trained end-to-end on seven specific genres has, by construction, never observed 74 of the categories it must distinguish at test time, and must instead infer what each unseen genre means from its English-language definition alone, then apply that understanding to Arabic input text.

We present **SynAnchor**, which converts each English definition into a bank of short Arabic example texts (*anchors*), ranks genres by embedding similarity, and invokes GPT-5 only when retrieval confidence is low. This design targets the cross-lingual gap (English definitions, Arabic inputs) and fine-grained sibling confusion without any labelled training data for test genres.

We do not use test gold labels for tuning or analysis; all reported test scores come from the official Codabench evaluation.¹

Our contributions are:

- **SynAnchor**, a training-free pipeline that bridges English genre definitions and Arabic inputs through synthesized anchor texts, margin-gated embedding retrieval, and selective LLM adjudication for ambiguous cases;
- A systematic comparison against supervised alternatives encoder fine-tuning with contrastive objectives, hierarchical prediction variants, and a synthetic-data AraBERT baseline under the same disjoint

¹Code: <https://github.com/ShahdAi/nasaq-aragenre-2026>.

train/development/test genre setting showing that strong development-set performance for fine-tuned encoders does not transfer to the 74-genre test inventory;

- Official test results of H-F1 0.6006 in Ara-Genre 2026, exceeding the organizers’ official baseline (0.2472) by approximately $2.4\times$ and outperforming all alternative configurations we evaluated on the test.

The remainder of this paper is organized as follows. Section 2 reviews related work. Section 3 describes the shared-task data. Section 4 presents the SynAnchor pipeline. Section 5 summarizes the experimental setup. Section 6 reports results against alternative configurations. Section 7 discusses generalization patterns. Section 8 presents a qualitative error analysis.

2 Related Work

Automatic genre identification has seen growing interest across languages. Kuzman and Ljubešić (2025) survey two decades of approaches, finding that Transformer-based models lead cross-lingual performance (F1: 0.61–0.84), though fewer than half of existing datasets are publicly available.

Kuzman et al. (2023) show that ChatGPT outperforms a fine-tuned XLM-RoBERTa on unseen genre data (micro-F1: 0.74 vs. 0.67), with performance driven by prompt language rather than text language supporting the use of English-prompted LLMs for non-English genre tasks.

For Arabic specifically, Ferhat et al. (2024) built a 3,400-document Arabic FTD corpus across 17 categories, finding CAMeLBERT the strongest classifier in their evaluation.

Bulygin and Sharoff (2018) instead trained on machine translated English and Russian data, showing genre signals transfer across translation.

Both approaches rely on prototypical genre examples matching texts against collected instances. SynAnchor follows the same intuition but replaces corpus collection with generation: anchors are produced directly from genre definitions, enabling classification over inventories unseen during development.

3 Dataset

The shared task provides Arabic texts in Modern Standard Arabic, Classical Arabic, and regional dialects, with disjoint specific-genre inventories

across splits and six recurring broad categories. Training contains 140 texts across seven genres, development 110 texts across six genres, and test 27,972 texts across 74 unseen genres.

Test texts cover diverse forms, including job advertisements, book blurbs, news, scripture, poetry, and dialectal lyrics, with several closely related sibling genres. English definitions are provided for all genres, but no labelled Arabic examples are available for test-only genres. Systems must output `broad_genre` and `specific_genre` fields following the official JSON schema.

4 System Overview

Figure 1 shows the submitted *flat* system; alternative configurations are evaluated in §6.

4.1 Anchor Generation

GPT-4o (Hurst et al., 2024) generates up to 12 *persona* and 12 *boundary* Arabic texts per specific genre from the official definitions (temperature 0.7 for diversity).

Boundary neighbors are the three nearest genres by BGE-M3 (Chen et al., 2024) cosine similarity of their English definition texts.

A GPT-4o quality control pass drops anchors that fit a sibling genre better than their assigned genre (prompt in Appendix B.2).

Formal genres use MSA; dialect appears only when definitions require it (e.g., Egyptian/Gulf lyrics, interactive posts; Table 2).

Generation runs once and is cached (1,531 anchors over 74 genres; mean 20.7 after filtering, range 10–24; offline prompts in Appendix B.2).

4.2 Retrieval and Decision

BGE-M3 embeds all anchors once. Each input is scored per genre by mean cosine similarity to its $k=3$ nearest anchors; genres are ranked globally.

If the score margin between rank-1 and rank-2 is ≥ 0.04 , top-1 is accepted; else GPT-5 (OpenAI, 2025) receives the input, the top-three genre definitions, and three cached anchors per candidate, and returns a (broad, specific) JSON pair (temperature 0; prompt verbatim in Appendix B.1).

Broad genre follows from the chosen specific genre; parse failure or invalid output falls back to anchor top-1 (7 cases on test).

4.3 Alternative Approaches Explored

Hierarchical variant: predict broad, then specific within that broad (GPT-4o (Hurst et al., 2024); Ta-




Figure 1: SynAnchor flat pipeline: offline anchor synthesis and margin-gated BGE-M3 retrieval with selective GPT-5 adjudication.

ble 1). **AraBERT baseline:** fine-tuned on synthetic definition-derived data (Antoun et al., 2020) (dev only; Table 1).

Contrastive fine-tuning (e5-large): maps texts and category documents into a shared embedding space (Wang et al., 2024) via MultipleNegativesRankingLoss (Henderson et al., 2017) with implicit in-batch negatives (Table 1).

Hierarchical contrastive fine-tuning with hard negatives (e5-base): extends this with three hard-negative sources (Robinson et al., 2021) and Leave-One-Genre-Out cross-validation (Arlot and Celisse, 2010), with two-stage hierarchical prediction at inference (Table 1).

5 Experimental Setup

Development (110 texts, six genres) guided all design choices, including the margin threshold; training genres (seven, unseen in dev) checked generalization; the test set (27,972 texts, 74 genres) was used only to obtain final Codabench scores for the configurations reported in Table 1; all design choices, including the margin threshold, were made on development. Test labels were never used for tuning or qualitative error counting.

The submitted configuration (synanchor.yaml) uses GPT-4o for anchor generation, BGE-M3 for retrieval, $k=3$, margin 0.04 (selected on dev from $\{0.02, 0.04, 0.06\}$), a top-3 shortlist, and GPT-5 adjudication at temperature 0 with max_tokens=4096. No text preprocessing was applied.

The official evaluation metric is hierarchical

macro-F1, computed as the mean of broad macro-F1 and specific macro-F1. The main replication settings used by the submitted system are reported in Appendix A.

6 Results

Table 1 summarizes main scores.

On test, SynAnchor attains H-F1 0.6006 on 74 genres (27,972 texts, Codabench).

The organizers’ official baseline, submitted under the “AraGenre” team name, scores a Hierarchical Macro-F1 of 0.2472 on the 74-genre test set (Codabench). On development, the organizers report 0.4658 for their embedding-similarity baseline (Baseline 1); the substantially lower test figure reflects the harder unseen-genre inventory.

The hierarchical GPT-4o variant reaches 0.5782 on test, below the submitted flat system (+0.0224 absolute H-F1).

Both contrastive fine-tuning approaches generalize poorly despite strong development-set scores: e5-large drops from 0.8716 (dev) to 0.358 (test), and the hierarchical e5-base variant drops from 0.9032 (dev) to 0.3152 (test).

Inference routing on test: 6,980 anchor-only ($\approx 25\%^2$), 20,992 GPT-5-adjudicated ($\approx 75\%$), and 7 fallback to anchor top-1 after adjudication failure.

7 Discussion

Table 1 shows a sharp generalization gap: development performance does not transfer to the disjoint 74-genre test set.

²Percentages rounded; exact count: 6,980/27,972.

System	Split	Hierarchical Macro F1	Specific Macro F1	Specific Weighted F1	Specific Accuracy	Broad Macro F1	Broad Weighted F1	Broad Accuracy
SynAnchor (submitted)	test	0.6006	0.4264	0.4548	0.4634	0.7748	0.8312	0.8298
Hierarchical variant (GPT-4o)	test	0.5782	—	—	—	—	—	—
Contrastive fine-tuning (e5-large)	test	0.358	0.1714	0.1802	0.2098	0.5447	0.6142	0.6054
Hierarchical contrastive w/ hard negatives (e5-base)	test	0.3152	0.1894	0.1886	0.2001	0.441	0.4909	0.4737
Organizers' baseline (AraGenre)	test	0.2472	0.1358	0.1467	0.1597	0.3585	0.4527	0.445
SynAnchor (anchor retrieval)	dev	0.873	—	—	—	—	—	—
AraBERT (synthetic fine-tune)	dev	0.7762	—	—	—	—	—	—
Contrastive fine-tuning (e5-large)	dev	0.8716	—	—	—	—	—	—
Hierarchical contrastive w/ hard negatives (e5-base)	dev	0.9032	—	—	—	—	—	—
Baseline 1:Definition embedding	dev	0.4658	0.4658	0.4999	0.4818	0.4658	0.4999	0.4818
Baseline 2:Definition RAG	dev	0.3899	0.3899	0.3890	0.4273	0.3899	0.3890	0.4273

Table 1: Main results. Full metrics are reported for the submitted test system; other systems report H-F1 only. Dev and test scores are not directly comparable due to disjoint genre inventories.

Both contrastive fine-tuning variants collapse relative to their dev scores, although they still exceed the organizers' official baseline on test (0.358 and 0.3152 vs. 0.2472).

Whatever these models learn about aligning Arabic text with English definitions during training, it does not appear to survive the jump to a much larger and more varied genre inventory at test time.

We were not able to evaluate the AraBERT baseline (0.7762 dev) on test, so a direct comparison is not possible.

SynAnchor encodes no genre-specific knowledge into trained weights; instead, it retrieves cached Arabic anchors at inference time and calls GPT-5 only when retrieval is ambiguous.

On test it reaches H-F1 0.6006, an absolute gain of 0.0224 over the hierarchical variant (0.5782) and approximately $2.4\times$ the organizers' baseline (0.2472).³

This is consistent with looking up genre knowledge from definitions at inference rather than baking it into fixed weights.

That said, SynAnchor and the hierarchical variant share the same GPT-4o/GPT-5 backbone and anchor infrastructure, so the flat-vs-hierarchical comparison is confounded by decision architecture rather than retrieval alone.

8 Error Analysis

Because test gold labels are not released, we cannot compute a confusion matrix on test; we instead report qualitative hypotheses grounded in development. The dev-to-test drop (H-F1 0.873 $\rightarrow$ 0.6006) may reflect anchor distribution shift the 1,531 anchors were generated from submission-time definitions and may under-cover the register and topic diversity of the test corpus and register mismatch,

since anchors default to MSA while test texts may be informal or mixed. The flat architecture also derives the broad genre from the specific one, so specific errors propagate to the broad label.

On development, where gold is available, 14 of 110 texts (12.7%) were misclassified (Table 3), concentrated in the sibling-confusion pattern above. Seven Advice Columns texts were predicted as Motivational Writing (both give encouraging second-person guidance), and six Religious Commentary texts were predicted as Analytical or Educational genres, since they discuss interpretive traditions in an academic register with no overt religious vocabulary precisely the case our *implicit* anchors target.

9 Conclusion

SynAnchor shows that definition-guided anchor synthesis, combined with margin-gated retrieval and selective LLM adjudication, can generalize to 74 unseen Arabic genres without labelled training data for those genres. On development, the anchor-based configuration achieved strong performance, while the hierarchical configuration provided no corresponding advantage. However, matched ablations were not conducted on test, and test gold labels are unavailable, limiting conclusions about individual design components. Limitations include dependence on proprietary LLM APIs for $\approx 75\%$ of test inputs, the lack of test gold for quantitative error analysis, and non-comparable dev/test splits. Future work includes matched flat versus hierarchical ablations, cheaper open-weight adjudicators, and stronger contrastive anchors for template siblings.

³Ratio $0.6006/0.2472 \approx 2.43$.

References

- Wissam Antoun, Fady Baly, and Hazem Hajj. 2020. [AraBERT: Transformer-based model for Arabic language understanding](#). In *Proceedings of the 4th Workshop on Open-Source Arabic Corpora and Processing Tools, with a Shared Task on Offensive Language Detection*, pages 9–15. European Language Resources Association.
- Sylvain Arlot and Alain Celisse. 2010. [A survey of cross-validation procedures for model selection](#). *Statistics Surveys*, 4:40–79.
- Mikhail V. Bulygin and Serge A. Sharoff. 2018. [Using machine translation for automatic genre classification in Arabic](#). In *Proceedings of the International Conference on Computational Linguistics (Dialogue)*.
- Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. 2024. [BGE-M3 embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation](#). *Preprint*, arXiv:2402.03216.
- Mo El-Haj, Saad Ezzini, Shadi Abudalfa, Mustafa Jarrar, Nguyen Minh Chi, and Nguyen Minh Quan. 2026. Aragenre 2026: A hierarchical definition-guided arabic genre classification shared task. In *Proceedings of the 4th Arabic Natural Language Processing Conference (ArabicNLP 2026)*, Budapest, Hungary. Association for Computational Linguistics.
- Zeyd Ferhat, Abir Betka, Riyadh Barka, Zineddine Kahoul, Selma Boutiba, Mohamed Tiar, Habiba Dahmani, and Ahmed Abdelali. 2024. [Functional text dimensions for Arabic text classification](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 352–360, Bangkok, Thailand. Association for Computational Linguistics.
- Nizar Y. Habash. 2010. *Introduction to Arabic Natural Language Processing*. Morgan & Claypool Publishers.
- Matthew Henderson, Rami Al-Rfou, Brian Strope, Yun-Hsuan Sung, László Lukács, Ruiqi Guo, Sanjiv Kumar, Balint Miklos, and Ray Kurzweil. 2017. [Efficient natural language response suggestion for smart reply](#). *Preprint*, arXiv:1705.00642.
- Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. 2024. Gpt-4o system card. *arXiv preprint arXiv:2410.21276*.
- Taja Kuzman and Nikola Ljubešić. 2025. [Automatic genre identification: A survey](#). *Language Resources and Evaluation*, 59(1):537–570.
- Taja Kuzman, Igor Mozetič, and Nikola Ljubešić. 2023. [ChatGPT: Beginning of an end of manual linguistic data annotation? use case of automatic genre identification](#). *Preprint*, arXiv:2303.03953.
- OpenAI. 2025. [GPT-5 system card](#). OpenAI technical report.
- Joshua Robinson, Ching-Yao Chuang, Suvrit Sra, and Stefanie Jegelka. 2021. [Contrastive learning with hard negative samples](#). In *Proceedings of the 9th International Conference on Learning Representations*.
- Liang Wang, Nan Yang, Xinyu Huang, Linjun Yang, Rangan Majumder, and Furu Wei. 2024. [Multilingual E5 text embeddings: A technical report](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics.
- Wenpeng Yin, Jamaal Hay, and Dan Roth. 2019. [Benchmarking zero-shot text classification: Datasets, evaluation and entailment approach](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*, pages 3914–3923. Association for Computational Linguistics.

A Replication and Tables

A.1 Reproducibility Details

The submitted configuration (`synanchor.yaml`) uses GPT-4o for anchor generation and BGE-M3 for retrieval. Up to 12 persona and 12 boundary anchors are generated per specific genre using temperature 0.7, followed by a GPT-4o quality-control pass.

The resulting anchor cache contains 1,531 anchors across 74 specific genres, with a mean of 20.7 anchors per genre after filtering and a range of 10–24 anchors per genre.

At inference time, BGE-M3 scores each input against the cached anchors using the mean cosine similarity of the $k=3$ nearest anchors for each genre.

The submitted margin threshold is 0.04, selected on development data from $\{0.02, 0.04, 0.06\}$.

Inputs below this margin are passed to GPT-5 using a top-3 candidate shortlist and three cached anchors per candidate.

GPT-5 is run at temperature 0 with `max_tokens=4096`. No text preprocessing is applied.

The broad genre is derived from the selected specific genre. If GPT-5 produces an invalid output or parsing fails, the system falls back to the anchor-based top-1 prediction.

The complete cached anchor set used by the submitted system should be provided as supplementary material to enable reproduction of the reported retrieval results.

A.2 Submission Format and Evaluation

The system outputs one JSON object per input instance with the required `broad_genre` and `specific_genre` fields, following the official starter-kit schema.

The final submitted predictions were evaluated using the official Codabench evaluation, which produced the reported Hierarchical Macro-F1 score of 0.6006.

A.3 Cached Anchor Examples

B Prompt Templates

B.1 Inference Adjudication (GPT-5)

Used when anchor margin < 0.04 .

{candidates_block} = top-3 definitions + three input-nearest anchors; {text} = input.

Inference adjudication (GPT-5)

```
You are an expert Arabic linguist. Decide which single
↳ genre best fits the Arabic text below. The candidates
↳ were pre-selected as the closest matches.

## Candidates (definitions + representative examples):
{candidates_block}

## Arabic Text:
{text}

## Decision principles:
1. Subject matter and communicative purpose outweigh
↳ writing style.
2. Attend to the Arabic linguistic register: vocabulary of
↳ classical scholarly or interpretive traditions can
↳ signal a genre even when its most obvious keywords are
↳ absent.
3. A text addressing a specific personal problem differs
↳ from a text offering broad inspiration; a text whose
↳ subject is faith/morality-interpretation differs from
↳ general analysis, regardless of tone.

Return ONLY:
{"broad_genre": "<exact name>", "specific_genre": "<exact
↳ name>"}
```

B.2 Anchor Generation (GPT-4o, cache v4)

Offline templates (temp. 0.7) for the cached anchors in Table 2.

Each genre receives two complementary anchor banks: **Persona** (diverse clear positives) and **Boundary** (hard positives near the nearest confusable neighbor).

Persona anchors

You are an expert Arabic writer. Generate {n} short,
↳ realistic Arabic texts (1-3 sentences each) that are
↳ clear examples of the genre defined below.

Genre: {specific_genre}
(broad category: {broad_genre})

Broad definition:
{broad_genre_definition}

Genre definition:
{specific_genre_definition}

Generate the texts as if written by DIFFERENT author
↳ personas, covering at least 5 of:

formal journalist,
casual social-media user,
traditional scholar,
community advisor,
official spokesperson,
everyday storyteller,
educator.

Vary register (MSA and informal), length, and subtopic.

At least 4 texts must be IMPLICIT: the genre recognizable
↳ from purpose/register/framing alone, without its most
↳ characteristic vocabulary.

- if the genre definition specifies a particular dialect,
↳ regional variety, register, or sociolect of Arabic, the
↳ texts must authentically use THAT variety's own function
↳ words, pronouns, verb conjugations, and idiomatic
↳ expressions throughout not MSA with a single topical
↳ word swapped in.

Return ONLY a JSON array of {n} strings.

No commentary, no markdown fences.

Boundary anchors

You are an expert Arabic writer. Generate {n} short Arabic
↳ texts that clearly belong to genre G but sit CLOSE to
↳ the boundary with genre N: they may share surface style
↳ or vocabulary with N, yet their subject matter and
↳ communicative purpose unambiguously match G's
↳ definition.

These are hard positive examples.

Genre G:
{specific_genre}
(broad category: {broad_genre})

G's definition:
{specific_genre_definition}

Genre N (nearby/confusable, for contrast only):
{neighbor_genre}

N's definition:
{neighbor_definition}

Return ONLY a JSON array of {n} strings.

No commentary, no markdown fences.

Quality-control pass.

Quality-control pass

For each anchor, decide whether it truly best fits its
↳ assigned genre, or whether it actually fits a DIFFERENT
↳ genre in the list better.

Given all genre definitions and the anchors for one assigned
↳ genre, return ONLY a JSON array of {n} booleans (true =
↳ keep, false = drop), same order as the anchors.

Specific genre	Broad	Anchor excerpt (cache v4)
classical_poetry	Creative	في خطوات الفقراء على الطرقات، تسمع أنين الزمن، وتبصر غزيرة الضمّت.
egyptian_song_lyrics	Creative	حبيبي يا نور العين، وينك ووين أحلامنا؟
gulf_song_lyrics	Creative	يا ليل، سامرني بصوت الغالي، وذكرني بأحلى لحظات العمر.
sales_jobs	Informative	فرصة ممتازة: مندوب مبيعات للانضمام لفريقنا الديناميكي... وتحقيق الأهداف البيعية الشهرية.
economy_news	Informative	تشير التقارير إلى أن النمو الاقتصادي في المنطقة سيتباطأ بسبب التوترات الجيوسياسية المستمرة.
math_encyclopedia	Informative	في الهندسة، المثلث هو شكل له ثلاث زوايا. مجموع زوايا المثلث دائماً ١٨٠ درجة.
literature_fiction_book_description	Informative	رواية «أغاني الغربة» تسرد قصة مهاجر يبحث عن هويته في بلد جديد.
twitter_posts	Interactive	👉 قالوا لي إن الابتسامه معديه. جربتها اليوم، وإنتوا؟
youtube_comments	Interactive	الطريقة اللي شرحت بيها الموضوع كانت واضحة جداً. أتمنى أتعلم منك المزيد.
app_reviews	Interactive	بصراحة، تجربة التطبيق كانت مخيبة للأمل، الإعلانات مزعجة وجودة الصوت سيئة.
high_school_textbooks	Learning	التاريخ هو مرآة الحضارات، من خلاله نتعرف على الأحداث والإنجازات التي شكلت حاضرنا ومستقبلنا.
values_student_writing	Learning	الإيتار يعني وضع احتياجات الآخرين قبل احتياجاتنا. هذه القيمة هي أساس بناء مجتمع متعاون.
constitutions	Legal	تعمل الدولة على تعزيز السلام والأمن الداخلي والخارجي، وتلتزم بحل النزاعات بطرق سلمية ودبلوماسية.
contracts	Legal	وفقاً لبنود التعاون، يتعهد الطرف الأول بتزويد الطرف الثاني بالمواد الخام المطلوبة في الوقت المحدد.
quran	Religious	وَلَقَدْ يَسَّرْنَا الْقُرْآنَ لِلذِّكْرِ فَهَلْ مِنْ مُدَكِّرٍ
hadith	Religious	قال نبينا العظيم: «الدين النصيحة». فلنكن نصحاء محبين للخير.
tafsir	Religious	تفسير قوله تعالى: «قُلْ هُوَ اللَّهُ أَحَدٌ»، يؤكد على توحيد الألهية وأهمية الإخلاص في العبادة.

Table 2: Verbatim cached Arabic anchors (test cache v4; submitted system). Punctuation normalized for rendering where the cache uses ellipsis or alternate quote glyphs.

Gold	Predicted	Development text (gist)
Advice Columns	Motivational Writing	ابدأ بخطوات صغيرة ولا تحاول تغيير كل شيء دفعة واحدة.
Advice Columns	Motivational Writing	أحياناً يكون الاستماع الجيد أهم من تقديم الحلول السريعة.
Religious Commentary	Analytical Reports	يشير بعض الشراح إلى أن الاختلاف بين الناس في الفهم والتفسير أمر طبيعي.
Religious Commentary	Educational Explanations	يلاحظ الباحثون أن اللغة المستخدمة في النصوص الوعظية تميل إلى الأسلوب العاطفي.

Table 3: Representative development-set misclassifications (gold vs. predicted). Both clusters are sibling confusions: practical advice mistaken for motivation, and academically-framed religious commentary mistaken for analytical or educational genres.