

Tarjih at StanceEval-2026: Argument-Guided Arabic Stance Detection

Samah Aloufi Joud Alrefaei Rafaa Alharbi
Wejdan Alsaedi Nada Almarwani

Department of AI and Data Science, CCSE, Taibah University
(e-mail: slhebi, 4454096, 4455019, 4454677, nmarwani, @taibahu.edu.sa)

Abstract

We present Tarjih, our proposed system for the Stance Detection in Arabic Language shared task (StanceEval-2026). Rather than predicting stance directly, Tarjih formulates the task as verification among competing arguments. An instruction-tuned LLM generates one *steelman* argument for each stance label, and a second blind pass evaluates each argument against the tweet without access to the gold label. Each argument–tweet is then jointly encoded with AraBERT to model their compatibility. From this pipeline, we derive three systems: a *gated classifier*, a *pair scorer*, and a *ranking read-out*. The pair scorer achieves the best seen-target test F_{avg2} of 0.73, while the gated classifier performs best on the unseen-target track with ≈ 0.77 . Overall, the lower test performance and ranking among the submitted systems indicate limited generalization across targets.

1 Introduction

Stance detection determines whether a text expresses a favorable, opposing, or neutral position toward a particular target. Unlike sentiment analysis, which primarily captures overall polarity, stance detection is target-dependent and may require inferring implicit positions from context, making it useful for analyzing political and social discourse (ALDayel and Magdy, 2021).

Early stance detection approaches relied on rule-based and supervised machine-learning methods with manually engineered lexical and syntactic features (Gera and Neal, 2025; Siddiqua et al., 2018; Liu et al., 2016). More recently, pretrained language models have enabled fine-tuning, few-shot learning, prompting-based methods, and Large Language Models (LLMs) for in-target and cross-target stance detection (Jamadi Khiabani and Zubiaga, 2025; Gera and Neal, 2025; Zhang et al., 2024; Cruickshank and Ng, 2026). Despite these advances, generalization to unseen targets remains

challenging, since models must learn transferable stance representations rather than patterns specific to the targets observed during training.

Although substantial progress has been made in stance detection, research has remained focused on English, while Arabic stance detection has received comparatively limited attention. The Arabic text on social networks introduces additional challenges due to the rich language morphology and dialectal variation with informal writing conventions such as spelling variation, abbreviations and non-standard grammatical forms Alkhathlan et al. (2025); Alturayef et al. (2023, 2022).

To address this limitation, Alturayef et al. (2022) introduced Mawqif, a multi-dimensional Arabic dataset consisting of 4,121 tweets across three targets. The Mawqif dataset was subsequently adopted in StanceEval 2024, the first shared task dedicated to Arabic stance detection, providing a common evaluation benchmark for stance detection across three predefined targets Alturayef et al. (2024). The submitted systems explored a range of approaches, including fine-tuning Arabic pretrained language models (Amal et al., 2024; Galal and Kaseb, 2024), multi-task learning with sentiment and sarcasm (Badran et al., 2024; Hasanaath and Alansari, 2024), ensemble, and LLM-based methods (Ibrahim, 2024; Alghaslan and Almutairy, 2024). More recently, Alkhathlan et al. (2025) proposed the ArabicStanceX dataset Alkhathlan et al. (2025) that includes 14,477 Saudi Arabic tweets covering 17 topics and investigated multi-topic, zero-shot, and few-shot stance detection. The results highlighted continuing challenges in predicting the NONE class and generalizing to unseen topics.

This direction was further extended by Albalawi et al. (2026b) who introduced Mawqif-XT, extending the Mawqif dataset with 996 manually annotated tweets from three additional targets and providing a dedicated benchmark for cross-target

Track	Split	#Tweets	Targets
Track 1	Train	3,502	Women Empowerment, Covid Vaccine, Digital Transformation
	Dev	619	Same as training
	Test	352	Women Driving
Track 2	Train	2,721	Covid Vaccine, Digital Transformation
	Dev	1,400	Women Empowerment
	Test	644	Ecars, Trimester

Table 1: Dataset statistics and target distribution across the two StanceEval-2026 tracks.

stance detection. The dataset was subsequently adopted in the StanceEval-2026 shared task, which evaluates Arabic stance detection under both seen-target and unseen-target settings (Albalawi et al., 2026a).

In this paper, we present Tarjih, our system submitted to StanceEval-2026, which formulates Arabic stance detection as *argument verification*. For each tweet, an instruction-tuned LLM generates *steelman* arguments for FAVOR, AGAINST, and NONE, and a second LLM evaluates the consistency with the source tweet. The resulting argument–tweet representations provide a common base for three submitted variants: a *gated classifier*, a *pair scorer*, and a *ranking read-out* that differ in how evidence is converted into the final stance prediction.

The main contributions of this work are: 1) an argument-verification formulation for Arabic stance detection; 2) a two-stage LLM generation and verification framework; 3) a comparison of three decision mechanisms under seen-target and unseen-target settings to examine how different evidence-selection strategies affect generalization.

2 Task Setup

StanceEval-2026 requires classifying the stance of an Arabic tweet toward a given target as FAVOR, AGAINST, or NONE. The task includes two tracks with different target splits, summarized in Table 1: Track 1 evaluates generalization to new targets within previously seen thematic domains, while Track 2 tests zero-shot generalization to entirely unseen targets. Performance is measured using F_{avg2} , the macro-averaged F_1 over the FAVOR and AGAINST classes. Further details are provided in (Albalawi et al., 2026a).




Figure 1: Tarjih overview: stance arguments are generated, verified against the tweet, and jointly encoded with AraBERT for final prediction.

3 Method

Figure 1 shows the overall framework. Given a tweet x and its target t , we first generate one argument for each possible stance: FAVOR, AGAINST, and NONE. A second LLM then evaluates how well each argument matches the tweet without seeing the gold stance label. Finally, each argument is paired with the tweet and encoded using AraBERT.

We build three systems on top of this shared pipeline: a *gated classifier*, a *pair scorer*, and a *ranking read-out*. The gated classifier and ranking read-out are submitted to both tracks, while the pair scorer is submitted only to the seen-target track.

3.1 Steelman Argument Generation and Verification

For each tweet–target pair, Qwen-2.5-14B-Instruct (Team, 2024; Yang et al., 2024) generates three one-sentence arguments:

$$R(x, t) = \{r_{\text{FAVOR}}, r_{\text{AGAINST}}, r_{\text{NONE}}\}. \quad (1)$$

The model is asked to produce the strongest possible argument for each stance, even when that stance is unlikely to be correct. We refer to this as *steelman generation*. Each argument must also be grounded in the tweet. This gives us three plausible competing interpretations rather than one strong explanation and two weak alternatives.

A second LLM pass acts as a blind verifier. It receives the tweet, target, and three generated arguments, but not the gold stance label. For each argument r_k , it produces a consistency score:

$$v_k \in [0, 1], \quad (2)$$

where a higher score means that the argument is more consistent with the tweet.

The verifier is used only as an additional signal, not as a stance classifier. Its highest-scoring argument matches the gold label in 71% of the training instances and 73% of the development instances. This shows that the verifier provides useful information but is not accurate enough to replace the gold labels. We therefore use its scores to identify weaker arguments or adjust their contribution during training, while supervision still comes from the gold stance labels. The full generation and verification prompts are provided in Appendix A.

3.2 Argument–Tweet Encoding

Each generated argument is encoded together with the tweet using a shared bert-large-arabertv02-twitter(Antoun et al.) encoder. For argument r_k , the input is

$$\langle t \oplus x, r_k \rangle. \quad (3)$$

Pairing the argument with the tweet is important because all three arguments are intentionally written to be plausible. The model therefore cannot rely only on how fluent or convincing an argument sounds. Instead, it must learn whether the tweet actually supports that argument. For example, a well-written AGAINST argument may appear convincing on its own even when the tweet clearly expresses FAVOR. Joint encoding allows the model to compare the argument directly with the evidence in the tweet.

In our experiments, encoding the arguments separately did not generalize well to unseen targets. We therefore use joint argument–tweet encoding in all three systems.

3.3 Prediction Systems

Gated classifier. The gated classifier combines information from the tweet and the three argument–tweet pairs. A learned gate assigns a weight to each argument, while the verifier helps mask arguments judged to be unreliable. The weighted argument representation is combined with the tweet representation and passed to a class-weighted classifier.

During training, we also use a ranking objective that encourages the argument corresponding to the gold stance to score higher than the other two. Because the verifier can make mistakes, the gate is trained before applying the verifier mask so that an incorrect verifier decision does not directly penalize the correct argument.

For the unseen-target track, we additionally augment the training data with extra targets and replace the original target with a generic phrase for a subset of training examples. This is intended to reduce reliance on target-specific patterns. The gated classifier is submitted to both tracks.

Pair scorer. The pair scorer uses a simpler decision rule. Instead of combining the three argument representations in a classifier, it assigns a score directly to each argument–tweet pair:

$$s_k = f(x, t, r_k). \quad (4)$$

The stance associated with the highest-scoring pair is selected:

$$\hat{y} = \arg \max_k s_k. \quad (5)$$

Training encourages the gold argument to receive a higher score than the other two while also learning to judge each pair independently. The verifier does not decide the correct label; it only reduces the influence of training examples whose generated argument is judged weak. The submitted system also uses class-prior factoring, class weighting, and minority oversampling.¹

Ranking read-out. The ranking read-out reuses the ranking head learned by the gated classifier. Instead of passing the combined representation to the final classifier, it directly selects the stance whose argument receives the highest ranking score. This system does not introduce additional prediction parameters and provides a simple comparison between combining all argument representations and directly selecting the best-supported argument.²

4 Results

Table 2 presents the development and test results in terms of F_{avg2} . On the seen-target track, the pair scorer achieves the best test score of 0.73, followed by the ranking read-out (0.72) and the gated classifier (0.70). This ordering differs from the development results, where the ranking read-out obtains the highest score. Although the reported development score for the pair scorer comes from an ablation setting and is therefore not directly comparable to the submitted configuration, its test performance suggests that directly scoring each argument–tweet pair can generalize well to a new target.

¹The pair scorer is submitted only to the seen-target track.

²The ranking read-out is submitted to both tracks.

System	Dev	Test
<i>Track 1: seen target</i>		
Gated classifier	.8788	.7033
Pair scorer	.8243 [†]	.7315
Ranking read-out	.8829	.7206
<i>Track 2: unseen target</i>		
Gated classifier	.8282	.7699
Ranking read-out	.8346	.7507

Table 2: F_{avg2} results. Development scores are average over three seeds. [†] the pair-scorer development score is from the available Pointwise ablation; the submitted system also uses prior factoring, class weighting, and minority oversampling.

A similar difference between development and test performance appears in the unseen-target track. The ranking read-out achieves the highest development score (0.83), but the gated classifier performs better on the test set, reaching ≈ 0.77 compared with 0.75 for the ranking read-out. The gated classifier also incorporates target augmentation in this track, which may have contributed to its stronger adaptation to targets not seen during training.

Overall, the best-performing development system is not the best-performing test system in either track. This highlights the difficulty of selecting models for target-generalization settings: strong performance on development targets does not necessarily translate into stronger performance on unseen targets.

5 Analysis

The three systems share the same argument-verification pipeline but differ in how the final prediction is produced. The gated classifier combines information from all three arguments, the pair scorer selects the highest-scoring argument-tweet pair, and the ranking read-out directly uses the ranking head learned by the gated model.

The test results do not follow the same ordering as the development results. In the seen-target track, the pair scorer achieves the best test score (0.73), while the ranking read-out performs best on development (0.88). In the unseen-target track, the gated classifier achieves the best test score (≈ 0.77), although the ranking read-out has the higher development score. Since the submitted systems also differ in some training settings, these results should be interpreted as system-level comparisons rather than evidence that one decision mechanism is always better.

The verifier agrees with the gold stance on about 71% of the training and development instances. This shows that its scores contain useful information, but are not reliable enough to replace the gold labels. We therefore use the verifier only as an auxiliary signal. Overall, the results show that the shared argument-tweet representation can support different prediction strategies across both evaluation tracks.

6 Conclusion

We presented Tarjih, an argument-verification approach for Arabic stance detection developed for StanceEval-2026. Tarjih generates one argument per stance, evaluates its consistency with the tweet, and jointly encodes each argument-tweet pair using AraBERT.

The pair scorer performed best on the seen-target track, while the gated classifier performed best on the unseen-target track. However, performance remained below the top submitted systems, highlighting the challenge of cross-targets generalization. These results suggest that argument verification is a promising direction, while further work is needed to improve its robustness and competitiveness.

Limitations

The proposed framework has several limitations. First, it relies on LLM-generated arguments, so errors made during generation or verification can affect the final prediction. The generated arguments were not manually reviewed, which means that possible biases or systematic errors may remain in the pipeline.

Second, the same model, Qwen-2.5-14B-Instruct, is used for both argument generation and verification. This may reduce the independence of the verification step. The verifier also agrees with the gold stance on only about 71% of the training and development instances, so we use it only as an auxiliary signal rather than as a source of labels.

Finally, the framework requires several LLM calls for each instance, making it more computationally expensive than an encoder-only model.

AI Usage Disclosure

We used Claude (Anthropic) as a coding assistant during system development, primarily to help draft training code, and evaluation utilities. All system design decisions, experiments, and interpretations of the results were carried out by the authors.

We also used AI tools for editing the manuscript language, including grammar, clarity, and conciseness. All technical content, claims, and analyses were written and verified by the authors.

References

- Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Saad Ezzini, Maram Kurdi, Asma Yamani, Ahmed Ashraf, Maged Al-Shaibani, and Nora Alturayef. 2026a. Stanceeval-2026: The second stance detection shared task. In *Proceedings of the Fourth Arabic Natural Language Processing Conference (ArabicNLP 2026)*, Budapest, Hungary. Association for Computational Linguistics.
- Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Maram Kurdi, Saad Ezzini, Asma Yamani, and Ahmed Ashraf. 2026b. [Mawqif-xt: An arabic benchmark dataset for cross-target stance detection](#). Preprint, arXiv:2608.09539.
- Abeer ALDayel and Walid Magdy. 2021. [Stance detection on social media: State of the art and trends](#). *Information Processing Management*, 58(4):102597.
- Mamoun Alghaslan and Khaled Almutairy. 2024. [MGKM at StanceEval2024 fine-tuning large language models for Arabic stance detection](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 816–822, Bangkok, Thailand. Association for Computational Linguistics.
- Ali Alkhathlan, Faris Alahmadi, Faris Kateb, and Hend Al-Khalifa. 2025. [Constructing and evaluating arabicstancex: a social media dataset for arabic stance detection](#). *Frontiers in Artificial Intelligence*, Volume 8 - 2025.
- Nora Alturayef, Hamzah Luqman, and Moataz Ahmed. 2023. [A systematic review of machine learning techniques for stance detection and its applications](#). *Neural Comput. Appl.*, 35(7):5113–5144.
- Nora Alturayef, Hamzah Luqman, Zaid Alyafeai, and Asma Yamani. 2024. [stanceeval2024 : The first arabic stance detection shared task](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 774–782, Bangkok, Thailand. Association for Computational Linguistics.
- Nora Saleh Alturayef, Hamzah Abdullah Luqman, and Moataz Aly Kamaleldin Ahmed. 2022. [Mawqif: A multi-label Arabic dataset for target-specific stance detection](#). In *Proceedings of the Seventh Arabic Natural Language Processing Workshop (WANLP)*, pages 174–184, Abu Dhabi, United Arab Emirates (Hybrid). Association for Computational Linguistics.
- Mezghani Amal, Rahma Boujelbane, and Mariem El-louze. 2024. [ANLP RG at StanceEval2024: Comparative evaluation of stance, sentiment and sarcasm detection](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 788–793, Bangkok, Thailand. Association for Computational Linguistics.
- Wissam Antoun, Fady Baly, and Hazem Hajj. Arabert: Transformer-based model for arabic language understanding. In *LREC 2020 Workshop Language Resources and Evaluation Conference 11–16 May 2020*, page 9.
- Mohamed Badran, Mo'men Hamdy, Marwan Torki, and Nagwa El-Makky. 2024. [AlexUNLP-BH at StanceEval2024: Multiple contrastive losses ensemble strategy with multi-task learning for stance detection in Arabic](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 823–827, Bangkok, Thailand. Association for Computational Linguistics.
- Iain J. Cruickshank and Lynnette Hui Xian Ng. 2026. [Prompting and fine-tuning open source large language models for stance classification](#). *ACM Trans. Intell. Syst. Technol.*, 17(4).
- Omar Galal and Abdelrahman Kaseb. 2024. [Team_Zero at StanceEval2024: Frozen PLMs for Arabic stance detection](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 783–787, Bangkok, Thailand. Association for Computational Linguistics.
- Parush Gera and Tempestt Neal. 2025. [Deep learning in stance detection: A survey](#). *ACM Comput. Surv.*, 58(1).
- Ahmed Hasanaath and Aisha Alansari. 2024. [StanceCrafters at StanceEval2024: Multi-task stance detection using BERT ensemble with attention based aggregation](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 811–815, Bangkok, Thailand. Association for Computational Linguistics.
- Michael Ibrahim. 2024. [CUFE at StanceEval2024: Arabic stance detection with fine-tuned llama-3 model](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 807–810, Bangkok, Thailand. Association for Computational Linguistics.
- Parisa Jamadi Khiabani and Arkaitz Zubiaga. 2025. [Cross-target stance detection: A survey of techniques, datasets, and challenges](#). *Expert Systems with Applications*, 283:127790.
- Can Liu, Wen Li, Bradford Demarest, Yue Chen, Sara Couture, Daniel Dakota, Nikita Haduong, Noah Kaufman, Andrew Lamont, Manan Pancholi, Kenneth Steimel, and Sandra Kübler. 2016. [IUCL at SemEval-2016 task 6: An ensemble model for stance detection in Twitter](#). In *Proceedings of the 10th International Workshop on Semantic Evaluation (SemEval-2016)*, pages 394–400, San Diego, California. Association for Computational Linguistics.

Umme Aymun Siddiqua, Abu Nowshed Chy, and Masaki Aono. 2018. Stance detection on microblog focusing on syntactic tree representation. In *Data Mining and Big Data*, pages 478–490, Cham. Springer International Publishing.

Qwen Team. 2024. [Qwen2.5: A party of foundation models](#).

An Yang, Baosong Yang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Zhou, Chengpeng Li, Chengyuan Li, Dayiheng Liu, Fei Huang, Guanting Dong, Haoran Wei, Huan Lin, Jialong Tang, Jialin Wang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Ma, and 40 others. 2024. Qwen2 technical report. *arXiv preprint arXiv:2407.10671*.

Bowen Zhang, Xianghua Fu, Daijun Ding, Hu Huang, Genan Dai, Nan Yin, Yangyang Li, and Liwen Jing. 2024. [Investigating chain-of-thought with chatgpt for stance detection on social media](#). *Preprint*, arXiv:2304.03087.

النص: "{text}"
الهدف: "{target}"

ملاحظة: قد يكون النص مكتوباً بالفصحى أو بلهجة عربية عامية. حلل المعنى المقصود بناءً على السياق، ولا تفترض أن الصياغة العامية تعني غموضاً أو حيادية. حلل موقف كاتب النص تجاه الهدف المذكور أعلاه.

لكل احتمال، اكتب مبرراً من جملة واحدة بوضوح لماذا قد* تُصنّف النص بهذا الشكل، حتى لو لم يكن هذا هو التصنيف الصحيح في رأيك. استشهد بجزء محدد من النص إن أمكن، ولا تختبر معلومات غير موجودة في النص.

صع كل مبرر بأسلوب عام قدر الإمكان بحيث يمكن فهم نمط الحجة حتى لو تغير الهدف.

المواقف:
1. مع (Favor):
2. ضد (Against):
3. محايد (None):

أخرج فقط بصيغة JSON:

```
{"favor": "...", "against": "...", "none": "..."}

```

Figure 2: Stage-A rationale generation prompt (original Arabic). Placeholders {text} and {target} are substituted per instance; the model returns one steelman rationale per label in the JSON schema shown. An English rendering is given in Appendix A.1.

A LLM Prompts

Both stages use Qwen-2.5-14B-Instruct with 4-bit NF4 quantization. Stage A generates the three steelman arguments, while Stage B evaluates them in a blind verification setting. The verifier is never shown the gold label, so all resulting verification scores and derived quantities remain available at test time.

The prompts were issued in Arabic. We provide faithful English renderings. The placeholders {text} and {target} are substituted for each instance.

A.1 Stage A: Steelman Argument Generation

For every tweet and target, the model receives the following instruction:

Text: "{text}"
Target: "{target}"

Note: the text may be written in Modern Standard Arabic or a colloquial Arabic dialect. Analyze the intended meaning from context; do not assume that colloquial phrasing implies vagueness or neutrality.

Analyze the stance of the text’s author toward the target above.

For each possibility, write a one-sentence argument explaining why the text *could* be classified this way, even if you believe this is not the correct label. Cite a specific part of the text where possible, and do not invent information absent from the text.

Phrase each argument as generally as possible, so that the argument pattern remains comprehensible

أنت حكم محايد. JSON فقط.

النص: "{text}"
الهدف: "{target}"

ملاحظة: قد يكون النص بالفصحى أو بلهجة عامية، وقد يحتوي على سخرية أو نفي ضمني. لا تحكم على الصياغة السطحية، بل على المعنى الفعلي المقصود.

المبررات التالية عن موقف النص من الموضوع:

- favor: "{rationale_favor}"
- against: "{rationale_against}"
- none: "{rationale_none}"

قّم مدى توافق كل تفسير مع المعنى الفعلي المقصود في النص، على مقياس من 0 إلى 1 (0 = لا يتوافق إطلاقاً، 1 = يتوافق تماماً).

أخرج فقط بصيغة JSON:

```
{"favor": 0.0, "against": 0.0, "none": 0.0}

```

Figure 3: Stage-B blind verification prompt (original Arabic). The bold first line is the system message; the three Stage-A rationales are inserted and the gold label is withheld. English rendering in Appendix A.2.

even if the target changes.

Stances: 1. Favor: 2. Against: 3. None:

Output only in JSON format:

```
{"favor": "...", "against": "...", "none": "..."}

```

Figure 2 shows the Arabic version used during generation phase. The prompt contains three design constraints. The *steelman requirement* forces the LLM to construct an argument for all three labels rather than revealing only its preferred stance. The *grounding requirement* discourages unsupported information. The *generality requirement* encourages interpretations that remain useful when the target changes.

Parallel prompts of the same structure were also used during development to generate sentiment, style, and sarcasm channels. These auxiliary channels were excluded from the submitted configurations because the corresponding tags were unavailable for the test set.

A.2 Stage B: Blind Argument Verification

The verifier receives the generated arguments but not the gold stance.

System message: “You are a neutral judge. JSON only.”

Text: "{text}"
Target: "{target}"

Note: the text may be MSA or dialectal and may contain sarcasm or implicit negation; judge the actually intended meaning, not the surface phrasing.

The following arguments concern the text’s stance toward the topic:

- favor: “{r_favor}”
- against: “{r_against}”
- none: “{r_none}”

Rate the consistency of each interpretation with the actually intended meaning of the text, on a scale from 0 to 1 (0 = not consistent at all; 1 = fully consistent).

Output only in JSON format:

```
{"favor": 0.0, "against": 0.0, "none": 0.0}
```

Figure 3 shows the Arabic version used during the verification phase. Thresholding the verification scores yields the argumenthood indicators a_k , which are used by the downstream systems as auxiliary information. The highest-scoring argument also defines the verifier opinion $\hat{y}^J$, but this opinion is never used as the gold training label.

Verifier agreement with the gold stance is 0.714 on the training split and 0.725 on the development split.

B Additional Implementation Details

All models use bert-large-arabertv02-twitter. The tweet stream uses a maximum length of 256 tokens, while each argument–tweet pair uses a maximum length of 224 tokens. The learning rate is 10^{-5} and the effective batch size is 32.

Because the models often reach their strongest development performance early in training, checkpoints are evaluated at fixed step intervals rather than only at epoch boundaries. The checkpoint with the strongest development score is retained.

Repeated runs of identical configurations varied by as much as 1.5 F_{avg2} points, motivating our use of three random seeds for development reporting.

For the gated classifier, the verifier-derived mask is applied to the gate. The gate-alignment objective is computed on the unmasked attention weights so that an incorrect verifier decision does not directly penalize the model for assigning weight to the gold argument.

For the pair scorer, the verifier does not determine the training label. Instead, it reduces the contribution of instances for which the generated gold-label argument is judged weak. The submitted configuration additionally removes the training

prior from the pair scores and uses class weighting and minority oversampling.

For the unseen-target gated-classifier submission, training is augmented with additional targets and with a subset of instances in which the original target is replaced by a generic phrase. This augmentation is intended to reduce reliance on target-specific lexical associations.

C Argument-Quality Analysis

We additionally inspect the generated arguments along five dimensions: integrity, length, verifier profile, inter-channel discriminability, and grounding in the original tweet. The verifier score profile is relatively stable between the training and development splits but changes on the held-out test target.

The purpose of this analysis is not to treat the verifier as an external gold standard, but to characterize how stable the generated argument channels remain across splits. This distinction is important because the verifier itself is imperfect and is used only as an auxiliary signal in the submitted systems.

Figure 4 presents two representative errors that illustrate different failure modes. For tweet (ID 2481), the verifier assigns similar scores to FAVOR and AGAINST (0.70 vs. 0.80), indicating low inter-channel discriminability. The final hashtag pushes the prediction toward AGAINST, although the main body of the tweet is largely supportive of continued vaccination. For tweet (ID 3984), the verifier favors AGAINST (0.70), but the selected argument is incorrectly grounded. The expression *كانهم حاطين لي هواء مو تطعيم* describes how quick and painless the injection felt rather than expressing a negative stance toward the vaccine. The closing phrase *الله ينفع به* also expresses hope for its benefit. This illustrates a grounding error in the interpretation of colloquial language.

Tweet: وأن كان تطعيم كورونا سنوي ما المانع من ذلك ؟ قدر الله هذا الامر وأن تعيش هذه الجائحة فلا سبيل بعد توفيق الله الا بإتباع تعليمات الاطباء فإن قرروا الجرعه الثالثه والرابعه حتى وان صارت سنويه فما المانع أن كانت امنه ولا اضرار جانبيه لها ومتوفره بالمجان ؟ #مطعم_جرعتين_وارفض_الثالثه !		
ID: 2481	Gold: FAVOR	Predicted: AGAINST
Channel	Generated Argument	Verifier
FAVOR	كاتب النص يبدو موافقاً على التطعيم السنوي ضد كوفيد-19 إذا كان آمناً ومجائياً، حيث يقول "وإن صارت سنويه فما المانع ان كانت امنه ولا اضرار جانبيه لها ومتوفره بالمجان".	0.70
AGAINST	لكن الكاتب يظهر عدم رغبة واضحة في تلقي الجرعة الثالثة أو الرابعة، كما يشير في نهاية النص #مطعم_جرعتين_وارفض_الثالثه" مما قد يشير إلى معارضته لتلقي المزيد من الجرعات.	0.80
NONE	النص يعبر عن موقف واضح حول التطعيم، وبالتالي قد لا يكون مناسباً تصنيفه بأنه محايد.	0.05
Tweet: تطعيم كورونا كأنهم حاطين لي هواء مو تطعيم ماامداني اجلس الا قالت لي خلصنا سحب الدم احس فيه أكثر،الله ينفع به		
ID: 3984	Gold: FAVOR	Predicted: AGAINST
Channel	Generated Argument	Verifier
FAVOR	الكاتب يشير إلى أنه بعد إتمام عملية سحب الدم، شعر بتأثير التطعيم، مما يدل على أنه يتوقع فائدة من التطعيم رغم الشكوي من عدم الشعور الفوري بتأثيره.	0.05
AGAINST	الكاتب يصف العملية بأنها "كأنهم حاطين لي هواء" مما يشير إلى شكه في فعالية التطعيم وربما عدم رضاه عن العملية.	0.70
NONE	الكاتب يعبر عن شكوكه حول فعالية التطعيم ولكنه ينتظر النتائج، مما يشير إلى عدم وجود موقف واضح أو حاسم تجاه التطعيم.	0.05

Figure 4: Representative qualitative errors for the COVID Vaccine target.