

JLShen at StanceEval-2026: Cross-Lab LLM Ensembles and a Break-Even Rule for the Excluded Class

Junlong Shen

University of Alberta, Edmonton, Canada
junlong6@ualberta.ca

Abstract

Our two systems for StanceEval-2026, Arabic stance detection on Mawqif-XT, placed 1st of 31 on Track 1 ($F_{avg2} = 0.8994$) and 2nd of 22 on Track 2 (0.9357). Track 1 averages six zero-shot readings of each tweet by frontier models from two providers, and calls a fine-tuned Arabic pipeline only to break exact ties. Track 2 is one reader given a short generated description of the target. The leaderboard publishes three aggregate scores over the same predictions, and reading them shows that about 60% of our errors on both tracks are tweets annotated NONE that we placed in a stance class. Because the score ignores that class, the gain from moving such a row is fixed before the row is read. We derive the precision a repair must reach, and nine rankings plus an oracle from the corpus’s own annotation columns all fall short of it.

1 Introduction

Stance detection asks whether a text is in FAVOR of, AGAINST, or expresses NONE towards a target. StanceEval-2026 (Albalawi et al., 2026a) evaluates Arabic stance detection on Mawqif-XT (Albalawi et al., 2026b) in two tracks: Track 1, whose blind test is a target held out from training, and Track 2, whose blind test is two further unseen targets. We entered both as team JLShen.

Neither entry is a fine-tuned Arabic classifier: Track 1 averages six zero-shot readings of each tweet by frontier models from two providers, with a fine-tuned pipeline settling only exact ties, and Track 2 is one reader whose prompt carries a short generated description of the target (§3). The recipe that won the 2024 edition is the weakest component of either.

Both systems then face the same decision: how many tweets to leave in the NONE class. F_{avg2} , the mean of the Favor and Against F1 scores, excludes NONE from the score but not from the data, so a tweet the annotators marked NONE becomes a false

positive inside a counted class once committed to a stance. We can measure that rather than assert it, because the evaluation phase scores the full fixed test set and publishes three aggregates over the same predictions.

Our contributions are:

1. A closed-form decoder from the three published leaderboard columns to any entry’s per-class F1 and error count, and from our own marginals to our exact confusion matrix (§5), which localises our whole deficit on both tracks to the excluded class.
2. A break-even rule for repairing predictions under such a metric, whose detection target is the *union* of labels whose metric delta shares a sign (§3); re-scoring rankings already paid for lifted the best measured precision from 0.333 to 0.500.
3. Nine rankings measured against that bar on the blind test, and an oracle bound from the corpus’s own annotation columns, none of them clearing it (Table 5).

Both ranked entries rebuild bit-exactly, offline and on CPU, from released code.¹

2 Background

Task and data. A system reads one tweet and one target string and emits one of FAVOR, AGAINST, NONE; a submission is one label per line in test-row order. Mawqif-XT is Arabic tweets, mostly Saudi dialect and Modern Standard Arabic, annotated for stance, sentiment and sarcasm. Track 1 releases 3,502 / 619 / 352 train / dev / test rows over four targets, Track 2 2,721 / 1,400 / 644 over five (Table 4, Appendix D). Every Track-1 test tweet carries the target *Women driving*, absent from training, so the nominally seen-target track is

¹<https://github.com/jlshen025/codabench>

near-transfer, and that inverted the ranking of every component we had built.

A released training row. One tweet, its target and its gold labels:

الإلكتروني للدفع الكامل التحول شيء أسوأ

'The worst thing is the complete switch to electronic payment.'
| target Digital transformation | gold NONE, none_reason
Not clear

A plain complaint is still gold NONE, the label the annotators reserve for rows whose stance they could not call. Appendix D adds a FAVOR and an AGAINST row.

Metric. From the shipped scoring program, $F_{avg2} = \frac{1}{2}(F_1^{FAV} + F_1^{AG})$ pooled over all rows, with F_{avg3} over all three classes and accuracy also reported. Our reimplementaion matched the server to four decimals on the development phase.

Related work. Mawqif (Alturayef et al., 2022) introduced target-specific Arabic stance detection with sentiment and sarcasm labels, and StanceEval-2024 (Alturayef et al., 2024) built the first shared task on it, won by AraBERT (Antoun et al., 2020) fine-tuned with multi-task sentiment and sarcasm heads (Badran et al., 2024). We built that recipe first. The 2026 edition adds target shift (Allaway and McKeown, 2020), which it does not survive.

3 System Overview

The rule both systems share. Because F_{avg2} ignores the NONE class, moving a row out of a stance pool changes exactly one F1 term, so the payoff is computable before the row is read. On our final Track-1 entry, moving a row from AGAINST to NONE is worth +0.001274 if it is gold NONE or gold FAVOR, and −0.001616 if it is gold AGAINST. A repair pays only if its precision exceeds

$$p^* = \Delta^+ / (\Delta^- + \Delta^+) = 0.559. \quad (1)$$

So the class to detect is not the excluded class (Figure 1). A gold-FAVOR row in the AGAINST pool is already missing from the FAVOR numerator, so moving it out pays exactly what a gold-NONE row pays. The target is the union *not*-AGAINST, which enlarges the payoff set from 22 to 34 of 187 pool rows and drops the lift a ranking needs from $4.8\times$ to $3.1\times$. Re-scoring submissions already paid for against this label raised the best measured precision from 0.333 to 0.500.

On Track 2 the repairs run the other way, out of the NONE pool into a stance class: flipping k rows

our predicted AGAINST pool, 187 rows




Figure 1: The metric pays the same for two different mistakes, so the detector’s target is their union: exact F_{avg2} deltas for one row moved out of our final Track-1 entry’s AGAINST pool, by that row’s gold label.

of which g are truly class X pays iff $g/k > F_1^X/2$. That threshold *rises* as the system improves, giving an objective stopping condition.

Track 1: a cross-lab voter ensemble. Track 1 averages six one-hot voters and adds a weak arbiter: $\hat{y} = \Pi_N(\arg \max[\frac{1}{6} \sum_v p_v + \epsilon p_{ft}])$ with $\epsilon = 10^{-3}$. Two frontier models from one provider (gpt-5.6-luna, gpt-5.5) read the tweet zero-shot with the target and a three-way instruction, 44 tweets per call, then read it again with $k=12$ demonstrations retrieved by embedding similarity. Two models from a second provider (claude-opus-4-8, claude-sonnet-4.5) read it under the zero-shot prompt. We call that pair *cross-lab* voters: they agree with the first provider’s voters on 0.776–0.812 of rows, below any within-provider variant built by changing the prompt, temperature, batch or reasoning depth.

The fine-tuned Arabic pipeline enters only as p_{ft} , so it decides nothing but exact ties. It is the weakest standalone component here and still the best arbiter available, because ties are settled by annotator convention and it alone was trained on these annotators. Π_N keeps the N rows with the largest NONE margin and commits the rest; N is one scalar tuned on the server.

Track 2: a grounded reader with routing.

Track 2 is one reader. For each target an LLM writes a compact Arabic card holding a definition and six to eight aspects, which goes into the system prompt so the reader can score indirect stance through an aspect. Cards are frozen after generation, so nothing is hand-tuned to a test target. gpt-5.6-sol labels all 644 rows zero-shot at temperature 0 in batches of about 40. The test inputs are also translated to English once and the same reader labels that view; the views agree on 93.7–

96.2% of rows, and each target takes the better-measured one. A stronger reader re-judged contested rows and overrode 73, in directions fixed in advance on the development split.

Resources beyond the training data. Neither ranked entry trains on any external stance corpus. Track 1’s arbiter fine-tunes public checkpoints, MARBERTv2 (Abdul-Mageed et al., 2021), AraBERT-Twitter (Antoun et al., 2020) and Qwen2.5-14B-Instruct (Qwen Team, 2024) with LoRA (Hu et al., 2022), on the organisers’ train and dev splits only. Both entries call proprietary frontier models over an API. An external Arabic stance corpus was tested on Track 2 and discarded.

4 Experimental Setup

Development used the organisers’ splits: Track 1 by cross-validation over train and dev, Track 2 by training on its two seen targets and evaluating on the labelled unseen-target dev split. The evaluation phase returns the exact final metric on the full fixed blind test with no held-out remainder, so we selected operating points and architecture on it directly; every number below is a server score unless marked otherwise, and all we took from the server is aggregate. LLM calls are greedy at temperature 0 with a fixed batch order. Appendix A gives commands, checksums and hyperparameters; Appendix B the prompts verbatim.

5 Results

Official results, and reading the board. Our Track-1 entry scored $F_{avg2} = 0.899400$ ($F_{avg3} = 0.664640$, accuracy 0.857955), ranked 1st of 31 scored entries in evaluation phase 29242; our Track-2 entry scored 0.935700 (0.866986, 0.913043), 2nd of 22 in phase 29244. Ranks are the organisers’ official standings, which leave three scored entries unranked; the boards were read live on 2026-09-07, since scores moved after close.

The board publishes F_{avg2} , F_{avg3} and accuracy over the same predictions. Since F_{avg3} averages the same two F1 terms plus a third,

$$F_1^{\text{NONE}} = 3F_{avg3} - 2F_{avg2}, \quad \text{err} = (1 - \text{acc})n \quad (2)$$

for *any* entry on the board (Table 1). With the gold marginals recovered and our own predicted marginals known locally, the three aggregates pin the diagonal of our own confusion matrix to a

Track 1 ($n=352$)				Track 2 ($n=644$)			
# team	F_{avg2}	$F1_N$	err	team	F_{avg2}	$F1_N$	err
1 JLShen	.899	.195	50	ZHR	.952	.812	43
2 ZHR	.886	.361	56	JLShen	.936	.730	56
3 NAMAA	.872	.273	65	SRK NLP Lab	.928	.695	66

Table 1: Most of our error mass sits in the class the score excludes: on Track 1 we rank first with the *lowest* F1 on NONE of the three ranked teams; on Track 2 the team above us holds a higher one. Official standings of each deciding phase (31 and 22 entries scored), with the F1 of NONE and the error count *decoded* by Eq. 2 from the published triples. Boards read live 2026-09-07 (phases 29242 / 29244); our rows (**JLShen**, bold) are submissions 878207 and 869013.

unique integer solution. Our Track-1 entry predicts 162 / 183 / 7 against gold 158 / 160 / 34, and **30 of its 50 errors are gold-NONE rows inside our stance predictions**. Track 2 predicts 234 / 343 / 67 against gold 216 / 336 / 92, with 34 of its 56 errors the same shape. Each submission is therefore an exact measurement of how many of its edits were right.

Table 2 gives both build-up ladders. Two rungs deserve naming.

The fine-tuned recipe does not survive target shift. The Track-1 pipeline reproducing the 2024 winning approach scores $F_{avg2} = 0.8978$ by cross-validation on seen targets, confirmed to 0.89779 on the development server, and 0.7762 on the held-out target, where a bare zero-shot reader scores 0.8393. Blending it in hurts monotonically (0.7992 at weight 0.25, 0.8215 at 0.40), and Track 2 shows the same ordering. Two or three training targets teach an encoder those targets, not a transferable notion of stance.

Cross-lab voters gave the largest single gain, in one integration form only. Adding the two second-provider models as separate voters was worth +0.0065 (0.8815 $\rightarrow$ 0.8881). Averaging them into one voter gives 0.8816, adding one 0.8787, adding a third 0.8819, and substituting them for the few-shot pair 0.8663. The ensemble is bracketed closed in every direction we varied, and an in-distribution gauge would have vetoed this win (Appendix F).

Track 1’s last two rungs involve no new model. Re-sweeping the NONE count on the changed voter set moved its optimum from 14 to 3, worth +0.0060; the sweep is interior, since $N=0$ scores

Track 1		Track 2	
system	F_{avg2}	system	F_{avg2}
fine-tuned pipeline	0.7762	fine-tuned encoders	0.6695
1 zero-shot LLM	0.8393	Fanar-9B LoRA	0.7425
2 voters + arbiter	0.8723	grounded zero-shot	0.9016
+ NONE op. point	0.8778	+ view routing	0.9169
+ 2 few-shot voters	0.8815	+ prior-calib. repair	0.9298
+ 2 cross-lab voters	0.8881	+ cross-family read	0.9317
+ voter swap	0.8912	+ view-vote ranking	0.9332
+ re-swept N	0.8972	+ annotator simulation	0.9357
+ 4 not-AG flips	0.8994		

Table 2: On both tracks the conventional fine-tuned Arabic recipe is the weakest component, and the last runs cost no new model. Every row is an independent scored submission on the blind test.

0.89322 and $N=8$ scores 0.89500. Moving four more rows out of the AGAINST pool under a tone-based ranking added +0.0022 and decoded to 3 of 4 correct, above the 0.559 bar.

Where the ceiling is. Since Eq. 2 localised our deficit to one class, we built detectors for it: prompts stating the corpus’s abstention rule, graded frontier-model detectors, supervised binary heads, hashtag rules, tone rankings, reader disagreement and a stacker over stored votes. Nine reached the blind test and decoded to precisions of 0.000 to 0.500 against break-evens of 0.514 to 0.559 (Appendix E). None clears its bar.

One mechanism covers the Track-1 failures. Every auxiliary signal validated on the three seen targets transferred at roughly a third of its validated strength to the held-out one. Graded detectors reached 0.50–0.65 on development data and 0.17–0.30 on the blind test, and both supervised binary heads put their weakest leave-one-target-out fold on Women empowerment, the training target closest to the test’s.

On Track 2 the cause is in the corpus rather than the models. Mawqif-XT ships a none_reason column we had not used: of the 392 gold-NONE rows in train and dev, 382 carry a reason, *Not clear* 370 times against *Not Related* 12. The abstain class is 97% “the annotators could not tell” and 3% off-topic, and our displacement, news-sharing and “no evaluative content” detectors all detect the 3%. Mean stance:confidence agrees, 0.919 for FAVOR against 0.734 for NONE. The right question is whether two annotators would disagree, and that reframing is real and still too weak: reader disagreement raises $P(\text{gold} = \text{NONE})$ from 3.8% to 17.6%, yet measured 1 correct row in 10 against

a 0.540 bar, and an oracle with *perfect* sarcasm labels forecasts 0.052.

Error analysis. Three shapes account for the residual. The largest is abstention rows committed to a stance, above. The second is adversarial quotation: 174 of the 352 Track-1 test rows carry the anti-driving campaign hashtag #لن_تقودي, which splits 123 AGAINST to 50 FAVOR because supporters quote it mockingly, while the four rows carrying the pro-driving slogan #قيادة_المراه_للسياره_ضروره are gold AGAINST, as our ensemble called them; a hashtag fixes the topic, never the stance. The third is convention: prompts stating the corpus’s true abstention rule made the readers more faithful to the guideline and cost 2.4 points, because a faithful annotator abstains and the metric never rewards it.

6 Conclusion

Our entries placed 1st of 31 and 2nd of 22, but the transferable part is the analysis. A task publishing several aggregates over one set of predictions hands every participant an exact per-class diagnostic. A metric that drops a class prices every repair before a row is read, against a target that is the union of labels whose delta shares a sign. Whether that class is recoverable from text is the next edition’s question.

Acknowledgments

We thank the StanceEval-2026 organisers for the task, the corpus, and the annotation columns that made §5 possible.

Ethics Statement

Our systems classify public tweets on socially sensitive topics, including women’s rights in Saudi Arabia. Stance classifiers on such topics can misrepresent individuals, particularly for the sarcastic and dialectal text where our systems are weakest, and we claim no fitness for any use beyond this shared task. The data is the organisers’ Mawqif-XT corpus under its research-only terms; the released code redistributes no tweet text. We did not attempt to recover per-example labels of the hidden test set: every submitted candidate had its rows chosen by a reader’s judgement about the released text, never by an optimiser over label assignments.

Limitations

Our operating points and several architectural choices were selected against the evaluation server, which returns the exact final metric on the full fixed test. That is unbiased here, but our scores are not an estimate of generalisation to a further unseen target, and the ≈ 0.004 separating our top few configurations is not a robust architectural preference. Our ceiling statements cover the instruments we built and the variables our oracle tests held fixed, not the task: the leading entries score 0.9906 and 0.9616, so the data supports far more than we reached, and we do not know how. Both systems depend on proprietary API models that change without notice, and one Track-1 voter has since been withdrawn by its provider. We release every voter’s per-row output so the aggregation reproduces without them, but a fresh run of the readers would be a different system. Single seed effects are uncontrolled for the LLM voters.

References

- Muhammad Abdul-Mageed, AbdelRahim Elmadany, and El Moatez Billah Nagoudi. 2021. [ARBERT & MARBERT: Deep bidirectional transformers for Arabic](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7088–7105, Online. Association for Computational Linguistics.
- Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Saad Ezzini, Maram Kurdi, Asma Yamani, Ahmed Ashraf, Maged Al-Shaibani, and Nora Alturayef. 2026a. StanceEval-2026: The second stance detection shared task. In *Proceedings of the Fourth Arabic Natural Language Processing Conference (ArabicNLP 2026)*, Budapest, Hungary. Association for Computational Linguistics.
- Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Maram Kurdi, Saad Ezzini, Asma Yamani, and Ahmed Ashraf. 2026b. [Mawqif-XT: An Arabic benchmark dataset for cross-target stance detection](#). Preprint, arXiv:2608.09539.
- Emily Allaway and Kathleen McKeown. 2020. [Zero-shot stance detection: A dataset and model using generalized topic representations](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 8913–8931, Online. Association for Computational Linguistics.
- Nora Alturayef, Hamzah Luqman, Zaid Alyafei, and Asma Yamani. 2024. [StanceEval 2024: The first Arabic stance detection shared task](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 774–782, Bangkok, Thailand. Association for Computational Linguistics.
- Nora Saleh Alturayef, Hamzah Abdullah Luqman, and Moataz Aly Kamaleldin Ahmed. 2022. [Mawqif: A multi-label Arabic dataset for target-specific stance detection](#). In *Proceedings of the Seventh Arabic Natural Language Processing Workshop (WANLP)*, pages 174–184, Abu Dhabi, United Arab Emirates (Hybrid). Association for Computational Linguistics.
- Wissam Antoun, Fady Baly, and Hazem Hajj. 2020. [AraBERT: Transformer-based model for Arabic language understanding](#). In *Proceedings of the 4th Workshop on Open-Source Arabic Corpora and Processing Tools, with a Shared Task on Offensive Language Detection*, pages 9–15, Marseille, France. European Language Resources Association.
- Mohamed Badran, Mo’men Hamdy, Marwan Torki, and Nagwa El-Makky. 2024. [AlexUNLP-BH at StanceEval2024: Multiple contrastive losses ensemble strategy with multi-task learning for stance detection in Arabic](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 823–827, Bangkok, Thailand. Association for Computational Linguistics.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. [LoRA: Low-rank adaptation of large language models](#). In *International Conference on Learning Representations (ICLR)*.
- Qwen Team. 2024. [Qwen2.5 technical report](#). Preprint, arXiv:2412.15115.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. [Self-consistency improves chain of thought reasoning in language models](#). In *International Conference on Learning Representations (ICLR)*.

A Reproduction

Both released archives rebuild their ranked submission bit-exactly from cached per-row model outputs, offline, on CPU, with numpy alone. They start from stored arrays rather than live API calls because frontier endpoints are not bit-stable over time and one Track-1 voter has since been withdrawn by its provider; a fresh read would be a different system that has to be re-measured rather than compared.

```
# Track 1, submission 878207 (352 labels)
cd stanceeval_t1/code
python build_na.py --k 4 --mode sent \
    --label final
# -> _eval/final.zip
# md5 f6e77e7391fd39e2055a1893548cab40

# Track 2, submission 869013 (644 labels)
```

```
cd stanceeval_t2/code
python reproduce_champion.py
# -> staging/sub_reproduced.zip
# md5 13d07fcb58e4a5eb48bf855ea143eb67
```

The Track-2 script asserts the byte-for-byte match itself and exits non-zero on any mismatch. Each archive also ships the per-row output of every model run as .npz arrays aligned to the released test-row order, with a *_meta.json beside each recording the model identifier, the grounding card, the row count and the wall-clock of that run, and the per-batch call log.

Software. Python 3.11.5, numpy 2.4.2, torch 2.5.1, transformers 5.9.0, scikit-learn 1.8, pandas with keep_default_na=False throughout, since the string "None" is a label and pandas otherwise reads it as a missing value. All API calls go to an OpenAI-compatible /chat/completions endpoint at temperature 0.

Arbiter training (Track 1 only). Four Qwen2.5-14B-Instruct LoRA adapters, two verbalizers (Arabic-letter and Latin-letter) $\times$ two seeds (42, 1): rank 16, $\alpha = 32$, dropout 0.05, attention modules only, learning rate 1×10^{-4} , 4 epochs at a fixed epoch count with no eval-based selection, max length 256, micro-batch 4. Two multi-task encoders (MARBERTv2 and AraBERT-Twitter) with auxiliary sentiment and sarcasm heads: learning rate 2×10^{-5} , batch 32, max length 128, 10 epochs, seed 42. The two families blend 0.15 encoders / 0.85 LoRA. Training used the organisers' train and dev splits only (4,121 rows), about 36 minutes per adapter on one H100 slice.

B Prompts

Whatever is the method ships verbatim. The Track-1 zero-shot voters ran under the following system prompt, with the tweets supplied as a numbered batch of 44:

You are an expert annotator for Arabic stance detection. For each numbered Arabic tweet and its TARGET topic, output the author's stance TOWARD THE TARGET as exactly one of: Favor, Against, None.

- Favor: the author supports, endorses, or is glad about the target.
- Against: the author opposes, criticizes, distrusts, mocks, or fears the target (or its mandate).
- None: no clear stance, off-topic, purely factual/news, or ambiguous.

Sarcasm and rhetorical questions usually signal Against. Judge the AUTHOR's own stance. Output EXACTLY one line per tweet formatted '<number>|<Label>' (Label = Favor, Against, or

None). No other text.

The convention-grounded prompt of §5, written from the corpus's own annotation columns, replaces the three definitions with the following and costs 2.4 points:

- Favor: the author supports the target. In ~95% of Favor cases this is stated EXPLICITLY; implicit support is a rare, marked case.
 - Against: the author opposes the target. In ~85% of Against cases this is stated EXPLICITLY.
 - None: you genuinely CANNOT DETERMINE the author's position toward the target, or the tweet is NOT ABOUT the target. None does NOT mean 'the author is neutral' - it means the stance is NOT RECOVERABLE from the text.
- Decision rule: if an explicit statement of support or opposition is present, label it. If you must rely on a chain of inference, irony, or guessing what the author 'really' means, the annotators most often recorded that as None. Sarcasm is NOT a reliable Against signal (it splits roughly evenly) and sarcastic tweets are 2.5x MORE likely to be labelled None.

The Track-2 reader assembles its system prompt around the frozen target card:

You are an expert annotator for Arabic stance detection. Classify the AUTHOR's stance TOWARD THE TARGET as exactly one of: Favor, Against, None.

TARGET: {target}
MEANING: {card.definition}
RELATED ASPECTS: {card.aspects}
Guidelines:

- Stance may be expressed INDIRECTLY via a related aspect or consequence - that still counts as stance toward the target.
- Favor: the author supports, praises, endorses, defends, or is glad about the target or its aspects.
- Against: the author opposes, criticizes, distrusts, mocks, fears, or rejects the target or its aspects. Sarcasm and rhetorical questions usually signal Against.
- None: use ONLY if the tweet is unrelated to the target, or purely factual/descriptive with no evaluative stance. Do NOT default to None merely because the target is not named literally.

Output EXACTLY one line per tweet formatted '<number>|<Label>'. No other text.

Cards are produced once per target by a separate call that returns a JSON object with an Arabic definition sentence and six to eight comma-separated aspects, and are then frozen. Regenerating them is not safe: one regeneration resolved the school-term target as a pregnancy trimester instead, and a stance-oriented rather than aspect-oriented card for the electric-cars target scored 0.9153 against 0.9169.

C Decoded residuals

The three published aggregates plus our own predicted marginals pin the diagonal of the confusion matrix to a unique integer solution. For Track 2 the off-diagonal retains one free parameter, closed by a single constructed submission that maps all 234 of our own FAVOR predictions to AGAINST; that makes the returned score a bijection of the parameter, and the server’s 0.364732 identifies it exactly. The probe applies one global map to a pool defined by our own predictions and returns one aggregate number, so it identifies no row.

	Track 1	Track 2
test rows	352	644
gold FAV/AG/NONE	158/160/34	216/336/92
our pred FAV/AG/NONE	162/183/7	234/343/67
correct per class	146/152/4	207/323/58
total errors	50	56
abstain rows in a stance pool	30	34

Table 3: Decoded residual of both ranked entries. On both tracks about 60% of the errors are rows the annotators marked NONE.

A matched marginal is not a matched assignment. A convention-anchored Track-2 base predicts 219 / 332 / 93 against gold 216 / 336 / 92, within three rows on every class, and scores 0.88797, which is 0.048 below our entry whose marginals are off by 18, 7 and 25. When a rival places more of a class than you do, that is equally consistent with a threshold gap and with better per-row identification, and only the second is worth chasing. One global-rule submission distinguishes them cheaply.

D Released splits and two further labelled examples

Two more rows from the released split, for contrast with the NONE row in §2. The first carries negative sentiment and no explicit endorsement and is still gold FAVOR, so the corpus resolves ambiguity toward stance rather than away from it.

كاشير تحطها موب فيها تليق بوظايف يكون المرأة تمكين
‘Women’s empowerment means jobs worthy of her, not putting her at a cash register.’ | target Women empowerment | gold FAVOR, favor_reason *F_Implicit*, sentiment Negative

الرجل حساب على لا يكون أن يجب المرأة تمكين
‘Women’s empowerment must not come at the expense of men.’
 | target Women empowerment | gold AGAINST, stance:confidence 0.70, sentiment Negative

split	target	tweets	Fav/Ag/NONE
<i>Track 1 — one held-out test target</i>			
train	Women empowerment	1,190	760/371/59
	Covid vaccine	1,167	509/508/150
	Digital transformation	1,145	879/142/124
dev	the same three targets	619	380/180/59
test†	Women driving	352	—
<i>Track 2 — two unseen test targets</i>			
train	Covid vaccine	1,373	599/598/176
	Digital transformation	1,348	1,035/167/146
dev	Women empowerment	1,400	894/436/70
test†	Electric cars	332	—
	School terms	312	—

Table 4: No test target occurs in its track’s training data, and both tracks train on a FAVOR-heavy distribution their blind tests do not share. Released splits of Mawqif-XT, counted from the organisers’ CSVs. †Test labels are not released; §5 decodes them to 158/160/34 and 216/336/92.

E Instruments against their break-even

Table 5 gives the precision of the rows each ranking moved, against that channel’s break-even. The blind-test rows are exact integer decodes of the submission that deployed the ranking; the oracle rows are upper bounds computed on the labelled split from the annotators’ own columns, and are not comparable to a deployed score.

instrument	prec.	break-even
<i>measured on the blind test</i>		
T1 tone (sentiment) ranking, $k=6$	0.500	0.559
T1 NONE-vote count, $k=18$	0.444	0.559
T1 graded NONE detector, $k=10$	0.400	0.559
T1 campaign-hashtag rule, $k=18$	0.278	0.559
T1 ensemble margin, $k=8$ (control)	0.000	0.559
T2 2-way fresh unanimity, $k=5$	0.400	0.514
T2 target displacement, 3 tiers	0.231	0.540
T2 reader disagreement, $k=10$	0.100	0.540
T2 3-way cross-lab unanimity, $k=4$	0.000	0.524
<i>oracle upper bound, labelled split</i>		
T2 gold stance:confidence<0.8	0.130	0.514
T2 gold sarcasm=Yes	0.052	0.514

Table 5: Every ranking we built falls below the precision the metric demands, and so does an oracle with the annotators’ own judgements.

F Full negative-result ledger

Every entry is a scored submission on the blind test, quoted against the best score at the time it was measured.

Track 1 (contemporaneous best 0.8815, 0.8881 or 0.8972). Blending the fine-tuned pipeline in

at weight 0.25/0.40/1.0: 0.7992/0.8215/0.7762. Convention-grounded prompts as a whole-decision replacement: 0.8574, 0.8664. Re-reading the 71 contested rows with a strong frontier model: 0.8576, 0.8602. Transfer-fit Qwen LoRAs as voters, trained leave-one-target-out and with target-name swapping: 0.8723–0.8787, the best-calibrated and best-decorrelated competent voter we built and still a drag. Per-tweet rather than per-batch demonstration retrieval, which we had believed was a bug: 0.8694; demonstration count $k=6$ 0.8725 and $k=24$ 0.857 against 0.8815 at $k=12$. English translation pivot: 0.8405, 0.8558. Multi-agent debate 0.8542; self-consistency at $K=8$ (Wang et al., 2023) 0.859; knowledge injection of target stance poles 0.8577; anti-abstention prompting 0.8532; batch-size and permutation variants 0.8662–0.8755; high reasoning effort, which agrees with its own base on 0.918 of rows. Alternative arbiters: few-shot 0.8758, encoder ensemble 0.8722, cross-lab 0.8788, mistral-large 0.8562. A seventh voter: 0.8819, 0.89123, 0.88806. Dropping any one voter: 0.87894–0.88516. Newer models from the same second provider: 0.88801, 0.88819. No arbiter at all: 0.89426. Boundary shift by the ensemble’s own margin, as a control for the tone ranking: 0.87285 with precision exactly 0 of 8, which shows the ensemble is confidently rather than marginally wrong about those rows. Non-flagship models mostly emitted degenerate all-NONE output on batched Arabic; two were unusable outright.

An in-distribution gauge mispredicted the largest gain. On the seen target closest to the Track-1 blind test, claude-opus-4-8 scored below the voters it was added to, which would have kept it out of the ensemble under any development-set promotion rule. It produced the run’s largest gain. On a transfer test, an in-distribution gauge should not veto a structurally novel voter.

Arbiter weight is a phantom axis. With k one-hot voters the vote granularity is $1/k$, so every arbiter weight in $(0, 1/6)$ yields byte-identical predictions; we measured 10^{-4} through 0.17 agreeing exactly. Only 0 and weights above 0.2 differ.

Track 2 (champion 0.9357). Higher reasoning effort 0.8985 against 0.9016 greedy. Self-consistency at $T=0.7$ agrees with greedy on 96–98% of rows. A strict-abstention prompt lost 0.0035 on test while gaining 0.0107 on the held-

out proxy target, so the optimal abstention point sign-flips across targets. Few-shot convention demonstrations from disjoint targets changed 8 of 312 rows and lost 0.0027. A third input view, a dialect-to-MSA rewrite, is decorrelated but not better: 0.9125 routed, 0.9156 as a three-view consensus, against 0.9169 for two-view routing. A second adjudication pass over already-adjudicated rows scored 0.0004 below its mechanical control, against +0.0102 for the first pass, so the gain is one-shot rather than iterable; re-running the mechanical control each round is what reveals this. Six convention-anchored readers applied as nested override tiers lost monotonically in dose, -0.0022 at 7 flips through -0.0193 at 44, which is the signature of an anti-informative instrument rather than an uninformative one. Soft fusion of ten test views with a metric-optimal plug-in threshold: 0.9231–0.9321 across six blend weights, all below the champion. Also ruled out on this data: multi-task auxiliary heads, CAMELBERT as a third encoder, entailment templating, knowledge-injection glosses, supervised contrastive loss, constructed synthetic augmentation, external-corpus augmentation both jointly and as intermediate training, a Qwen2.5-32B voter, verbalizer averaging, and per-class threshold calibration. Every encoder variant we tried improved on the seen targets and hurt under leave-one-target-out.

Cross-lab agreement is not free evidence. On Track 2 the pattern inverts the Track-1 result: precision *fell* as reader agreement rose. Two-way fresh unanimity scored 2 of 5, three-way cross-laboratory unanimity 0 of 4, and a two-of-three extension 2 of 13. Readers that share a pre-training prior agreeing is the same evidence counted twice. The Track-1 gain came from cross-provider voters entering a *vote*, not from treating their agreement as a confidence signal.