

SRK AI Lab at StanceEval-2026: Multi-LLM Few-Shot Prompting with Evidence-Guided Auditing for Arabic Stance Detection

V. T. Rushi Kannan¹, Sara Renjit², Hitesh Tekchandani¹

¹Department of Artificial Intelligence and Data Science, Koneru Lakshmaiah University

²Department of Computer Science, Indian Institute of Information Technology Kottayam

¹{2310080010, hitesh.t}@klh.edu.in, ²sararenjit@iiitkottayam.ac.in

Abstract

This paper presents the SRK AI Lab system for Arabic stance detection in the StanceEval-2026 shared task using the Mawqif-XT benchmark. We employ a few-shot multi-LLM framework with GPT-5.5 Instant, Gemini 3.6 Flash with Extended Thinking, and Qwen 3.7 Plus to classify target–post pairs as *Favor*, *Against*, or *None*. The models independently predict stance, followed by GPT-5.5 Instant auditing based on target–post relation, linguistic context, and stance evidence. The audit results are combined with cross-model comparison to select the final label. Our system achieves an official Favg2 of 0.8420 in Track 1 and 0.9192 in Track 2, ranking 8th and 3rd, respectively. These results demonstrate the performance of multi-LLM few-shot prompting with GPT-5.5 auditing for Arabic stance detection.

1 Introduction

Stance detection asks whether a text supports, opposes, or takes no clear position toward a particular target. Unlike general sentiment classification, stance depends on the relationship between the text and its target (Mohammad et al., 2016; AlDayel and Magdy, 2021). This makes Arabic social-media stance detection challenging, as opinions may be implicit and posts often contain informal, dialectal, and context-dependent language (Khouja, 2020). StanceEval-2026 is the second Arabic stance detection shared task and uses the Mawqif-XT benchmark. Given an Arabic social-media post and a target, systems predict one of three labels: *Favor*, *Against*, or *None* (Albalawi et al., 2026a,b). The task covers both seen and unseen targets, providing a setting to study stance classification and cross-target generalization. The implementation and experimental resources of our system are publicly available online.¹

¹https://github.com/Rushikannan2/SRK_AI_Lab_Team

We develop a few-shot multi-LLM framework using GPT-5.5 Instant, Gemini 3.6 Flash with Extended Thinking, and Qwen 3.7 Plus for target-specific stance detection. The framework combines independent model predictions with GPT-5.5 auditing and cross-model comparison to support the final stance decision.

2 Background

Stance detection determines whether a text expresses a *Favor*, *Against*, or *None* position toward a given target (Mohammad et al., 2016; AlDayel and Magdy, 2021). Unlike sentiment analysis, stance depends on the relation between the text and the target. This is particularly challenging for Arabic social media, where stance may be implicit and language is often informal or dialectal (Khouja, 2020). The Mawqif dataset introduced an Arabic target-specific stance resource (Alturayef et al., 2022), followed by StanceEval 2024, the first Arabic stance detection shared task (Alturayef et al., 2024). Recent work has also explored LLMs for Arabic stance detection (Al Hariri and Abu Farha, 2024). StanceEval-2026 uses the Mawqif-XT benchmark to evaluate stance detection on both seen and unseen targets (Albalawi et al., 2026b,a).

Track 1 focuses on three seen targets: COVID-19 Vaccine, Digital Transformation, and Women Empowerment. It contains 3,502 training and 619 development instances, with the training distribution shown in Table 1.

Track 2 tests cross-target generalization. Its training data contain 2,721 instances from COVID-19 Vaccine and Digital Transformation, while the development and evaluation sets include targets not seen during training (Albalawi et al., 2026a). The training distribution is shown in Table 2.

Arabic pretrained models such as ARBERT and MARBERT have shown strong performance across

Target	Train	Favor	Against	None
COVID-19 vaccine	1,167	43.62	43.53	12.85
Digital Transformation	1,145	76.77	12.40	10.83
Women Empowerment	1,190	63.87	31.18	4.96

Table 1: Track 1 training distribution by target and stance (%).

Target	Train	Favor	Against	None
COVID-19 vaccine	1,373	43.63	43.55	12.82
Digital Transformation	1,348	76.78	12.39	10.83

Table 2: Track 2 training distribution by target and stance (%).

Arabic NLP tasks (Abdul-Mageed et al., 2021). More recently, instruction-based LLMs have been applied to Arabic stance detection (Al Hariri and Abu Farha, 2024). Our work follows this direction by combining predictions from multiple LLMs with GPT-5.5 auditing.

3 System Overview

We use a few-shot multi-LLM framework for Arabic target-specific stance detection. GPT-5.5 Instant, Gemini 3.6 Flash with Extended Thinking, and Qwen 3.7 Plus independently predict *Favor*, *Against*, or *None* using task instructions and class-balanced demonstrations, following the in-context learning paradigm (Brown et al., 2020). Prompt-based LLMs have also been explored for Arabic stance detection (Al Hariri and Abu Farha, 2024). The three predictions are passed to a GPT-5.5 audit, and the overall workflow is shown in Figure 1.

GPT-5.5 audits the model predictions using the target–post relationship and relevant linguistic and stance evidence. The audit observations are then considered together with the three predictions through cross-model comparison, allowing both agreement and disagreement to inform the final decision. The same overall workflow is used for Track 1 (seen targets) and Track 2 (unseen targets).

The independent predictions provide complementary evidence, while GPT-5.5 auditing draws on model-based feedback and critique (Madaan et al., 2023; Gou et al., 2024). Unlike multi-agent




Figure 1: Overview of the multi-LLM workflow for seen- and unseen-target stance detection.

debate, the models do not exchange reasoning or engage in iterative discussion. Instead, the final decision considers the three model predictions, audit observations, and evidence from the target–post pair.

4 Experimental Setup

4.1 Dataset Usage

For both tracks, training instances were used to construct the few-shot demonstrations, while development instances were kept separate. No development instances or evaluation instances were used as demonstrations. The demonstrations were fixed before evaluation and remained unchanged during inference. Track 1 evaluates seen targets, whereas Track 2 evaluates targets unseen during training (Albalawi et al., 2026a).

4.2 Preprocessing

For few-shot demonstrations, training instances were formatted as *Target–Post–Stance* examples while preserving the original text. Hashtags, emojis, punctuation, and other textual cues were retained. Evaluation instances were used as provided, without additional preprocessing.

4.3 Prompting Strategy

Zero-shot and few-shot prompting were considered during system development. Zero-shot prompts included the task definition, target, Arabic post, and three stance-label descriptions. Few-shot prompts additionally included class-balanced examples of *Favor*, *Against*, and *None*. Instructions and label descriptions were in English, with the target in English and the post in Arabic.

For the final few-shot configurations, GPT-5.5 Instant and Gemini 3.6 Flash with Extended Thinking used 24 demonstrations (8 per class), while Qwen 3.7 Plus used 150 (50 per class). All demonstrations were drawn from the training data of the corresponding track, fixed before evaluation, and contained no development instances or evaluation instances. Demonstration counts were independent of inference chunk sizes. For Track 2 Gemini, the *Label* field was used as the stance prediction; the other requested fields were retained as auxiliary outputs. Prompt templates are provided in the Appendix.

4.4 Model Configuration

The final system uses GPT-5.5 Instant, Gemini 3.6 Flash with Extended Thinking, and Qwen 3.7 Plus as independent predictors of *Favor*, *Against*, or *None*. Their predictions are generated independently and examined by GPT-5.5 using the target–post relationship and relevant linguistic and stance evidence. The audit does not modify the original model outputs. Audit observations, predictions, and target–post evidence are considered together during cross-model comparison and final decision analysis. The final decision selects the most supported stance label; for unresolved ambiguity, the fixed priority order was Gemini 3.6 Flash with Extended Thinking, GPT-5.5 Instant, then Qwen 3.7 Plus. This priority served as a secondary criterion and did not automatically override the available evidence.

4.5 Audit Procedure

The GPT-5.5 audit was applied to all three model outputs, including agreement and disagreement cases. It receives the target, Arabic post, and predictions with available reasoning or justification, and examines the target–post relationship, author’s stance, and linguistic or contextual evidence such as negation, sarcasm, irony, rhetorical questions, and implicit meaning. The audit may provide a

Track	Favg2	Favg3	Accuracy
Track 1 (Seen)	0.8420	0.6229	0.7813
Track 2 (Unseen)	0.9192	0.8374	0.8882

Table 3: Official StanceEval-2026 results of the final multi-LLM system.

stance assessment, but this is not treated as a fourth independent model prediction; its observations support the cross-model comparison and final decision.

4.6 Implementation and Inference

Evaluation instances were processed in chunks through the respective chat interfaces. For each model, all chunks were processed within the same ongoing chat conversation rather than separate conversations. Track 1 used 50-instance chunks for all models. Track 2 used 50-instance chunks for GPT-5.5 Instant and Qwen 3.7 Plus, and 30-instance chunks for Gemini 3.6 Flash with Extended Thinking. No API- or SDK-based inference pipeline was used.

4.7 Evaluation Metrics

The primary ranking metric is Favg2, the macro-average of F1 for *Favor* and *Against*. Favg3 is the macro-average of F1 over all three classes, including *None* (Albalawi et al., 2026a). Accuracy is also reported. Official submission results are distinguished from post-submission analyses.

5 Results

5.1 Official Results

Our final multi-LLM system combines GPT-5.5 Instant, Gemini 3.6 Flash with Extended Thinking, and Qwen 3.7 Plus. Using the complete prediction, audit, and final-decision workflow described above, the system achieved an official Favg2 of 0.8420 in Track 1 and 0.9192 in Track 2. Favg3 and accuracy are also reported in Table 3 to provide additional performance measures across the three stance labels.

The system performs strongly in both evaluation settings, including the unseen-target setting in Track 2.

5.2 Retained Configuration Results

Selected zero-shot and few-shot configurations with retained standalone evaluation records are reported separately from the final multi-LLM results

Tr.	Config.	Favg2	Favg3	Acc.
1	Qwen – Zero-shot	0.7991	0.6148	0.7472
1	Qwen – Few-shot (150)	0.8174	0.6168	0.7614
1	Gemini – Zero-shot	0.8260	0.6519	0.7670
1	Gemini – Few-shot (24)	0.8366	0.6494	0.7727
Multi-LLM + GPT-5.5 Instant Audit + Final Decision				
1	Decision	0.8420	0.6229	0.7813
2	Gemini – Zero-shot	0.8758	0.7886	0.8354
2	Gemini – Few-shot (24)	0.9185	0.8410	0.8929
Multi-LLM + GPT-5.5 Instant Audit + Final Decision				
2	Decision	0.9192	0.8374	0.8882

Table 4: Selected Track 1 and Track 2 configurations. Fav2, Fav3, and accuracy are reported. The final system combines three LLM predictions with GPT-5.5 auditing and cross-model final decision analysis.

in Table 4. These results provide additional context on individual model configurations evaluated during system development. The final row reports the complete multi-LLM system. Only retained standalone configurations are reported; unavailable standalone results are not inferred or reconstructed.

The selected configurations show that few-shot prompting improves Fav2 for the reported Qwen and Gemini configurations. The final multi-LLM system achieves the highest Fav2 among the configurations shown for both tracks. These results are presented as development observations rather than component-wise ablations, since the individual contributions of cross-model comparison and GPT-5.5 auditing were not separately evaluated.

5.3 Error Analysis

We analyzed representative Track 1 and Track 2 development cases with available gold labels for qualitative error analysis, as test-set gold labels were unavailable. The cases highlight stance challenges involving target, stance-holder, and intended meaning beyond surface sentiment.

Table 5 shows unanimous errors and model-

Tr.	Error Type	G	GPT	Qwen	Gemini	FD
1	Sarcasm / reversal	F	A	A	A	A
1	Speaker attribution	N	A	F	F	F
2	Target scope	F	F	F	A	F
2	Rhetorical ambiguity	N	A	A	A	A

Table 5: Representative development cases showing unanimous errors and model disagreement. G denotes the Gold Label; F, A, and N denote Favor, Against, and None; FD denotes the Multi-LLM Final Decision.

disagreement cases. GPT, Qwen, and Gemini provide individual predictions, while FD denotes the multi-LLM Final Decision based on GPT-5.5 auditing and cross-model analysis. Track 1 includes sarcasm/reversal and speaker attribution cases, while Track 2 includes target scope and rhetorical ambiguity.

6 Conclusion

We presented the SRK AI Lab system for StanceEval-2026, achieving official Fav2 scores of 0.8420 and 0.9192 in Tracks 1 and 2, ranking 8th and 3rd, respectively. The framework combines few-shot multi-LLM predictions, GPT-5.5 auditing, cross-model comparison, and final decision analysis. These results highlight the value of combining complementary model predictions with evidence-guided auditing for target-specific Arabic stance detection. Future work will explore fine-tuned Arabic LLMs and improved prompt optimization.

7 Limitations

The approach may be sensitive to LLM-specific errors in implicit or context-dependent cases. GPT-5.5 serving as both predictor and auditor may introduce dependence, and the audit and final decision stages were not separately evaluated through component-wise ablations.

8 Ethical Considerations

Arabic stance detection may reflect dialectal and demographic biases. LLMs may also misclassify sensitive, implicit, sarcastic, or ambiguous content.

9 Acknowledgments

The authors thank the StanceEval-2026 organizers and Mawqif-XT contributors for their support.

References

- Muhammad Abdul-Mageed, AbdelRahim Elmadany, and El Moatez Billah Nagoudi. 2021. [ARBERT & MARBERT: Deep bidirectional transformers for Arabic](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7088–7105, Online. Association for Computational Linguistics.
- Youssef Al Hariri and Ibrahim Abu Farha. 2024. [SMASH at StanceEval 2024: Prompt engineering LLMs for Arabic stance detection](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 800–806, Bangkok, Thailand. Association for Computational Linguistics.
- Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Saad Ezzini, Maram Kurdi, Asma Yamani, Ahmed Ashraf, Maged Al-Shaibani, and Nora Alturayef. 2026a. [StanceEval-2026: The second stance detection shared task](#). In *Proceedings of the Fourth Arabic Natural Language Processing Conference (ArabicNLP 2026)*, Budapest, Hungary. Association for Computational Linguistics.
- Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Maram Kurdi, Saad Ezzini, Asma Yamani, and Ahmed Ashraf. 2026b. [Mawqif-XT: An Arabic Benchmark Dataset for Cross-Target Stance Detection](#). *Preprint*, arXiv:2608.09539.
- Abeer AlDayel and Walid Magdy. 2021. [Stance detection on social media: State of the art and trends](#). In *Information Processing & Management*, 58(4):102597.
- Nora Alturayef, Hamzah Luqman, Zaid Alyafeai, and Asma Yamani. 2024. [StanceEval 2024: The first Arabic stance detection shared task](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 774–782, Bangkok, Thailand. Association for Computational Linguistics.
- Nora Saleh Alturayef, Hamzah Abdullah Luqman, and Moataz Aly Kamaleldin Ahmed. 2022. [Mawqif: A multi-label Arabic dataset for target-specific stance detection](#). In *Proceedings of the Seventh Arabic Natural Language Processing Workshop (WANLP)*, pages 174–184, Abu Dhabi, United Arab Emirates (Hybrid). Association for Computational Linguistics.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D. Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, and 12 others. 2020. [Language models are few-shot learners](#). In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901. Curran Associates, Inc.
- Zhibin Gou, Zhihong Shao, Yeyun Gong, Yelong Shen, Yujiu Yang, Nan Duan, and Weizhu Chen. 2024. [CRITIC: Large language models can self-correct with tool-interactive critiquing](#). In *The Twelfth International Conference on Learning Representations (ICLR 2024)*, Vienna, Austria.
- Jude Khouja. 2020. [Stance prediction and claim verification: An Arabic perspective](#). In *Proceedings of the Third Workshop on Fact Extraction and VERification (FEVER)*, pages 8–17, Online. Association for Computational Linguistics.
- Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Shashank Gupta, Bodhisattwa Prasad Majumder, Katherine Hermann, Sean Welleck, Amir Yazdanbakhsh, and Peter Clark. 2023. [Self-Refine: Iterative refinement with self-feedback](#). In *Advances in Neural Information Processing Systems*, volume 36, pages 46534–46594. Curran Associates, Inc.
- Saif Mohammad, Svetlana Kiritchenko, Parinaz Sobhani, Xiaodan Zhu, and Colin Cherry. 2016. [SemEval-2016 task 6: Detecting stance in tweets](#). In *Proceedings of the 10th International Workshop on Semantic Evaluation (SemEval-2016)*, pages 31–41, San Diego, California. Association for Computational Linguistics.

A Prompts and Demonstration Details: Track 1

A.1 GPT-5.5 Instant and Gemini 3.6 Flash with Extended Thinking Prompts

We used the same task prompt for GPT-5.5 Instant and Gemini 3.6 Flash with Extended Thinking in the Track 1 few-shot setting. The prompt defines the three stance labels (*Favor*, *Against*, and *None*) and emphasizes target-specific stance, including implicit stance, negation, sarcasm, rhetorical questions, and irony. Both models used 24 class-balanced demonstrations, with eight examples per stance label, provided before the evaluation instances and kept fixed during inference.

The exact prompt is reproduced below.

You are an expert Arabic stance detection system.

Task:

Given an Arabic tweet and a target topic, classify the stance of the tweet toward ONLY the given target.

Possible labels:

- Favor
- Against
- None

Definitions

Favor

- Supports the target.
- Defends the target.
- Encourages or praises the target.
- Expresses positive opinion toward the target.
- Implicit support also counts as Favor.

Against

- Opposes the target.
- Criticizes the target.
- Rejects the target.
- Expresses negative opinion toward the target.
- Implicit criticism also counts as Against.

None

Use None ONLY when:

- The tweet is unrelated to the target.
- The stance cannot be determined.
- The tweet is purely factual.
- The tweet asks a neutral question.
- The tweet discusses both sides without taking a position.
- The target is merely mentioned without expressing support or opposition.

Reasoning Guidelines

Before predicting:

1. Determine the main topic of the tweet.
2. Determine whether the target is actually being discussed.
3. Identify the author's true opinion.
4. Detect negation.
5. Detect sarcasm.
6. Detect rhetorical questions.
7. Detect irony.
8. Ignore emotional intensity.
9. Ignore personal attacks unless they indicate stance.
10. Prefer the IMPLIED stance over literal wording.

Important

Do NOT classify based only on positive or negative sentiment.

Sentiment and stance are different.

Examples:

A positive tweet about an opponent may still be Against.

A negative tweet about an opponent of the target may actually be Favor.

Always reason about the TARGET, not the tweet sentiment.

If evidence clearly supports Favor or Against, never predict None.

Use None only when no stance can reasonably be inferred.

Before producing the final label, internally verify whether the opposite label could also be supported.

If the opposite label is better supported, change the prediction.

If Favor or Against is reasonably supported, avoid predicting None.

Output only the final answer.

Output format

Return ONLY

ID,Label

Example

182,Favor

Do not output explanations.

Do not output reasoning.

Do not output markdown.

The few-shot demonstrations used with GPT-5.5 Instant and Gemini 3.6 Flash with Extended Thinking are available in the project repository:

[Track-1 GPT-5.5 Instant prompt and demonstrations](#)

[Track-1 Gemini prompt and demonstrations.](#)

Now classify the following unseen examples.

Important: - Predict the stance toward the given target only. - Distinguish stance from sentiment. - Infer implicit stance

when supported by the tweet. - Predict None only when no Favor or Against stance can reasonably be inferred. - Return exactly one label for every ID.

Output format:

ID,Label

Now classify the following unseen examples. Use the fixed demonstrations together with the same stance definitions and decision rules used throughout this conversation.

Keep the same task instructions and label definitions.

A.2 Qwen 3.7 Plus Prompt

The following prompt was used for Qwen 3.7 Plus in the Track 1 few-shot setting. It used 150 labeled demonstrations, with 50 examples each for *Favor*, *Against*, and *None*, before the unlabeled evaluation examples.

You are an expert Arabic stance detection researcher participating in an official benchmark.

Task:

Your task is to classify the stance expressed in an Arabic social media post toward a given target.

Possible labels:

- Favor
- Against
- None

Definitions:

Favor:

The author supports, promotes, defends, encourages, praises, or expresses a positive stance toward the target.

Against:

The author opposes, criticizes, rejects, mocks, discourages, or expresses a negative stance toward the target.

None:

The text does not clearly support or oppose the target. It may be:

- factual,
- informational,

- asking a question,
- quoting someone,
- ambiguous,
- mentioning the topic without expressing the author's own stance.

Important guidelines:

1. Focus on the author's opinion, NOT on sentiment alone.
2. Consider sarcasm, irony, rhetorical questions, and implicit meaning.
3. Ignore hashtags, URLs, usernames, and emojis unless they change the stance.
4. Infer the stance from the complete context rather than isolated keywords.
5. The target may never have appeared in training. Generalize using the meaning of the text instead of memorizing specific topics.
6. Do not guess. If the stance is genuinely unclear, choose None.

Demonstrations:

You will first receive labeled demonstrations.

The demonstrations contain 50 examples for each stance label (*Favor*, *Against*, and *None*), for a total of 150 demonstrations.

Study the demonstrations carefully and learn how Favor, Against, and None are distinguished.

Do not answer yet.

Wait until I provide unlabeled test examples.

The few-shot demonstrations used with Qwen 3.7 Plus in Track 1 are available in the project repository. The complete prompt and demonstrations can be accessed here:

[Qwen Track 1 prompt and demonstrations.](#)

Inference Instructions:

You have already studied the labeled demonstrations.

Now classify the following unseen examples using the same decision rules.

Instructions:

1. Carefully read both the Target and the Arabic Tweet.

2. Infer the author's stance toward the given target.
3. Focus on the author's actual opinion, not just sentiment.
4. Handle sarcasm, rhetorical questions, implicit support, and implicit opposition carefully.
5. If the stance is genuinely unclear or only factual, predict None.
6. Be consistent with the demonstrations you have already learned.
7. Do not explain your reasoning.

Output format (exactly):

ID,Label

Example:

101,Favor

102,None

103,Against

Output only the predictions.

B Prompts and Demonstration Details: Track 2

B.1 GPT-5.5 Instant Prompt

The following prompt was used for GPT-5.5 Instant in the Track 2 few-shot setting. It defines the three stance labels and emphasizes target-specific stance, including implicit stance, negation, sarcasm, rhetorical questions, and irony.

You are an expert Arabic stance detection system.

Task:

Given an Arabic tweet and a target topic, classify the stance of the tweet toward ONLY the given target.

Possible labels:

- Favor
- Against
- None

Definitions

Favor

- Supports the target.
- Defends the target.
- Encourages or praises the target.

- Expresses positive opinion toward the target.
- Implicit support also counts as Favor.

Against

- Opposes the target.
- Criticizes the target.
- Rejects the target.
- Expresses negative opinion toward the target.
- Implicit criticism also counts as Against.

None

Use None ONLY when:

- The tweet is unrelated to the target.
- The stance cannot be determined.
- The tweet is purely factual.
- The tweet asks a neutral question.
- The tweet discusses both sides without taking a position.
- The target is merely mentioned without expressing support or opposition.

Reasoning Guidelines

Before predicting:

1. Determine the main topic of the tweet.
2. Determine whether the target is actually being discussed.
3. Identify the author's true opinion.
4. Detect negation.
5. Detect sarcasm.
6. Detect rhetorical questions.
7. Detect irony.
8. Ignore emotional intensity.
9. Ignore personal attacks unless they indicate stance.
10. Prefer the IMPLIED stance over literal wording.

Important

Do NOT classify based only on positive or negative sentiment.

Sentiment and stance are different.

Examples:

A positive tweet about an opponent may still be Against.

A negative tweet about an opponent of the target may actually be Favor.

Always reason about the TARGET, not the tweet sentiment.

If evidence clearly supports Favor or Against, never predict None.

Use None only when no stance can reasonably be inferred.

Before producing the final label, internally verify whether the opposite label could also be supported.

If the opposite label is better supported, change the prediction.

If Favor or Against is reasonably supported, avoid predicting None.

Output only the final answer.

Output format

Return ONLY

ID,Label

Example

182,Favor

Do not output explanations.

Do not output reasoning.

Do not output markdown.

The Track 2 GPT-5.5 Instant prompt was used with 24 few-shot demonstrations, comprising eight examples for each stance label (*Favor*, *Against*, and *None*). The complete prompt and demonstrations are available at [Track 2 GPT-5.5 Instant prompt and demonstrations](#).

The following instruction was then used to classify the unseen examples:

Now classify the following unseen examples.

Use the demonstrations together with the same stance definitions and decision rules used throughout this conversation.

Important:

- Predict the stance toward the given target only.
- Distinguish stance from sentiment.

- Infer implicit stance when supported by the tweet.
- Predict None only when no Favor or Against stance can reasonably be inferred.
- Return exactly one label for every ID.

Output format:

ID,Label

Now classify the following unseen examples. Use the fixed demonstrations together with the same stance definitions and decision rules used throughout this conversation.

Keep the same task instructions and label definitions.

B.2 Gemini 3.6 Flash with Extended Thinking Prompt

The following prompt was used for Gemini 3.6 Flash with Extended Thinking in the Track 2 few-shot setting. It emphasizes author-level stance, target-specific interpretation, and the handling of sarcasm, irony, mockery, rhetorical questions, emojis, and other pragmatic signals. The prompt also requires the model to provide a label, confidence, difficulty, and a concise justification.

You are one of the world's leading Arabic computational linguists and an expert annotator for the MAWQIF / StanceEval 2026 Shared Task.

Your objective is NOT sentiment classification.

Your objective is to determine the AUTHOR'S TRUE STANCE toward the TARGET.

Possible labels ONLY:

Favor

Against

None

CRITICAL RULES

The author's REAL opinion is more important than the literal wording.

Never classify based only on hashtags.

Never classify based only on sentiment.

Always infer the intended stance.

If the author mocks, ridicules, laughs at, or sarcastically imitates one side, then the stance is usually toward the opposite side.

HIGH PRIORITY PHENOMENA

Pay EXTRA attention to:

- Sarcasm
- Irony
- Mockery
- Satire
- Self-deprecating jokes
- Hyperbole
- Rhetorical questions
- Emoji meaning
- Cultural expressions
- Quoted speech
- Repeated hashtags
- Hashtag reversal

Positive wording does NOT always mean Favor.

Negative wording does NOT always mean Against.

EMOJI INTERPRETATION

Treat emojis as evidence.

may completely reverse the intended stance.

OUTPUT FORMAT

Think carefully.

Then output exactly:

Label: Favor / Against / None

Confidence: XX%

Reasoning:

(3–6 concise sentences explaining the linguistic evidence, sarcasm, implied meaning, and why alternative labels were rejected.)

Before assigning a label, explicitly evaluate:

Interpretation A:

The tweet is literal.

Interpretation B:

The tweet is sarcastic.

Interpretation C:

The tweet is mocking the opposite side.

Interpretation D:

The tweet is neutral humor.

Choose whichever interpretation is best supported by the text.

Only then assign the stance.

Before giving your final answer, ask yourself:

"Could this tweet be intentionally saying the opposite of what it literally says?"

If YES:

Re-evaluate the stance.

If sarcasm or mockery better explains the tweet than the literal wording, classify using the implied stance rather than the literal words.

Finally report:

Label:

Confidence:

Difficulty:

Easy

Medium

Hard

Reason:

If Difficulty = Hard, briefly explain what makes the tweet ambiguous (sarcasm, implicit meaning, cultural context, humor, quoting, etc.).

The Track 2 Gemini prompt used 24 few-shot demonstrations, with eight examples each for *Favor*, *Against*, and *None*. The complete prompt and demonstrations are available at [Track-2 Gemini prompt and demonstrations](#).

B.3 Qwen 3.7 Plus Prompt

The following prompt was used for Qwen 3.7 Plus in the Track 2 few-shot setting. It was followed by 150 labeled demonstrations, comprising 50 examples for each stance label (*Favor*, *Against*, and *None*).

You are an expert Arabic stance detection researcher participating in an official benchmark.

Task:

Your task is to classify the stance expressed in an Arabic social media post toward a given target.

Possible labels:

- Favor
- Against
- None

Definitions:**Favor:**

The author supports, promotes, defends, encourages, praises, or expresses a positive stance toward the target.

Against:

The author opposes, criticizes, rejects, mocks, discourages, or expresses a negative stance toward the target.

None:

The text does not clearly support or oppose the target. It may be:

- factual,
- informational,
- asking a question,
- quoting someone,
- ambiguous,
- mentioning the topic without expressing the author's own stance.

Important guidelines:

1. Focus on the author's opinion, NOT on sentiment alone.
2. Consider sarcasm, irony, rhetorical questions, and implicit meaning.
3. Ignore hashtags, URLs, usernames, and emojis unless they change the stance.
4. Infer the stance from the complete context rather than isolated keywords.
5. The target may never have appeared in training. Generalize using the meaning of the text instead of memorizing specific topics.
6. Do not guess. If the stance is genuinely unclear, choose None.

You will first receive labeled demonstrations.

Study them carefully and learn how Favor, Against, and None are distinguished.

Do not answer yet.

Wait until I provide unlabeled test examples.

The Track 2 Qwen prompt and few-shot demonstrations are available in the project repository: [Qwen Track 2 prompt and demonstrations](#).

B.4 GPT-5.5 Instant Audit Prompt

The GPT-5.5 Instant audit procedure was used to independently evaluate the predictions generated by GPT, Qwen 3.7 Plus, and Gemini 3.6 Flash with Extended Thinking for both Track 1 and Track 2. The audit was performed against the original target and Arabic tweet, with the final label determined from the linguistic evidence in the tweet.

B.4.1 Track 1

The following prompt was used for GPT-5.5 Instant to audit the predictions for the Track 1 test examples.

You are GPT Audit, the final expert auditor for Arabic stance detection.

I have predictions from three independent models:

- GPT
- Qwen
- Gemini

for StanceEval-2026 Track 1 and Track 2.

Your job is to audit every prediction against the ORIGINAL Target and Arabic Tweet and determine the most correct final label.

IMPORTANT MODEL PRIOR:

Based on prior development observations, Gemini was given a slight prior when models disagreed.. Therefore, give Gemini a SLIGHTLY HIGHER PRIOR when models disagree.

However:

- Gemini is NOT automatically correct.
- GPT can be correct.
- Qwen can be correct.
- The tweet and target are the ultimate evidence.
- Never choose Gemini simply because it is Gemini.
- Never use majority voting blindly.

Your objective is to produce the BEST POSSIBLE FINAL LABEL, not to maximize agreement with any model.

LABELS**Favor:**

- Supports the target.
- Defends the target.
- Praises or encourages the target.
- Expresses positive stance toward the target.
- Includes implicit support.

Against:

- Opposes the target.
- Criticizes the target.
- Rejects the target.
- Expresses negative stance toward the target.
- Includes implicit opposition.

None:

- Target is unrelated.
- Stance genuinely cannot be determined.
- Tweet is purely factual.
- Neutral question.
- Balanced discussion without taking a position.
- Target is merely mentioned without stance.

IMPORTANT:

Do NOT confuse sentiment with stance.

Always determine:

"What is the author's position toward THIS SPECIFIC TARGET?"

AUDIT PROCESS

For EVERY example:

1. Read the Target.
2. Read the original Arabic Tweet.
3. Read GPT's prediction.
4. Read Qwen's prediction.
5. Read Gemini's prediction.
6. Independently determine the stance from the tweet.
7. Compare your interpretation with all three predictions.
8. Determine whether the current prediction should be KEPT or CHANGED.
9. Assign a difficulty level.
10. Assign confidence.
11. Give a concise but meaningful audit reason.

Pay special attention to:

- Negation
- Sarcasm
- Irony
- Rhetorical questions
- Mockery
- Implicit stance
- Target/opponent relationships
- Positive sentiment toward an opponent
- Negative sentiment toward an opponent
- Tweets where the target is only mentioned
- Tweets that appear factual but contain implicit stance

MODEL DISAGREEMENT POLICY

When all three agree:

- Verify the tweet independently.
- If clearly correct → KEEP.
- If the consensus is wrong → CHANGE.

When GPT + Qwen agree and Gemini disagrees:

- Do NOT automatically choose GPT/Qwen.
- Give Gemini some additional consideration because Gemini has performed best in my testing.
- Then decide strictly from the tweet evidence.

When GPT + Gemini agree and Qwen disagrees:

- Give meaningful weight to GPT+Gemini agreement.
- Still verify the tweet.

When Qwen + Gemini agree and GPT disagrees:

- Give meaningful weight to Qwen+Gemini agreement.
- Still verify the tweet.

When Gemini is the ONLY model predicting a label:

- Carefully inspect the tweet.
- Gemini receives a slight prior advantage because of its observed performance.
- But override Gemini if the actual linguistic evidence supports another label.

When all three disagree:

- Ignore voting.
- Independently classify the tweet from the Target + Tweet.

IMPORTANT:

The models are hypotheses.

The Target + Tweet is the evidence.

DIFFICULTY

Use exactly:

EASY

- Explicit stance.
- Very little ambiguity.
- Example: "I support X" / "X is terrible."

MEDIUM

- Stance is inferable but requires some interpretation.
- Some implicit meaning or contextual understanding.

HARD

- Sarcasm, irony, negation, rhetorical question, implicit stance, conflicting signals, target confusion, or genuine ambiguity.

CONFIDENCE

HIGH:

Very strong evidence for the final label.

MEDIUM:

Reasonably strong evidence but some interpretation is required.

LOW:

Genuine ambiguity or insufficient evidence.

MANDATORY OUTPUT FORMAT

First provide a DETAILED COMPARISON TABLE.

Use EXACTLY these columns:

| ID | Target | Tweet | GPT | Qwen
| Gemini | GPT/Qwen/Gemini Agreement |
Audit Label | Change? | Difficulty |
Confidence | Audit Decision | Reason |

For "GPT/Qwen/Gemini Agreement", use:

- ALL AGREE
- GPT+Qwen
- GPT+Gemini
- Qwen+Gemini
- ALL DIFFER
- OTHER

For "Change?":

- KEEP
- CHANGE

For "Audit Decision":

Use one of:

- KEEP GPT
- KEEP QWEN
- KEEP GEMINI
- KEEP CONSENSUS
- CHANGE TO FAVOR
- CHANGE TO AGAINST
- CHANGE TO NONE
- further REVIEW

The Reason must explain the decision in 1–3 concise sentences.

VERY IMPORTANT TABLE REQUIREMENT

The table must make it immediately obvious:

1. What GPT predicted.
2. What Qwen predicted.
3. What Gemini predicted.
4. Whether they agreed.
5. Which prediction you finally selected.
6. Whether the original prediction should change.
7. How difficult the example is.
8. How confident you are.
9. Why you selected that label.

Do NOT hide the three model predictions inside the Reason column.

AFTER THE TABLE

Provide a SECOND SUMMARY TABLE:

Category	Count
Total Examples	

KEEP	
CHANGE	
EASY	
MEDIUM	
HARD	
HIGH Confidence	
MEDIUM Confidence	
LOW Confidence	
All 3 Agree	
GPT+Qwen Agree	
GPT+Gemini Agree	
Qwen+Gemini Agree	
All 3 Different	

Then provide:

MODEL-WISE AUDIT
| Model | Predictions Retained |
Predictions Overruled | Approx. Audit
Reliability | Main Strengths | Main
Failure Patterns |
| - | - | - | - | - | - |
| GPT | | | | | |
| Qwen | | | | | |
| Gemini | | | | | |

Be HONEST here.

Do not claim numerical accuracy unless it can actually be calculated from the provided evidence.

CHANGE ANALYSIS

Then provide:

CHANGES RECOMMENDED
| ID | Original GPT | Original Qwen
| Original Gemini | Final Label |
Difficulty | Confidence | Why Changed |
| - | - | - | - | - | - | - |

Only include examples where your final audited decision differs from the relevant model prediction or consensus.

Then provide:

GEMINI DISAGREEMENT CASES
List the cases where Gemini disagrees with GPT/Qwen and explain:

- When Gemini appears correct.
- When Gemini appears incorrect.
- When the case is genuinely ambiguous.

This is especially important because Gemini has generally performed best in my previous testing.

FINAL RECOMMENDATION

Finally provide:

FINAL AUDIT RECOMMENDATION
Summarize:

1. Which predictions should be kept.
2. Which predictions should be changed.

3. Which model appears strongest overall based ONLY on the audited examples.
4. Whether Gemini's apparent advantage is supported by these examples.
5. Where GPT is stronger.
6. Where Qwen is stronger.
7. Which examples should receive further REVIEW.
8. Whether the final audited labels appear reliable enough for submission.

IMPORTANT:

Do NOT simply say "Gemini is best."

Actually compare the predictions and evidence.

CRITICAL RULE

Never change a label merely because another model disagrees.

Never keep a label merely because two models agree.

Never automatically trust Gemini.

Never automatically trust GPT.

Never automatically trust Qwen.

Use the original Target + Arabic Tweet as the primary evidence, while giving Gemini a modest prior advantage because of its empirically observed performance.

Be HONEST, CONSERVATIVE, and EXPLICIT.

I will now provide the examples containing:

Target + Arabic Tweet + GPT prediction + Qwen prediction + Gemini prediction.

Audit every example in the required tables.

The complete Track 1 GPT-5.5 Instant audit prompt is available in the project repository: [GPT-5.5 Instant Track 1 audit prompt](#).

B.4.2 Track 2

The same GPT-5.5 Instant audit procedure was applied to the Track 2 predictions using the original target and Arabic tweet as the primary evidence.

You are GPT Audit, the final expert auditor for Arabic stance detection.

I have predictions from three independent models:

- GPT
- Qwen
- Gemini

for StanceEval-2026 Track 1 and Track 2.

Your job is to audit every prediction against the ORIGINAL Target and Arabic Tweet and determine the most correct final label.

IMPORTANT MODEL PRIOR:

Based on prior development observations, Gemini was given a slight prior when models disagreed.. Therefore, give Gemini a SLIGHTLY HIGHER PRIOR when models disagree.

However:

- Gemini is NOT automatically correct.
- GPT can be correct.
- Qwen can be correct.
- The tweet and target are the ultimate evidence.
- Never choose Gemini simply because it is Gemini.
- Never use majority voting blindly.

Your objective is to produce the BEST POSSIBLE FINAL LABEL, not to maximize agreement with any model.

LABELS

Favor:

- Supports the target.
- Defends the target.
- Praises or encourages the target.
- Expresses positive stance toward the target.
- Includes implicit support.

Against:

- Opposes the target.
- Criticizes the target.
- Rejects the target.
- Expresses negative stance toward the target.
- Includes implicit opposition.

None:

- Target is unrelated.
- Stance genuinely cannot be determined.
- Tweet is purely factual.
- Neutral question.
- Balanced discussion without taking a position.
- Target is merely mentioned without stance.

IMPORTANT:

Do NOT confuse sentiment with stance.

Always determine:

"What is the author's position toward THIS SPECIFIC TARGET?"

AUDIT PROCESS

For EVERY example:

1. Read the Target.
2. Read the original Arabic Tweet.
3. Read GPT's prediction.
4. Read Qwen's prediction.
5. Read Gemini's prediction.
6. Independently determine the stance from the tweet.
7. Compare your interpretation with all three predictions.
8. Determine whether the current prediction should be KEPT or CHANGED.
9. Assign a difficulty level.
10. Assign confidence.
11. Give a concise but meaningful audit reason.

Pay special attention to:

- Negation

- Sarcasm
- Irony
- Rhetorical questions
- Mockery
- Implicit stance
- Target/opponent relationships
- Positive sentiment toward an opponent
- Negative sentiment toward an opponent
- Tweets where the target is only mentioned
- Tweets that appear factual but contain implicit stance

MODEL DISAGREEMENT POLICY

When all three agree:

- Verify the tweet independently.
- If clearly correct → KEEP.
- If the consensus is wrong → CHANGE.

When GPT + Qwen agree and Gemini disagrees:

- Do NOT automatically choose GPT/Qwen.
- Give Gemini some additional consideration because Gemini has performed best in my testing.
- Then decide strictly from the tweet evidence.

When GPT + Gemini agree and Qwen disagrees:

- Give meaningful weight to GPT+Gemini agreement.
- Still verify the tweet.

When Qwen + Gemini agree and GPT disagrees:

- Give meaningful weight to Qwen+Gemini agreement.
- Still verify the tweet.

When Gemini is the ONLY model predicting a label:

- Carefully inspect the tweet.
- Gemini receives a slight prior advantage because of its observed performance.
- But override Gemini if the actual linguistic evidence supports another label.

When all three disagree:

- Ignore voting.
- Independently classify the tweet from the Target + Tweet.

IMPORTANT:

The models are hypotheses.

The Target + Tweet is the evidence.

DIFFICULTY

Use exactly:

EASY

- Explicit stance.
- Very little ambiguity.
- Example: "I support X" / "X is terrible."

MEDIUM

- Stance is inferable but requires some interpretation.
- Some implicit meaning or contextual understanding.

HARD

- Sarcasm, irony, negation, rhetorical question, implicit stance, conflicting signals, target confusion, or genuine ambiguity.

CONFIDENCE

HIGH:

Very strong evidence for the final label.

MEDIUM:

Reasonably strong evidence but some interpretation is required.

LOW:

Genuine ambiguity or insufficient evidence.

MANDATORY OUTPUT FORMAT

First provide a DETAILED COMPARISON TABLE.

Use EXACTLY these columns:

ID	Target	Tweet	GPT	Qwen
Gemini	GPT/Qwen/Gemini Agreement	Audit Label	Change?	Difficulty
Confidence	Audit Decision	Reason		

For "GPT/Qwen/Gemini Agreement", use:

- ALL AGREE
- GPT+Qwen
- GPT+Gemini
- Qwen+Gemini
- ALL DIFFER
- OTHER

For "Change?":

- KEEP
- CHANGE

For "Audit Decision":

Use one of:

- KEEP GPT
- KEEP QWEN
- KEEP GEMINI
- KEEP CONSENSUS
- CHANGE TO FAVOR
- CHANGE TO AGAINST
- CHANGE TO NONE
- further REVIEW

The Reason must explain the decision in 1–3 concise sentences.

VERY IMPORTANT TABLE REQUIREMENT

The table must make it immediately obvious:

1. What GPT predicted.
2. What Qwen predicted.
3. What Gemini predicted.
4. Whether they agreed.

5. Which prediction you finally selected.
6. Whether the original prediction should change.
7. How difficult the example is.
8. How confident you are.
9. Why you selected that label.

Do NOT hide the three model predictions inside the Reason column.

AFTER THE TABLE

Provide a SECOND SUMMARY TABLE:

Category Count
- - :
Total Examples
KEEP
CHANGE
EASY
MEDIUM
HARD
HIGH Confidence
MEDIUM Confidence
LOW Confidence
All 3 Agree
GPT+Qwen Agree
GPT+Gemini Agree
Qwen+Gemini Agree
All 3 Different

Then provide:

```
### MODEL-WISE AUDIT
| Model | Predictions Retained |
Predictions Overruled | Approx. Audit
Reliability | Main Strengths | Main
Failure Patterns |
| - | - : | - : | - | - | - |
| GPT | | | | | |
| Qwen | | | | | |
| Gemini | | | | | |
```

Be HONEST here.

Do not claim numerical accuracy unless it can actually be calculated from the provided evidence.

CHANGE ANALYSIS

Then provide:

```
### CHANGES RECOMMENDED
| ID | Original GPT | Original Qwen |
Original Gemini | Final Label |
Difficulty | Confidence | Why Changed |
|
| - | - | - | - | - | - | - |
```

Only include examples where your final audited decision differs from the relevant model prediction or consensus.

Then provide:

GEMINI DISAGREEMENT CASES

List the cases where Gemini disagrees with GPT/Qwen and explain:

- When Gemini appears correct.
- When Gemini appears incorrect.
- When the case is genuinely ambiguous.

This is especially important because Gemini has generally performed best in my previous testing.

FINAL RECOMMENDATION

Finally provide:

FINAL AUDIT RECOMMENDATION

Summarize:

1. Which predictions should be kept.
2. Which predictions should be changed.
3. Which model appears strongest overall based ONLY on the audited examples.
4. Whether Gemini's apparent advantage is supported by these examples.
5. Where GPT is stronger.
6. Where Qwen is stronger.
7. Which examples should receive further REVIEW.
8. Whether the final audited labels appear reliable enough for submission.

IMPORTANT:

Do NOT simply say "Gemini is best."

Actually compare the predictions and evidence.

CRITICAL RULE

Never change a label merely because another model disagrees.

Never keep a label merely because two models agree.

Never automatically trust Gemini.

Never automatically trust GPT.

Never automatically trust Qwen.

Use the original Target + Arabic Tweet as the primary evidence, while giving Gemini a modest prior advantage because of its empirically observed performance.

Be HONEST, CONSERVATIVE, and EXPLICIT.

I will now provide the examples containing:

Target + Arabic Tweet + GPT prediction + Qwen prediction + Gemini prediction.

Audit every example in the required tables.

The complete Track 2 GPT-5.5 Instant audit prompt is available in the project repository: [GPT-5.5 Instant Track 2 audit prompt](#).

Development notes. The materials in this directory document earlier prompt and demonstration experiments conducted during system development. Demonstration sets were iterated using training data only. The final demonstration sets were selected before official evaluation and were kept fixed throughout evaluation.