

CUNI at AlexandriaX-2026: Finetuning LLMs for dialectal Arabic Translation

Josef Jon, Ondřej Bojar

Charles University, Faculty of Mathematics and Physics
{jon,bojar}@ufal.mff.cuni.cz

Abstract

We describe the Charles University (CUNI) submission to the AlexandriaX-2026 Subtask 1, a shared task on English-to-Arabic dialectal conversation translation covering 13 Arabic varieties. Our approach fine-tunes Gemma-4-31B-it with QLoRA on the Alexandria training data at the conversation level, using rich contextual prompts that include domain metadata, participant names and speaker labels. We explore multiple training data mixes and hyperparameters, yielding a pool of candidate translations for the test set. These are reranked with a genetic-algorithm-enhanced minimum Bayes risk (GA-MBR) decoding using sentence-level spBLEU as the fitness function. Fine-tuning improves over the unmodified base model by approximately 5.2 ChrF points on the development set (from 45.2 to 50.4 ChrF for the best single checkpoint, or 25.1 to 32.6 spBLEU); GA-MBR reranking over a 40-candidate development pool yields 51.3 ChrF / 34.0 spBLEU, a further improvement of 0.9 ChrF points over the best pool member. Our submission ranked second in BLEU/chRF++ the unconstrained track.

1 Introduction

The AlexandriaX-2026 Challenge Subtask 1 (El Mekki et al., 2026) is a shared task on translating English conversations into 13 Arabic dialectal varieties (listed in Table 1). The task uses the UBC-NLP/Alexandria dataset (EL Mekki et al., 2026), which provides multi-turn bilingual conversations in English and one of these varieties.

There are two aspects of the task that should be kept in mind by the participants. First, translations must preserve conversational coherence across turns: a good translation of one turn may depend on what was said in previous turns, and also on the people involved in the conversation, their relationship, and the domain and topic. All of this

additional information is part of the dataset. Second, dialectal appropriateness is important — the same English sentence should be rendered differently depending on the target variety.

Our submission builds on QLoRA fine-tuning and MBR reranking in combination with a genetic algorithm to modify the sentences in order to improve automatic scores. The main evaluation metric of the task is spBLEU with a single, human-made reference. This makes the conclusions based on the final ranking weaker (compared to human evaluation), as differences in the score might not reflect an actual translation quality change.

The paper is organized as follows. Section 2 describes the task, dataset, and related work. Section 3 presents our system. Section 4 covers the experimental setup. Section 5 reports results.

2 Background

2.1 Task Setup

The AlexandriaX-2026 Subtask 1 is a shared task on translating English conversations into 13 Arabic dialectal varieties.

Evaluation. The evaluation metrics are spBLEU (SentencePiece BLEU with the flores200 model; Goyal et al. 2022; Papineni et al. 2002; Team et al. 2022) and character n -gram F-score. ChrF, i.e., word n -gram order equal to 0 is used for our experiments; ChrF++, i.e. word n -gram order equal to 2, is used in the official evaluation (Popović, 2015).

Tracks. The challenge offers constrained and unconstrained tracks distinguished by allowed training data and model size. Our submission participates in the **unconstrained** track: we fine-tune Gemma-4-31B-it (Team, 2026), which exceeds the 5B parameter limit of the constrained track.

2.2 Datasets

Alexandria. The Alexandria dataset (EL Mekki et al., 2026) consists of multi-turn conversations in English, generated by an LLM, paired with their dialectal Arabic translation created by humans, across 13 target varieties. Conversations are short (2–4 turns) and annotated with a domain label (e.g. *business, technology, healthcare*) and participant names. The dataset is split into training, development, and public/private test portions; Table 1 summarises the sizes. We use the development split (3,963 conversations, 12,250 turns) for all reported evaluations.

WMT23 news data. We created a synthetic dialectal Arabic dataset by translating the English source sides of WMT23 test sets (Kocmi et al., 2023). We translated each English sentence into all 11 task dialects using Gemma-4-{12,31}B-it, Jais-2-70B-Chat and gpt-5.4-nano, then scored every translation with a GEMBA-style judge (Kocmi and Federmann, 2023) (Gemma-4-31B-it rating translation accuracy and dialect authenticity on a 0–100 scale) and retained only the best translation for each source sentence, and only if the translation scored ≥ 90 on both dimensions. Because the same model (Gemma-4-31B-it) acts as one of the translators and the judge, the filter may over-select outputs produced by this model.

Synthetic conversations. We generated 43,978 synthetic bilingual conversations using GPT-5.4-mini, seeded from Alexandria metadata (domain labels, participant roles, turn counts). These synthetic conversations cover also the Libyan and Sudanese dialects, that are excluded from the original Alexandria training dataset.

Dialect	Train		Dev		Public Test	
	Conv	Turns	Conv	Turns	Conv	Turns
Egyptian (EG)	982	3,108	352	1,113	366	1,118
Jordanian (JO)	1,730	5,501	352	1,113	347	1,107
Lebanese (LB)	2,915	8,906	371	1,118	369	1,106
Libyan (LY)	—	—	—	—	374	1,109
Moroccan (MA)	815	2,573	354	1,110	353	1,115
Mauritanian (MR)	1,748	5,515	352	1,114	357	1,112
Omani (OM)	2,066	6,280	381	1,109	387	1,118
Palestinian (PS)	4,669	14,933	352	1,110	370	1,109
Saudi (SA)	2,699	8,470	363	1,110	351	1,113
Sudanese (SD)	—	—	—	—	345	1,106
Syrian (SY)	1,869	6,071	347	1,119	353	1,114
Tunisian (TN)	665	2,034	378	1,116	379	1,109
Yemeni (YE)	988	3,089	361	1,118	355	1,106
Total	21,146	66,480	3,963	12,250	4,706	14,442

Table 1: Alexandria dataset split sizes by dialect. Libyan and Sudanese Arabic appear only in the test set.

2.3 Related Work

Arabic dialectal machine translation. Translating into Arabic dialects has received growing attention alongside the development of dialectal corpora such as MADAR (Bouamor et al., 2018). Most prior work focuses on dialect identification (Salameh et al., 2018; Abdelali et al., 2021) or translation *from* dialects into MSA or English (Zbib et al., 2012; Sajjad et al., 2013), rather than translation *into* dialects. The Alexandria Challenge targets the harder direction of producing dialectally appropriate output from English. More recently, large language models finetuned for Arabic have pushed quality further (Nagoudi et al., 2022a,b; Anwar et al., 2026; Jon et al., 2026).

Minimum Bayes risk decoding. Minimum Bayes risk (MBR) decoding selects the hypothesis that maximises expected utility under a distribution approximated by a set of samples (Kumar and Byrne, 2004). In case of reference-based metric used as utility, it selects the hypothesis that scores the best using all the other hypotheses as the references. Applied to MT, MBR consistently outperforms MAP (beam search) decoding (Eikema and Aziz, 2020; Freitag et al., 2022). An extension of this idea is to use a genetic algorithm to increase the number of high-quality candidate translations (Jon and Bojar, 2023).

3 System overview

3.1 Model Selection

We evaluated three Arabic-capable instruction-tuned models as base systems without fine-tuning (*Base* columns in Table 2): Gemma-4-31B-it and Gemma-4-12B-it (Team, 2026), and Jais-2-8B-Chat (Anwar et al., 2026). Gemma-4-31B-it and Gemma-4-12B-it score similarly out of the box (45.2 and 45.4 ChrF respectively), while Jais-2-8B-Chat lags at 33.1 ChrF, with near-zero score on Mauritanian Arabic (2.4 ChrF). We selected Gemma-4-31B-it as our primary model.

3.2 Conversation-level Translation

We translate each conversation as a single unit rather than turn by turn. The model receives all English turns in one prompt and outputs all Arabic turns in one pass, giving it access to full discourse context.

Rich context prompting. We prepend a context header to each conversation and label every turn,

based on the additional info provided in the input:

You are a professional, very precise translator and a native speaker. Translate inputs based on the instructions and always print out only the text of the best possible translation, no explanations. Keep the same formatting (e.g. markup, lines, spacing) as the original. Do not translate untranslatable parts of the input (URLs, code and similar).

Translate the following text into {dialect} dialect, only print out the translation, do not add any explanations:

[Domain:{domain}]|Participants:{participants}

{1st_speaker}: {turn_1}

{2nd_speaker}: {turn_2}

...

3.3 Fine-tuning

We fine-tuned the models with (Q)LoRA (Dettmers et al., 2023) with multiple hyperparameter configs: (LoRA rank $r \in \{8, 32, 64, 128\}$, $\alpha=2r$, full bfloat16 or 4-bit quantisation, $lr \in \{2e-5, 6.36e-5, 1e-4\}$). Training used effective batch size 64 (4×16 gradient accumulation steps), maximum sequence length 768 tokens, and up to 8,000 steps with checkpoints evaluated every 150 steps on a 150 dev set conversations per dialect. All 11 dialects were included simultaneously.

Training data. We experimented with five dataset mixes: alx (Alexandria training split only); alx+wmt23 (Alexandria augmented with GEMBA-filtered WMT23 dialectal translations); alx+synth (Alexandria augmented with the synthetic GPT-5.4-mini conversations); alx+wmt23 (Alexandria augmented with GEMBA-filtered WMT23 dialectal translations and the MADAR dataset) and alx+test (Alexandria train + public test set, to see if more in-domain data improve the scores, was not used in the final submission).

3.4 GA-MBR Ensemble

We further increased the evaluation scores by using genetic-algorithm minimum Bayes risk decoding (GA-MBR; Jon and Bojar 2023). We constructed a pool of candidate translations from multiple checkpoints and data configurations and combined them to obtain better MBR scores.

Each candidate is treated as a GA gene. The fitness function is sentence-level spBLEU (Goyal et al., 2022), computed against all remaining candidates as pseudo-references. The candidate translations are mutated by insertion, deletion or replacement operations and crossed-over with each other over multiple generations.

4 Experimental Setup

Preprocessing. Conversations are formatted as described in Section 3, loss is computed only on the Arabic output tokens.

Software and hardware. Training uses transformers (Wolf et al., 2020), PEFT (Mangrulkar et al., 2022), and TRL (von Werra et al., 2020) on a single NVIDIA GPU (either A100, A40 or H100, depending on the node). Translations are generated with vLLM (Kwon et al., 2023). 4-bit QLoRA quantisation used bitsandbytes (Dettmers et al., 2023). Evaluation uses sacrebleu (Post, 2018).

Decoding and post-processing. At inference time we use greedy decoding with thinking mode disabled (enable_thinking=False). The model is prompted to produce output in the same numbered-turn format as the training target.

GA-MBR As the search for optimal hyperparameters for the GA would be too resource-consuming, we set the mutation rate to 2 (2 tokens of each candidate on average are mutated) and crossover probability to 0.2 for all the experiments. The GA runs for 200 generations for the dev set experiments and 140 generations for the final submission. The GA operates independently at the *turn* level: each source turn is processed as a separate instance with all candidate translations of that turn as the initial population and pseudo-references (each initial candidate is copied 20 times to increase the population size, resulting in population sizes of 800 and 960 for our experiments). Candidates are tokenized into flores200 SentencePiece subword pieces (the same tokenizer used by spBLEU), which form the genes for crossover and mutation. Crossover performs single-point splicing of two parent gene sequences; mutation substitutes individual subword tokens drawn from the union of all candidates’ subword vocabularies. No semantic, fluency, or adequacy constraint beyond spBLEU fitness is applied. The final output for each turn is the best-scoring evolved sentence after all generations; this sentence may be a novel combination of subwords not present verbatim in any of the original candidates. Turns across a conversation are assembled independently, so different turns may originate from different candidates or from evolved hybrids.

We also experimented with conversation-level processing for the GA, where we concatenated the

whole conversation and treated it as a single gene for scoring and mutation/crossover, but this approach yielded worse final scores in our preliminary experiments.

The 48 test-set candidates came from multiple checkpoints of Gemma-4-31B-it fine-tuned on alx, alx+synth, alx+wmt23 or alx+wmt23+madar at various learning rates and steps, selected based on dev set chrF scores, see Table 6 in Appendix B. The development-set ablation pool contained 40 candidates selected by the same key.

Evaluation metrics. We report **ChrF** (Popović, 2015) and **spBLEU** (corpus-level BLEU with SentencePiece flores200 tokenisation).

5 Results

Base model performance. Gemma-4-12B-it and Gemma-4-31B-it score nearly identically without fine-tuning (45.4 vs. 45.2 ChrF overall). Jais-2-8B-Chat scores substantially lower (33.1 ChrF), with particularly low scores on Mauritanian (2.4 ChrF) and Omani (20.6). These low ChrF values show the model is producing outputs far from the references for these varieties, possibly due to limited exposure to them in pre-training.

Effect of fine-tuning. QLoRA fine-tuning yields consistent ChrF gains across all three models. Jais-2-8B-Chat shows the largest score increase overall (33.1 → 46.7 ChrF, +13.6), with especially large gains on Mauritanian (+30.3) and Omani (+28.1) — varieties where the base model produced near-zero scores. Gemma-4-31B-it gains 5.2 ChrF points (45.2 → 50.4), while Gemma-4-12B-it gains 3.9 points (45.4 → 49.3). The fine-tuned Gemma-4-31B-it reaches the highest ChrF overall (50.4 / 32.6 spBLEU), though the margin over Gemma-4-12B-it is small (1.1 ChrF), which may suggest that data quality and quantity have a stronger effect than model size in this range.

Other datasets. None of the additional datasets led to an increase in dev set spBLEU scores over the best single checkpoint trained only on the Alexandria dataset, the detailed results are shown in Table 5 in Appendix A. Some of the checkpoints trained either on alx+synth or alx+wmt23 configuration were however part of the top 48 candidates in the final test set pool for the GA-MBR.

Cross-dialect variation. Syrian Arabic consistently receives the highest scores across all fine-

tuned systems (53–56 ChrF), while Mauritanian Arabic scores lowest (30–38 ChrF). Whether this reflects low translation, low reference quality, or distribution mismatch between the training data and these varieties cannot be determined from automatic scores.

GA-MBR. Table 3 shows the effect of GA-MBR reranking over a 40-candidate development pool. MBR (equivalent to GA-MBR with $g=0$) already yields a higher score than the highest-scoring individual pool member (33.9 vs. 32.6 spBLEU). Evolutionary optimisation for 200 generations adds a marginal further increase (34.0 spBLEU), suggesting that score improvement from the GA step is small relative to MBR pool diversity.

Test set results. On the private test set (Table 4), the best single checkpoint scores 31.42 spBLEU. MBR reranking over a 48-candidate pool raises this to 32.47, and GA-MBR optimisation for 140 generations adds a further 0.05 points (32.52). The candidate checkpoints were selected based on their dev set ChrF scores, the list of the checkpoints is shown in Appendix B, Table 6.

6 Conclusion

We presented the CUNI submission to the AlexandriaX-2026 Subtask 1 on English-to-dialectal Arabic conversation translation. Our system combines QLoRA fine-tuning of Gemma-4-31B-it in conversation-level translation mode with a GA-MBR over a pool of candidate translations.

Fine-tuning with Alexandria dataset yields a 5.2 ChrF score increase over the base model (45.2 → 50.4). GA-MBR reranking over 40 development candidates adds a further 0.9 ChrF (51.3), with the gain largely attributable to MBR pool diversity rather than the evolutionary search itself. Mauritanian Arabic consistently receives the lowest automatic scores. Our final submission was ranked second based on spBLEU.

Limitations

Our evaluation relies primarily on automatic reference-based metrics, with only a single human reference per source, and does not include human evaluation of adequacy, fluency, or dialectal authenticity. Consequently, improvements in spBLEU or ChrF may not always correspond to improvements in translation quality.

Dialect	G4-12b				G4-31b				Jais-2-8b-chat			
	Base		SFT		Base		SFT		Base		SFT	
	ChrF	spBLEU	ChrF	spBLEU	ChrF	spBLEU	ChrF	spBLEU	ChrF	spBLEU	ChrF	spBLEU
EG	50.7	31.8	53.2	36.9	50.6	31.6	53.4	37.4	43.1	25.8	50.4	34.0
JO	50.4	29.7	54.0	37.1	50.1	29.3	55.5	38.5	39.8	20.4	51.8	34.5
LB	46.6	26.8	49.7	32.9	47.1	26.7	50.0	33.4	35.5	18.3	46.8	29.3
MA	40.8	21.1	43.7	25.3	40.8	21.1	45.1	26.9	30.8	14.5	41.7	23.5
MR	30.6	9.4	37.2	16.6	30.5	9.1	37.9	15.9	2.4	0.5	32.7	10.6
OM	43.2	21.8	49.2	28.6	41.9	20.4	53.8	35.9	20.6	7.1	48.7	29.5
PS	49.5	29.6	52.4	34.8	49.0	29.1	53.3	36.0	36.0	18.4	50.3	32.1
SA	50.7	29.5	54.4	36.0	49.7	28.6	55.7	37.7	44.5	23.7	51.4	32.5
SY	52.7	32.1	55.4	38.3	51.2	30.7	55.6	38.8	48.2	28.6	52.9	35.5
TN	40.9	22.1	47.4	29.9	42.8	24.6	48.2	30.8	35.1	16.7	44.9	26.6
YE	44.7	23.9	46.4	27.3	44.6	23.8	47.6	28.7	31.2	13.1	44.6	25.0
ALL	45.4	25.3	49.3	31.3	45.2	25.1	50.4	32.6	33.1	16.4	46.7	28.1

Table 2: Per-dialect ChrF and spBLEU on the Alexandria development set for both the base models and the best finetuned checkpoints for each base model (selected by an overall chrF on a subset of the dev set, per-dialect best scores might be higher).

System	ChrF	spBLEU
Base G4-31B-it	45.20	25.07
SFT G4-31B-it	50.42	32.61
MBR $g=0$ (40 dev candidates)	51.18	33.87
GA-MBR $g=200$ (40 dev candidates)	51.26	33.96

Table 3: Development-set ablation (11 dialects, 12,250 turns). The pool contains 40 candidates from multiple checkpoints and data configurations, none trained on the test set. GA-MBR $g=0$ is MBR initialisation without evolutionary optimisation.

System	spBLEU
Best single checkpoint	31.42
MBR ($g=0$, 48 candidates)	32.47
GA-MBR* ($g \sim 100$, 48 candidates)	32.5
GA-MBR ($g=140$, 48 candidates)	32.52

Table 4: Private test set spBLEU scores. The candidate pool contains 48 translations from the best checkpoints across alx, alx+synth, alx+wmt23 and alx+wmt23+madar data configurations. Asterisk shows the final submission, where the GA process has not fully finished and ran for around 100 generations for most of the examples.

Acknowledgments

This work was partially supported by SVV project number 260 821, by the EU Horizon project ARIEN (Grant agreement ID: 101121329) and by Czech Ministry of Education, Youth and Sports (grant MŠMT OP JAK Mezišektorová spolupráce CZ.02.01.01/00/23_020/0008518).

References

Ahmed Abdelali, Hamdy Mubarak, Younes Samih, Sabit Hassan, and Kareem Darwish. 2021. [QADI](#):

[Arabic dialect identification in the wild](#). In *Proceedings of the Sixth Arabic Natural Language Processing Workshop*, pages 1–10, Kyiv, Ukraine (Virtual). Association for Computational Linguistics.

Mohamed Anwar, Abed Alhakim Freihat, George Ibrahim, Mostafa Awad, Abdelrahman Sadallah, Gurpreet Gosal, Gokulakrishnan Ramakrishnan, Sarath Chandran, Biswajit Mishra, Rituraj Joshi, Ahmed Frikha, Etienne Goffinet, Abhishek Maiti, Ali El Filali, Sarah AlBarri, Samujjwal Ghosh, Rahul Pal, Parvez Mullah, Awantika Shukla, and 41 others. 2026. [Jais 2: A family of arabic-centric open large language models](#). *Preprint*, arXiv:2608.13580.

Houda Bouamor, Nizar Habash, Mohammad Salameh, Wajdi Zaghouani, Owen Rambow, Dana Abdulrahim, Ossama Obeid, Salam Khalifa, Fadhil Eryani, Alexander Erdmann, and Kemal Oflazer. 2018. [The MADAR Arabic dialect corpus and lexicon](#). In *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan. European Language Resources Association (ELRA).

Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. Qlora: Efficient finetuning of quantized llms. *arXiv preprint arXiv:2305.14314*.

Bryan Eikema and Wilker Aziz. 2020. [Is MAP decoding all you need? the inadequacy of the mode in neural machine translation](#). In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 4506–4520, Barcelona, Spain (Online). International Committee on Computational Linguistics.

Abdellah El Mekki, AbdelRahim A. Elmadany, Samar M. Magdy, Saad Ezzini, Mo El-Haj, Mustafa Jarrar, Zaid Alyafeai, Bernard Ghanem, and Muhammad Abdul-Mageed. 2026. AlexandriaX 2026: The First Shared Task on Dialectal Arabic Machine Translation. In *Proceedings of The Fourth Arabic Natural Language Processing Conference: Shared Tasks*, Budapest, Hungary. Association for Computational Linguistics.

- Abdellah EL Mekki, Samar M. Magdy, Houdaifa Atou, Ruwa AbuHweidi, Baraah Qawasmeh, Omer Nacar, Thikra Al-hibiri, Razan Saadie, Hamzah A. Alsayadi, Nadia Ghezael Hammouda, Alshima Mohammed Alkhazimi, Aya Hamod, Al-Yas Yaqoob Al-Ghafri, Wesam El-Sayed, Asila Ismail al Sharji, Mohamad Ballout, Anas Belfathi, Karim Ghaddar, Serry Sibae, and 28 others. 2026. [Alexandria: A multi-domain dialectal Arabic machine translation dataset for culturally inclusive and linguistically diverse LLMs](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 32567–32592, San Diego, California, United States. Association for Computational Linguistics.
- Markus Freitag, David Grangier, Qijun Tan, and Bowen Liang. 2022. [High quality rather than high model probability: Minimum Bayes risk decoding with neural metrics](#). *Transactions of the Association for Computational Linguistics*, 10:811–825.
- Naman Goyal, Cynthia Gao, Vishrav Chaudhary, Peng-Jen Chen, Guillaume Wenzek, Da Ju, Sanjana Krishnan, Marc’ Aurelio Ranzato, Francisco Guzmán, and Angela Fan. 2022. [The Flores-101 evaluation benchmark for low-resource and multilingual machine translation](#). *Transactions of the Association for Computational Linguistics*, 10:522–538.
- Josef Jon and Ondřej Bojar. 2023. [Breeding machine translations: Evolutionary approach to survive and thrive in the world of automated evaluation](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2191–2212, Toronto, Canada. Association for Computational Linguistics.
- Josef Jon, Rawan Bondok, and Ondřej Bojar. 2026. [Current state of LLMs for Arabic dialectal machine translation](#). In *Proceedings of the 2nd Workshop on NLP for Languages Using Arabic Script*, pages 329–363, Rabat, Morocco. Association for Computational Linguistics.
- Tom Kocmi, Eleftherios Avramidis, Rachel Bawden, Ondřej Bojar, Anton Dvorkovich, Christian Federmann, Mark Fishel, Markus Freitag, Thamme Gowda, Roman Grundkiewicz, Barry Haddow, Philipp Koehn, Benjamin Marie, Christof Monz, Makoto Morishita, Kenton Murray, Masaaki Nagata, Toshiaki Nakazawa, Martin Popel, and 3 others. 2023. [Findings of the 2023 conference on machine translation \(WMT23\): LLMs are here but not quite there yet](#). In *Proceedings of the Eighth Conference on Machine Translation*, pages 1–42, Singapore. Association for Computational Linguistics.
- Tom Kocmi and Christian Federmann. 2023. [Large language models are state-of-the-art evaluators of translation quality](#). In *Proceedings of the 24th Annual Conference of the European Association for Machine Translation*, pages 193–203, Tampere, Finland. European Association for Machine Translation.
- Shankar Kumar and William Byrne. 2004. [Minimum Bayes-risk decoding for statistical machine translation](#). In *Proceedings of the Human Language Technology Conference of the North American Chapter of the Association for Computational Linguistics: HLT-NAACL 2004*, pages 169–176, Boston, Massachusetts, USA. Association for Computational Linguistics.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the ACM SIGOPS 29th Symposium on Operating Systems Principles*.
- Sourab Mangrulkar, Sylvain Gugger, Lysandre Debut, Younes Belkada, Sayak Paul, Benjamin Bossan, and Marian Tietz. 2022. PEFT: State-of-the-art parameter-efficient fine-tuning methods. <https://github.com/huggingface/peft>.
- El Moatez Billah Nagoudi, AbdelRahim Elmadany, and Muhammad Abdul-Mageed. 2022a. [AraT5: Text-to-text transformers for Arabic language generation](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 628–647, Dublin, Ireland. Association for Computational Linguistics.
- El Moatez Billah Nagoudi, AbdelRahim Elmadany, and Muhammad Abdul-Mageed. 2022b. [TURJUMAN: A public toolkit for neural Arabic machine translation](#). In *Proceedings of the 5th Workshop on Open-Source Arabic Corpora and Processing Tools with Shared Tasks on Qur’an QA and Fine-Grained Hate Speech Detection*, pages 1–11, Marseille, France. European Language Resources Association.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. [Bleu: a method for automatic evaluation of machine translation](#). In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318, Philadelphia, Pennsylvania, USA. Association for Computational Linguistics.
- Maja Popović. 2015. [chrF: character n-gram F-score for automatic MT evaluation](#). In *Proceedings of the Tenth Workshop on Statistical Machine Translation*, pages 392–395, Lisbon, Portugal. Association for Computational Linguistics.
- Matt Post. 2018. [A call for clarity in reporting BLEU scores](#). In *Proceedings of the Third Conference on Machine Translation: Research Papers*, pages 186–191, Brussels, Belgium. Association for Computational Linguistics.
- Hassan Sajjad, Kareem Darwish, and Yonatan Belinkov. 2013. [Translating dialectal Arabic to English](#). In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 1–6, Sofia, Bulgaria. Association for Computational Linguistics.

Mohammad Salameh, Houda Bouamor, and Nizar Habash. 2018. [Fine-grained Arabic dialect identification](#). In *Proceedings of the 27th International Conference on Computational Linguistics*, pages 1332–1344, Santa Fe, New Mexico, USA. Association for Computational Linguistics.

Gemma Team. 2026. [Gemma 4 technical report](#). *Preprint*, arXiv:2607.02770.

NLLB Team, Marta R. Costa-jussà, James Cross, Onur Çelebi, Maha Elbayad, Kenneth Heafield, Kevin Hefernan, Elahe Kalbassi, Janice Lam, Daniel Licht, Jean Maillard, Anna Sun, Skyler Wang, Guillaume Wenzek, Al Youngblood, Bapi Akula, Loic Barrault, Gabriel Mejjia Gonzalez, Prangthip Hansanti, and 20 others. 2022. [No language left behind: Scaling human-centered machine translation](#). *Preprint*, arXiv:2207.04672.

Leandro von Werra, Younes Belkada, Lewis Tunstall, Edward Beeching, Tristan Thrush, Nathan Lambert, Shengyi Huang, Kashif Rasul, and Quentin Galouédec. 2020. [TRL: Transformers Reinforcement Learning](#).

Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Remi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, and 3 others. 2020. [Transformers: State-of-the-art natural language processing](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 38–45, Online. Association for Computational Linguistics.

Rabih Zbib, Erika Malchiodi, Jacob Devlin, David Stalard, Spyros Matsoukas, Richard Schwartz, John Makhoul, Omar F. Zaidan, and Chris Callison-Burch. 2012. [Machine translation of Arabic dialects](#). In *Proceedings of the 2012 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 49–59, Montréal, Canada. Association for Computational Linguistics.

A Different dataset mixes

Table 5 compares the four training data configurations for Gemma-4-31B-it. The overall scores are close: alx, alx+wmt23, and alx+synth all fall within 0.3 ChrF of each other (50.4, 50.2, and 50.2 respectively), with no configuration consistently dominating at the dialect level. alx achieves the highest aggregate ChrF and performs best on Omani, Saudi, and Palestinian varieties; alx+wmt23 leads on Egyptian, Syrian, and Tunisian; and alx+synth scores highest on Lebanese. Adding the public test set (alx+test) reduces development set performance by about 0.6 ChrF.

B Final candidate pool

Table 6 shows the composition of the initial list of translations (checkpoints used to produce these translations) used for the final GA-MBR for the submission.

Dialect	alx		alx+wmt23		alx+synth		alx+test	
	ChrF	spBLEU	ChrF	spBLEU	ChrF	spBLEU	ChrF	spBLEU
EG	53.4	37.4	53.8	37.7	53.4	37.4	53.7	37.5
JO	55.5	38.5	55.2	38.3	55.0	38.2	55.4	38.9
LB	50.0	33.4	49.4	32.3	50.9	34.5	50.0	33.0
MA	45.1	26.9	45.2	26.8	44.8	26.5	45.6	27.3
MR	37.9	15.9	37.5	16.2	37.5	16.5	36.7	15.6
OM	53.8	35.9	51.3	32.1	51.5	32.5	49.0	28.1
PS	53.3	36.0	52.8	35.1	53.1	35.7	53.2	35.4
SA	55.7	37.7	54.8	36.7	55.0	36.9	54.5	35.8
SY	55.6	38.8	56.4	38.9	55.5	38.4	55.6	38.2
TN	48.2	30.8	49.4	32.3	48.6	32.0	48.6	31.3
YE	47.6	28.7	47.5	28.4	47.6	28.7	46.8	27.7
ALL	50.4	32.6	50.2	32.3	50.2	32.6	49.8	31.7

Table 5: Per-dialect ChrF and spBLEU on the Alexandria development set (11 dialects, LY and SD excluded) for Gemma-4-31B-it fine-tuned on four training data configurations: alx (Alexandria training split only), alx+wmt23 (augmented with GEMBA-filtered WMT23 dialectal translations), alx+synth (augmented with GPT-5.4-mini synthetic conversations), and alx+test (Alexandria train + public test set, included for completeness — see text). Each column shows the best checkpoint per configuration. Bold marks the highest value in each row.

Data mix	LR	Steps
alx	5e-5 8e-5	450, 600, 600, 600, 600, 750, 900, 900, 1050, 1050, 1200, 1350, 1500 250, 600
alx+synth	8e-5 9e-5	300, 300, 350, 900 300, 450
alx+wmt23	5e-5 8e-5	1650, 2100, 2400, 3000, 3900, 4200, 5400, 5550, 5700, 5850, 6000, 6150, 6300, 6450, 6600, 6750, 6900, 7050, 7200, 7350, 7500, 7650, 7800, 7950 1650
alx+wmt23+madar	8e-5	1500, 1650

Table 6: All 48 candidates in the test-set GA-MBR pool (Gemma-4-31B-it). +synth adds GPT-4-mini synthetic conversations; +madar adds the MADAR corpus. Multiple entries with the same LR come from different training runs. All candidates use $r = 128$.