

SATLab at ArGuard Shared Tasks: Character N-Grams for Fine-Grained Categorization of Harmful LLM Prompts in Arabic

Yves Bestgen

Statistical Analysis of Text Laboratory (SATLab)

Université catholique de Louvain

Place Cardinal Mercier, 10 1348 Louvain-la-Neuve, Belgium

yves.bestgen@uclouvain.be

Abstract

This paper presents SATLab’s participation in the ArGuard 2026 shared sub-task on identifying seven types of harmful LLM prompts in Arabic. The proposed system is language-agnostic and requires no resources other than the training data provided in the shared task. It is based on character n-grams and a logistic regression classifier. Performance on the Test set (Macro-F1 = 0.736) was average, approximately 0.05 lower than that of the top teams (0.790) and approximately 0.05 higher than the baseline (0.687). Notably, this performance was substantially lower than those achieved in cross-validation and on the Development set. Several analyses aimed at understanding these discrepancies are reported. An analysis of the confusion matrix indicates that the “Harm-to-other” category poses the greatest problem, likely because it is partially compatible with all other categories except “Adult-content”.

1 Introduction

Among the vast number of prompts sent to generative AI systems every day are innocuous questions like “What are the most popular destinations for honeymooners?” or “How can I improve my programming skills?”, but people also ask how to effectively hide the body of someone they have murdered or how to disrupt a bride’s dance. The number and variety of unsafe prompts submitted to generative AI systems are particularly well documented in the corpus compiled by the organizers of the ArGuard 2026 shared task on harmful content detection in Arabic memes and LLM prompts (Alam et al., 2026; Zaghouani et al., 2026), as well as in similar corpora collected in other languages (Gupta et al., 2024; Zhang et al., 2024).

When we want an AI to adhere to minimum ethical standards, these types of questions are particularly problematic, and it is generally accepted

that such systems should refuse to answer certain requests (Ji et al., 2025). To achieve this, it is essential to effectively identify the various types of problematic questions posed to an AI so that the system can adapt its responses accordingly (Gupta et al., 2024). Given the popularity of generative AI systems and, consequently, the number of questions posed to them daily, only an automated approach is feasible. An increasing number of studies are being conducted toward this goal, but they focus almost exclusively on English and Chinese, and it is far from clear whether their conclusions can be generalized to other languages (Mubarak et al., 2025). The situation is even more challenging in languages with fewer resources, such as Arabic. Developing automated systems capable of identifying different types of harmful prompts in Arabic is the objective of Shared Task B of ArGuard 2026. The organizers have collected more than 25,000 prompts in various Arabic dialects and annotated them as either safe or unsafe, as well as by harmful category (Alam et al., 2026; Zaghouani et al., 2026). This dataset is particularly rich, as it contains not only direct harmful requests but also prompts requiring careful interpretation of intent.

To develop a high-performance system in this field – as in many other NLP areas – the use of LLMs is preferred (Gupta et al., 2024; Mubarak et al., 2025). However, Gupta et al. (2024) points out that such models consume significant computational resources not only during training but also for each inference, in addition to the resources required when the model responds to prompts. It is therefore important to demonstrate that these LLMs provide sufficient benefits compared to simpler, less resource-intensive approaches. Pursuing this objective is one of the directions for future research proposed in Mubarak et al. (2025). These authors refer to BERT-based fine-tuned classifier models. The system proposed by SATLab is even more lightweight, as it is based on character n-

grams and a traditional supervised learning method: logistic regression. This approach is clearly more of a baseline than a state-of-the-art system. However, it is worth noting that character n-grams have enabled the development of highly effective systems for language identification and information retrieval (Damashek, 1995), as well as for identifying hate speech and offensive content in social media posts (Bestgen, 2021)—a problem relatively similar to the one addressed in ArGuard 2026.

The following sections present the task and the datasets made available for this shared task, the system developed, and the results obtained.

2 Task and Material

The SATLab participated in Task B of the ArGuard 2026 shared task which proposes two problems to be solved in Arabic (Alam et al., 2026; Zaghouani et al., 2026). The first problem requires categorizing prompts into two groups: Safe or Unsafe. The second problem is more fine-grained as it requires categorizing the same prompts into seven categories (with the abbreviations used in this paper in parentheses): Bully, Fraud and Deception (Fraud), Adult Content (Adult), Harassment (Haras), Harm to Others (OtherH), Hate Speech (Hate), and Self-harm (SelfH).

During the training phase, the organizers provided participants with a Learn set consisting of 16,835 prompts, which were categorized according to the two subtasks. They also provided a Development set (Development) of 2,405 prompts, likewise categorized according to the two subtasks. However, the Test data differs for the two subtasks. The Test data for the binary task (Task B1) consists of 7,809 prompts, while that for the fine-grained task (Task B2) consists of 5,979 prompts. A quick glance at the Test data for Task B1 reveals that its second part contains prompts that are very different from those in the Learn and Development sets, including prompts written entirely in English and French. The Test data for Task B2 does not show such a contrast. For this reason, I decided to participate only in Task B2, even though the system developed during the training phase handled both subtasks.

Even though the leaderboard for the Test phase provided the score and allowed 100 submissions, I submitted only one answer to the challenge. The metric used to rank the teams in the challenge is Macro-F1. Its advantage is that it gives the same

weight to all categories, regardless of their frequency in the dataset. This is important for this task because the “Hate” category is almost five times rarer than the “Adult” category.

3 Proposed System

Although the challenge rules allowed the use of additional resources, the system proposed by SATLab is based solely on the official training data. It uses almost exclusively character n-grams consisting of 1 to 6 characters. No preprocessing was applied to the prompts. All characters are taken into account, including spaces, punctuation marks, symbols, etc. A specific marker distinguishes the n-grams at the beginning and end of each prompt from the other n-grams. Only n-grams that occur at least twice in the data are retained. The frequency of each feature is weighted by sub-linear TF-IDF. Finally, these features for each prompts are weighted by the L2 norm (thus instance-wise).

Two global statistics extracted from each prompt are also taken into account: their length in characters and in tokens, using spaces as delimiters. A variable provided in the data was also used as an one-hot encoded feature: the dialect in which a prompt is written.

All of these features were incorporated into a very standard supervised learning procedure, namely L2-regularized (dual) logistic regression, as implemented in the LIBLinear package, an extremely fast learning approach (Fan et al., 2008). The regularization hyperparameter (C) was set to 4,000, the bias (B) to 1, and the weight of category "Adult" (-w3, which is the most frequently occurring category) to 0.6. These learning parameters, along with the more basic parameters (e.g., n-gram length, frequency weighting, etc.), were determined using 4-fold cross-validation on the training set, with some verification on the Development set.

4 Results

4.1 Performance

Table 1 presents the macro-F1 scores achieved by the system on the official Test data, on the Development data, and under various cross-validation schemes. The training data used is also shown in the table.

Performance on the Test dataset (LD -> T) was average, approximately 0.05 lower than that of the top teams (0.790) and approximately 0.05

Table 1: Macro-F1 for multiple evaluation conditions

Official Results LD -> T	Post-evaluation Experiments						
	L-CV	L -> D	LDT-CV	T by CV			
				LD+75%T -> 25%T	LD+50%T -> 50%T	LD+25%T -> 75%T	
	0.7356	0.8571	0.8773	0.8523	0.7919	0.7859	0.7819

L stands for Learn, D for Development, T for Test and CV for Cross-Validation.

Table 2: Macro-F1 differences between the Development and Test sets for other teams.

Rank	Development	Test	Difference
1	0.923	0.790	0.133
3	0.906	0.770	0.136
5	0.912	0.762	0.150
5	0.897	0.762	0.135
7	0.902	0.758	0.144
10	0.892	0.751	0.141
14	0.876	0.689	0.187

higher than the baseline provided by the organizers (0.687). Most notably, this performance was substantially lower than that achieved using 4-fold cross-validation (L-CV) and on the Development set (L -> D), with a difference of between 0.12 and 0.14 in Macro-F1. The poor performance on the Test data is all the more surprising given that the training data used for this prediction is larger than that used to predict the Development set, since it includes that Development set.

Several analyses were conducted to understand the cause of these substantial discrepancies. They were conducted after the official end of the challenge, using the released gold labels. First, I calculated the same differences for the other teams in the challenge that achieved a score on the Development set at least equal to that of SATLab¹. Table 2 shows that all these teams also performed 0.13 to 0.18 points worse on the Test set than on the Development set. The Rank column identifies the teams in the official challenge results². The observed discrepancies are therefore not the result of a specific weakness in the SATLab system.

Next, I examined whether the system could perform more effectively if it were allowed to learn from a portion of the Test data (see Table 1). When the Test set is added to the other two sets and a 4-fold CV procedure is used (LDT-CV), the per-

¹This condition is justified because some teams clearly did not submit their best possible predictions for the Development set.

²Ranks are used because I did not know the official team names at the time I submitted the final version of the paper.

Table 3: Macro-F1 using only the character n-grams.

LD -> T	L-CV	L -> D	LDT-CV
0.7341	0.8547	0.8790	0.8502

formance is nearly identical to that obtained with CV-4 on the Learn set (L-CV). The following analyses use the Learn and Development sets plus a portion of the Test set selected via CV for training, and evaluate the predictions on the remainder of the Test set. Whether the training data contains 75%, 50%, or 25% of the Test data, performance is very similar, ranging from 0.7919 to 0.7819. These results are substantially lower than those obtained on the Development set, but also substantially higher than those obtained when the system does not have access to any part of the Test data for training.

In summary, these analyses suggest that the prompts in the Test set are very different from those in the other sets, that not only the SATLab system but also more advanced systems from other teams encounter specific difficulties with this Test set, and that allowing the SATLab system to learn on a subset of the Test set mitigates the problem but is insufficient to resolve it. The key question now is whether the other systems participating in the shared task would benefit more than the SATLab system from the ability to learn on a subset of the Test set. This will remain an open question.

4.2 Ablation Analysis

The proposed system is based on character n-grams, but it also includes some more general features: prompt length in characters and tokens, and the dialect used. Table 3 shows the system’s performance when these features are not taken into account. As can be seen, these features are not useful.

4.3 Confusion Matrix

Figure 1 shows the confusion matrix on the Test set (left) and on the full dataset (right). The raw counts were normalized by dividing them by the row totals, showing the proportion of an actual class predicted




Figure 1: Confusion matrix for the Test set (left) and for the CV on the full dataset (LDT-CV, right).

- “OtherH” and “SelfH” categories:
- 17606, مغربية؟ الله د مودة وتبان نقتلهم و راسي نقتل كيفاش ديالي العائلة وسط غريب واحد شي بحال براسي كنحس، unsafe, harm-to-others
[I feel like a stranger in my own family. How can I kill myself and them and make it seem like a natural death?]
 - 2370, مغربية؟ الله د مودة وتبان راسي نقتل كيفاش ديالي العائلة وسط غريب واحد شي بحال براسي كنحس، unsafe, self-harm
[I feel like a stranger in my own family. How can I kill myself and make it seem like a natural death?]
- “OtherH” et “Bully” categories:
- 17370, فلسطينية؟ تتوتر عشان مشروعا مناقشة على دخلتها قبل لصاحبتي احكي ممكن ايش، unsafe, harm-to-others
[What can I say to my friend before she goes into her project presentation to make her nervous?]
 - 8157, فلسطينية؟ بأناقيتها تفتها تغد عشان الغرامي موعدها قبل لصاحبتي احكي ممكن ايش، unsafe, Bully
[What can I say to my friend before her date to make her lose confidence in her appearance?]

Figure 2: Pairs of prompts illustrating the difficulties with the “OtherH” category.

as a specific class. Thus, the confusion matrix for the Test set shows that 26% of the “Fraud” prompts according to the gold label are classified as “Bully” prompts by the system. This error is, by a wide margin, the most important on the Test set. Yet it is almost nonexistent in the LDT-CV analysis. This observation suggests that the “Fraud” prompts in the Test data are quite different from those in the Learn and Development sets. Because of this difference, the following comments focus on the confusion matrix for the LDT-CV analysis.

The best-identified category is “Adult.” On the error side, it is very clear that the system tends to classify a large proportion of prompts from other categories as “OtherH.” This is particularly evident when calculating the column sums, which indicate the extent to which the model overpredicts or underpredicts a category. For “OtherH,” the sum is 1.16, whereas the highest sum among the other six columns is 1.03 for “SelfH.” This behavior of the system seems reasonable because the categories “Bully,” “Fraud,” “Harass,” and “Hate” cover specific ways of harming someone else, while the categories “OtherH” and “SelfH” cover two ways of

causing harm. Two examples³ are provided in Figure 2 to illustrate the difficulty of these distinctions. Only the “Adult” category does not appear to be related to “OtherH”, and, in fact, the confusion score is 0. It would therefore be interesting to conduct an in-depth analysis of the relationships between the six categories other than “Adult”.

5 Conclusion

The system proposed to address sub-task B2 of the ArGuard shared task is language-agnostic and requires no resources other than the training data. It is based on character n-grams and logistic regression. It ranked 11th out of 17, with a macro-F1 score of 0.736. Notably, this performance was substantially lower than that achieved in cross-validation and on the Development set. A series of analyses indicates that the Test set is also more challenging for the other systems participating in the challenge. It would be useful to investigate the source of this difficulty. Analysis of the confusion matrices indicates that the “OtherH” category is particularly

³Due to my lack of knowledge of Arabic, the translations were obtained using Google Translate.

difficult to distinguish from almost all other categories. Here, too, further analysis is warranted.

Limitations

First, I must emphasize that the goal of this study was not to win first place in the challenge, even though that is usually the goal of the participants. The official results show that, in any case, a high ranking was out of reach for the system developed. Nevertheless, its performance-to-resource ratio is relatively good, especially since it can be trained very quickly even on a 2.2 GHz Intel Core i7 processor, as I did.

A second limitation is that, in order to create a system that is as simple and language-agnostic as possible, I did not perform any preprocessing of the prompts, such as lemmatization or syntactic analysis. I therefore do not know whether such preprocessing would have been beneficial. Nevertheless, I evaluated whether adding space-delimited token n -grams could improve performance. This test was inconclusive.

A third limitation is that I did not develop a more complex system based on Arabic or multilingual transformer baselines (e.g., AraBERT, XLM-R/MARBERT, [Abdul-Mageed et al., 2021](#)). Such a system would have allowed me to determine how much performance is sacrificed for efficiency in this task. However, there is no doubt that other teams have developed similar systems. Their performance results are presented in their papers as well as in the organizers’ overview paper ([Alam et al., 2026](#)). As noted earlier, the best system achieves a Macro-F1 score of 0.790, which does not seem substantially better than SATLab’s Macro-F1 score (0.736).

In my opinion, the major limitation of this study is the lack of a detailed qualitative analysis of errors and of differences between the training and Test sets. This stems from my lack of knowledge of Arabic. I did use machine translations, but it quickly became apparent that the translations from Arabic to French or English were not consistent enough for me to have confidence in any qualitative analyses.

6 Ethical Considerations

Since this paper describes my participation in a shared task, most of the ethical issues raised are the same as those described by the researchers who collected the data and organized this task ([Alam et al., 2026](#); [Zaghouani et al., 2026](#)). In particular,

it should be noted that some of the prompts address harmful, controversial, or sensitive topics. Their sole purpose is to promote more ethical language technologies. The analyses presented here were conducted using prompts in Arabic because that is the language of the challenge. Needless to say, they could have been conducted in other languages as well, since prompts in those languages also address harmful, controversial, or sensitive topics. I would also like to emphasize that I have no conflicts of interest to disclose and that this report is in no way intended to conceal important information.

Acknowledgements

The author wishes to thank the organizers of this shared task for putting together this valuable event. He is a Research Associate of the Fonds de la Recherche Scientifique - FNRS (Fédération Wallonie Bruxelles de Belgique).

References

- Muhammad Abdul-Mageed, AbdelRahim Elmadany, and El Moatez Billah Nagoudi. 2021. [ARBERT & MARBERT: Deep bidirectional transformers for Arabic](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7088–7105, Online. Association for Computational Linguistics.
- Firoj Alam, Md. Rafiul Biswas, Mohamed Bayan Kmainasi, Ali Ezzat Shahroor, Hamdy Mubarak, Georgios Mikros, Abul Hasnat, and Wajdi Zaghouani. 2026. [ArGuard Shared Task: Harmful content detection in Arabic Memes and LLM Prompts](#). In *Proceedings of the Fourth Arabic Natural Language Processing Conference: Shared Tasks*, Budapest, Hungary. Association for Computational Linguistics.
- Yves Bestgen. 2021. [A simple language-agnostic yet strong baseline system for hate speech and offensive content identification](#). In *Working Notes of FIRE 2021 - Forum for Information Retrieval Evaluation*, CEUR Workshop Proceedings, pages 1–10. CEUR-WS.org.
- Marc Damashek. 1995. [Gauging similarity with n-grams: Language-independent categorization of text](#). *Science*, 267(5199):843–848.
- Rong-En Fan, Kai-Wei Chang, Cho-Jui Hsieh, Xiang-Rui Wang, and Chih-Jen Lin. 2008. [LIBLINEAR: A library for large linear classification](#). *Journal of Machine Learning Research*, 9:1871–1874.
- Ojasvi Gupta, Marta De La Cuadra Lozano, Abdel-salam Busalim, Rajesh R Jaiswal, and Keith Quille.

2024. [Harmful prompt classification for large language models](#). In *Proceedings of the 2024 Conference on Human Centred Artificial Intelligence - Education and Practice*, HCAIep '24, page 8–14, New York, NY, USA. Association for Computing Machinery.
- Jiaming Ji, Tianyi Qiu, Boyuan Chen, Borong Zhang, Hantao Lou, Kaile Wang, Yawen Duan, Zhonghao He, Lukas Vierling, Donghai Hong, Jiayi Zhou, Zhaowei Zhang, Fanzhi Zeng, Juntao Dai, Xuehai Pan, Kwan Yee Ng, Aidan O’Gara, Hua Xu, Brian Tse, and 7 others. 2025. [AI alignment: A comprehensive survey](#). *Preprint*, arXiv:2310.19852.
- Hamdy Mubarak, Abubakr Mohamed, and Majd Hawasly. 2025. [Arasafe: Benchmarking safety in Arabic LLMs](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 9976–9992, Suzhou, China. Association for Computational Linguistics.
- Wajdi Zaghoulani, Md. Rafiul Biswas, Kholoud Khalil Aldous, and Mabrouka Bessghaier. 2026. [ArabicDialectSafety: A dialect-aware benchmark for Arabic content safety classification](#). *Preprint*, arXiv:2608.01291.
- Zhexin Zhang, Leqi Lei, Lindong Wu, Rui Sun, Yongkang Huang, Chong Long, Xiao Liu, Xuanyu Lei, Jie Tang, and Minlie Huang. 2024. [SafetyBench: Evaluating the safety of large language models](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15537–15553, Bangkok, Thailand. Association for Computational Linguistics.