

IslamicEval 2026: The Second Shared Task of Capturing LLMs Hallucination in Islamic Content

Rahaf Alharbi¹, Abdulelah Alturki¹, Watheq Mansour², Rana Malhas³,
Hamdy Mubarak⁴, Kareem Darwish⁴, Tamer Elsayed³, Walid Magdy¹

¹ School of Informatics, The University of Edinburgh, UK

² The University of Queensland, Australia ³ Qatar University, Qatar ⁴ QCRI, HBKU, Qatar

{R.A.A.Alharbi-3, A.A.A.Alturki}@sms.ed.ac.uk,

uqwmans3@uq.edu.au, rana.malhas@qu.edu.qa,

{hmubarak, kadarwish}@hbku.edu.qa,

telsayed@qu.edu.qa, wmagdy@inf.ed.ac.uk

Abstract

We describe the main findings of the second IslamicEval shared task on hallucination detection in LLM-generated Islam related text (IslamicEval 2026). While the first occurrence of the task in 2025 separately evaluated hallucination detection/correction in quoted Quranic verses and Hadiths alongside a distinct QA retrieval subtask, this year integrates these objectives into four fine-grained subtasks with a much larger provided training set. Given an Islamic question and an LLM response, systems must: 1) detect all cited Quran and Hadith references (including verse numbers and narrators); 2) identify hallucinated citations; 3) correct the hallucinated content; and 4) assess the relevance of every citation to the prompt. Sixteen teams participated in the final phase across the four subtasks, utilizing diverse pipelines that highlight both the challenge of the task and the necessity of mitigation strategies in Islamic AI applications.

1 Introduction

Large Language Models (LLMs) are becoming an important part of Arabic natural language processing applications, with a range of Arabic-focused and multilingual models now being used across tasks such as question answering, retrieval, and content generation (Sengupta et al., 2023; Bari et al., 2024; Fanar Team, 2025). Alongside these advances, hallucination, the generation of text that appears plausible but is factually incorrect, unsupported, or fabricated, remains a key challenge.

This challenge is particularly sensitive in domains where accuracy and authenticity are critical, such as religion. As users increasingly consult LLMs for religious guidance, the risk of generating fabricated or ungrounded Islamic content has become a critical concern (Abdelaal et al., 2026; Elmahjub et al., 2026). Hallucinated citations can lead to misattributed verses, paraphrased text pre-

sented as verbatim citations, or entirely fabricated Hadiths, risking misrepresentation of canonical sources and contributing to the spreading of religious misinformation (Fawzi et al., 2026).

Recent benchmarks such as IslamicMMLU (Abdelaal et al., 2026), IslamicLegalBench (Elmahjub et al., 2026), MAWARITH (Boucekif et al., 2026), and QIAS 2025 (Boucekif et al., 2025) have expanded the evaluation of LLMs on Islamic content, mainly through factual knowledge and specific reasoning tasks. However, they do not address hallucination in generated Islamic content as a direct evaluation target. Instead, this issue usually appears only indirectly through patterns of factual errors.

IslamicEval 2025 (Mubarak et al., 2025) was the first shared task to make hallucinated citations the central evaluation target. IslamicEval 2026 introduces a few improvements. First, it moves from a single label per quotation to a more granular segment-level evaluation, where each citation is fragmented into its quoted text (*ayah* for Quranic verses or *matn* for the reported prophetic saying), its *isnad* (the chain of transmitters, where a narrator is named), and its claimed source (the collection, chapter, or verse attribution supplied by the model), and each segment receives its own labels. Errors that a single label would conceal, such as an accurately quoted text paired with an incorrect attribution, are now separately evaluable, directly exposing common LLM misattribution patterns. Second, it introduces a relevance subtask to verify whether a cited Quran and/or Hadith is relevant to the question’s answer. This subtask was missing last year; we noticed that cited Quran and Hadith can be correctly cited (non-hallucinated) but still irrelevant to the question, which is a different type of relevance-related hallucination. Third, it addresses issues such as variations in attested Hadith wording and the omission of diacritics by providing

a list of potential corrections rather than a single specific one. Fourth, questions in earlier Arabic Islamic question-answering datasets are usually hand-written (Malhas and Elsayed, 2020) or taken from textbooks (Alnefaie et al., 2023; Boucekif et al., 2025). These structured sources drift away from natural user queries (Bhatia et al., 2026). Therefore, we built our question pool from real Islamic questions from two interaction logs, namely Fanar and AnsariChat (Fanar Team, 2025; Kadous et al., 2026), thereby providing a realistic sample of what users ask LLMs in practice. Finally, we provide a substantially larger set of training, development, and test data, made possible by employing LLM-assisted annotations validated against expert-annotated gold sets in place of the exhaustive manual annotation used in IslamicEval 2025, and generating a fully annotated set of approximately 6,000 LLM responses across dataset splits, quadrupling the amount of released data compared to IslamicEval 2025.

2 Task Description

IslamicEval 2026 has four subtasks, each evaluated on a separate blind test set to allow teams to participate in any subtask they want. We restricted the teams to using LLM models that are at most 13B parameters.

- **Subtask 1: Span Detection.** Given an LLM-generated response to a user question, participants determine the spans of the claimed fragments of religious texts, and label each by type: *ayah*, Hadith *matn*, *isnad*, or claimed source. Spans are represented by character indexes into the response. Evaluation is based on character-level macro-averaged F1 over the four segment classes.
- **Subtask 2: Hallucination Identification.** For each segment of the gold citations, participants need to categorize it as correct or incorrect. The quoted text (*ayah* or *matn*) is correct only if it reproduces an attested text verbatim, while the *isnad* and claimed source are judged against the quoted text and are N/A when the quoted text is incorrect. Evaluation is based on accuracy per segment type, macro-averaged across types.
- **Subtask 3: Correction.** Given a gold citation span labeled incorrect, recover the canonical text it distorts, or mark the span as uncorrectable if none exists. Correction covers Quranic ayahs and Hadith *matn* within the six canonical collec-

tions released with the dataset. A span may have more than one possible correction, and matching any of them is credited.

- **Subtask 4: Answer Relevance.** Given a user question in Arabic, an LLM response to that question, and the correctly extracted Qur’anic and Hadith-*matn* citation spans from that response, classify each citation using a binary label: **relevant (1)** or **non-relevant (0)**. A citation span is considered **relevant** if it *directly* or *indirectly answers* the question. Citation spans that are merely topically related but *do not answer* the question, as well as citation spans that are entirely *unrelated*, should be classified as **non-relevant**. A citation span may be authentic and accurately quoted yet still be labeled non-relevant if it does not answer the question. The relevance labels are defined as follows.

– Relevant (1)

- * **with direct answer:** The citation span *explicitly* answers the question, and its meaning aligns with the question’s intent and context, or
- * **with indirect answer:** This applies in two cases: First: The citation span does not explicitly answer the question, but it *implicitly* provides the requested information, and its meaning aligns with the question’s intent and context. Second: The citation span does not directly or implicitly answer the question on its own, but it provides clear textual support for the reasoning in the system’s response such that an answer to the question can be inferred.

- **Non-relevant (0):** The citation span does not provide a *direct* or *indirect* answer to the question. If a citation is only topically related but does not actually answer the question or support the answer reasoning in a meaningful way, it must be marked as non-relevant. Topical overlap alone is insufficient.

We note that *partial support* is permitted; a relevant citation span may provide either a complete answer to the question, or an essential component of the answer. Figure 1 shows an example for all subtasks.

3 Data and Resources

The dataset is built from two sets of real user interactions rather than a curated question bank.

Question	ما حكم حمل الجنازة للدفن؟				
Translation: What is the ruling on carrying the funeral bier for burial?					
Response	حمل الجنازة والدفن هو من السنن المؤكدة في الإسلام، وقد حث عليه النبي محمد صلى الله عليه وسلم. قال الله تعالى في كتابه الكريم: {وَلَا تَسْرُوا الْقُفُلَ يَنْكُرُ} [البقرة: 237] ^{ق1}. وقد ورد في الحديث الشريف عن أبي هريرة رضي الله عنه أن النبي صلى الله عليه وسلم قال: "من شهد الجنازة حتى يصلى عليها فله قيراط، ومن شهدا حتى يُدفن فله قيراطان، قيل: وما القيراطان؟ قال: أصغرهما مثل أحد" [رواه البخاري ومسلم] ^{ق2}. وفي حديث آخر عن أبي سعيد الخدري رضي الله عنه، أن النبي صلى الله عليه وسلم قال: "إذا وُضعت الجنازة واحتملها الناس بأكتافهم، فإن كانت صالحة قالت: قدموني قدموني، وإن كانت غير صالحة قالت: يا ويلها، أين تذهبون بها؟ يسمع صوتها كل شيء إلا الإنسان، ولو سمعه صغى" [رواه البخاري] ^{ق3} ...				
...					
Citation	Segment	Subtask 1 Span Detection	Subtask 2 Hallucination Identification	Subtask 3 Correction	Subtask 4 Answer Relevance
C1 <i>Ayah citation</i>	Ayah	127–161	Correct	—	0
	Claimed source	164–175	Correct	—	—
C2 <i>Hadith citation</i>	Isnad	203–261	Correct	—	—
	Hadith matn	263–378	Correct	—	1
	Claimed source	381–399	Correct	—	—
C3 <i>Hadith citation</i>	Isnad	416–481	N/A	—	—
	Hadith matn	483–657	Incorrect	إذا وُضعت الجنازة واحتملها الرجال على أكتافهم، فإن كانت صالحة قالت: قدموني ...	1
	Claimed source	660–672	N/A	one of four accepted attested wordings	—

Figure 1: An annotated training-set example (Response R003070, Question Q00405) across the four subtasks. The second Hadith citation is fully correct; the third alters the attested wording, so its *isnad* and claimed source are N/A, and a correction is supplied; the correctly quoted Ayah is judged irrelevant to the question, showing that a correct citation is not necessarily relevant evidence.

Questions from the chat logs of Fanar (Fanar Team, 2025), an Arabic-centric AI platform, form the training and development sets, and questions from AnsariChat (Kadous et al., 2026), an Islamic retrieval-augmented chatbot, form the blind test set. Drawing the test set from a different source ensures no test question appears in the train/dev sets and that systems are evaluated on a question/response distribution not seen during development.

Prompts from both sources were filtered for Islamic relevance, keeping only standalone questions answerable from the Quran and Hadith, and for appropriateness, removing prompts whose framing demeans a group, solicits harm, contains obscene language or personal data. From each filtered pool, 500 questions were sampled, and responses were collected from 12 models: four Arabic-centric, six multilingual, and two frontier models, GPT-5.4 and Gemini-3.1-Pro. All models received the same Arabic prompt that explicitly asks them to support their answers with Quranic ayahs and Hadiths and to name the source of each citation. The generated answers were then cleaned of degenerate repetition, truncation, and heavy code-switching, which would otherwise degrade citation quality through dupli-

cated or truncated spans. After cleaning, 5,674 LLM responses to Fanar questions and 5,629 LLM responses to AnsariChat questions remain. Following citation extraction and release validation, we released 5,190 annotated responses to Fanar questions as the training and development sets, and a sample of 3,620 responses to AnsariChat questions was annotated as the blind test set, a size set by the annotation budget and by the scale of the released development split. Within the test set, 620 responses are reserved for Subtask 1, and the citations of the remaining 3,000 are partitioned among Subtasks 2, 3, and 4 so that the annotation of one subtask cannot leak the labels of another. Annotation runs as an LLM pipeline over the four subtasks in sequence, where each stage annotates the output of the previous one. For each subtask, a sample of the data was annotated by three Islamic-studies experts and used as a gold set to select the frontier model and prompt configuration most aligned with expert annotations, which are subsequently used to annotate the full set. The size of each gold set and its agreement with the automatic labels are reported per subtask in the annotation setup subsections. Table 1 summarises the dataset composition

Subtask & Element	Train	Dev	Test
Subtask 1: Span Detection			
Responses	4,706	484	620
<i>citation-free</i>	583	72	61
Quran citations	8,195	698	1,299
Hadith citations	7,048	588	1,102
Isnad spans	2,634	204	448
Claimed-source spans	14,902	1,238	2,355
Subtask 2: Hallucination Identification (<i>correct / incorrect</i>)			
Ayah	4,492 / 3,703	222 / 476	3,539
Matn	2,842 / 4,206	121 / 467	3,076
Isnad	808 / 130	16 / 14	465
Claimed source	6,781 / 1,126	283 / 146	3,295
Subtask 3: Correction (<i>correctable / not correctable</i>)			
Ayah	1,677 / 2,026	219 / 257	944
Matn	1,029 / 3,177	108 / 359	1,054
Subtask 4: Answer Relevance (<i>relevant / not relevant</i>)			
Ayah	2,953 / 576	76 / 82	432
Matn	2,301 / 401	75 / 49	333

Table 1: Dataset composition across splits and subtasks. For Subtasks 2–4, Train and Dev report slashed label counts formatted as *correct/incorrect* (Subtask 2), *correctable / not correctable* (Subtask 3), and *relevant / not relevant* (Subtask 4). Test reports the total number of unlabelled instances for evaluation. Subtask 2 excludes *isnad* and claimed source spans marked N/A due to incorrect text.

across the training, development, and test splits for all four subtasks.

3.1 Annotation Setup and Guidelines

3.1.1 Subtask 1

To annotate spans, an LLM detects and segments citations, followed by a deterministic matching step that locates each span in the response. Gemini 3.5 Flash (Google DeepMind, 2026b) was selected for this stage because it achieved higher precision and fewer false-positive Hadith detections than GPT-5.5 on a manually annotated validation set. It is prompted to find every intended *ayah* or Hadith in the response, and to return its type, full span, claimed source, and, for Hadith, the *matn* and *isnad* sub-spans. Because the model may return the text slightly altered (e.g., missing diacritics or quotation marks), we locate each fragment in the original response to record its exact start and end positions. A response is discarded whenever any of its fragments cannot be found, guaranteeing that every released span is an exact substring. For the blind test set, responses were annotated independently by Gemini 3.5 Flash and GPT-5.5 (OpenAI, 2026), retaining only spans where the two models agreed on the citation type and boundaries. The automatic annotation was validated against a gold set of 199 responses containing 692 citations (343 Quran and 349 Hadith), with 12 citation-free

responses retained to test whether the model correctly abstains when no citation is present. Gold spans were copied verbatim following the boundary guidelines in Appendix A. Against this gold set, the pipeline reached a citation-detection F1 of 98.6% and abstained correctly on all citation-free responses. It segmented every fragment reliably, with span F1 of 99.9% on the *ayah*, 98.6% on the *matn*, 97.4% on the *isnad*, and 99.5% and 98.1% on the Quranic and Hadith claimed source, respectively. The model rarely misses citations, and the residual errors are at the boundary level, such as quotation marks or emphasis markers swept into the span.

3.1.2 Subtask 2

The labels are produced by Gemini 3.1 Pro (Google DeepMind, 2026a), the model selected for this subtask because it achieved the strongest overall agreement with the expert majority labels, which receives each citation through a dedicated Quran or Hadith prompt with its segments as separate fields, so that the quoted text is labelled first and the *isnad* and claimed source are then labelled against it. The labelling rule is strict: the quoted text is correct only if every word matches an attested wording, meaning any documented transmission for Hadith and the canonical verse for Quran, with nothing omitted from the middle; a contiguous partial quotation is accepted, and Quranic diacritics are judged only where the citation supplies them. The authenticity grade of a Hadith does not affect the label, as a weak or forged report quoted exactly is still an existing text, whereas an altered or invented one is not. The decision rules are provided in Appendix B.

The automatic labels were validated against an expert gold set of 350 randomly sampled citations (250 Hadith and 100 Quran), independently annotated by three Islamic-studies experts. Each expert assessed every segment, assigning *correct* or *incorrect* to the quoted text, the *isnad*, and the claimed source. The majority vote across the three experts serves as the gold label. The automatic labels substantially agree with this gold set on Hadith and almost perfectly on Quran, with every segment scored only where the citation supplies it. On Hadith, agreement is substantial on the *matn* ($\kappa = 0.70$, $n = 248$) and the claimed source (0.70 , $n = 241$), reflecting the difficulty of judging wording against multiple attested transmissions, and rises to almost perfect on the *isnad* (0.86 , $n = 108$). A manual review of the Hadith disagreements shows that



Figure 2: One Hadith is attested in four wordings that are semantically equivalent but differ in wording, with the differences highlighted in red. All four are accepted corrections for a span that distorts this Hadith.

they stem primarily from differing reference texts—where the judge and human experts matched the *matn* to different attested transmissions—or from minor wording variations, rather than cases where a non-existent citation was marked correct. On Quran, where the text is fixed and a quotation either matches the verse or does not, agreement reaches 0.88 on the *ayah* ($n=99$) and 0.97 on the claimed source ($n=97$).

3.1.3 Subtask 3

Corrections are produced by GPT-5.4, selected for this stage through validation with expert judgments. The model was prompted to permit a correction only when a contiguous sequence of words from the span appears in a real canonical text after normalization, and to decline spans that share only topic or scattered words with the corpora. Every proposed correction is subsequently grounded against the released corpora by snapping it to the closest text, which fixes the scope of correction to the provided corpora and guards against potential hallucinated corrections. Grounded Hadiths are additionally expanded to retrieve attested variants through word-based Jaccard similarity, so that a participant matching any documented wording is credited. Figure 2 shows one Hadith with four attested variants, all of which are acceptable corrections for the hallucinated span. The corrector was validated on a gold set of 100 incorrect spans (50 Quran and 50 Hadith). Each span was corrected by GPT-5.4 and Gemini 3.1 Pro, and three Islamic-studies experts judged every model’s corrections, either accepting or rejecting it and supplying their own, including marking the span uncorrectable. Experts agree reliably on whether a span is recoverable (Fleiss $\kappa = 0.688$), but often propose different valid corrections, so the gold for a span is a set of accepted corrections rather than a single string; a candidate is majority-backed when at least two annotators support it. GPT-5.4 was chosen as the

corrector model as it agrees with a majority-backed candidate on 90% of spans, reproduces every correction on which the experts unanimously agree, and its residual errors lean toward the safer failure of declining rather than forcing a correction.

3.1.4 Subtask 4

To establish a reliable and scalable answer relevance annotation pipeline for thousands of question-response-citation triplets, we first conducted a pilot study on a sample of 185 triplets with two Islamic Studies domain experts following structured guidelines with illustrative examples. These guidelines (excluding examples) were then adapted into a zero-shot prompt for three frontier LLMs acting as judges: Gemini 3.5 Flash, GPT-5.4, and Claude Sonnet. To elevate label reliability beyond the moderate agreement observed between human experts (Cohen’s $\kappa = 0.497$), we adopted a unanimous agreement protocol across the three models. On the pilot sample, requiring 3-way unanimous LLM agreement raised the overall inter-annotator reliability between experts and models to Fleiss’ $\kappa = 0.587$, while agreement among the three LLMs reached $\kappa = 0.617$. Deploying these judges across the Subtask 4 dataset yielded substantial inter-annotator agreement, with Fleiss’ κ scores of 0.699 (train), 0.734 (dev), and 0.659 (test). Accordingly, for the published official Subtask 4 dataset, we strictly adopted this full agreement strategy to guarantee maximal label quality and faithfulness.

4 Evaluation Measures

Subtask 1. Character-level macro-averaged F1 is calculated over the four segment classes (*ayah*, *matn*, *isnad*, and claimed source), following IslamicEval 2025 but extended to the segment level: each character of a response is labelled with the class of the span covering it, and F1 is computed per class and macro-averaged. Predictions are matched to the gold by response identifier, and character-level F1 is computed independently for each segment class. Any span predicted where the gold has none counts against precision, so over-detection in citation-free responses is penalised directly.

Subtask 2. Predictions are matched to the gold by response, annotation, and segment identifiers, and scored with accuracy per segment type (*ayah*, *matn*, *isnad*, claimed source). Segments whose gold label is N/A are excluded from scoring, and a

missing prediction for a scored segment counts as incorrect. The official score is the macro-average across all four categories, preventing frequent classes from dominating rare ones.

Subtask 3. Subtask 3 is scored for overall accuracy, and separately for Quran and Hadith. A prediction is counted correct if it exactly matches any of the acceptable corrections for the span, or marks it uncorrectable when the span is baseless. Declining and correcting are scored on the same scale, so over-correcting fabricated spans costs as much as failing to recover real ones. Quran matching is diacritic sensitive apart from default diacritics, while Hadith is diacritic insensitive. Both system answers and gold corrections are normalized before comparison to avoid penalizing superficial orthographic differences, and all gold corrections are grounded in the shared corpora.

Subtask 4. Macro-averaged F1-score is computed using a two-step process to ensure a fair assessment of citation relevance. For each question-response pair, an initial F1-score is calculated over the system relevance labels; then, macro-average is calculated across all questions to compute the final score.

5 Evaluation Setup

We used CodaBench¹ platform to host competitions for all subtasks. To enable local evaluation, we released evaluation scripts, training, and development sets on our GitHub codebase². The maximum allowed number of submissions per team was 100 during the development phase and 30 during the test phase. However, participants were aware that only the last three valid submissions would be considered for the final evaluation. While we made the leaderboard available during the development phase, we kept it hidden during the test phase. Results of the last three runs were shared with participants by the end of the test phase, and they were requested to describe them in their system description papers. The test phase received relatively strong engagement, with 9 teams in Subtask 1 (16 submissions), 8 teams in Subtask 2 (20 submissions), 4 teams in Subtask 3 (10 submissions), and 7 teams in Subtask 4 (18 submissions).

¹<https://www.codabench.org>

²github.com/Watheq9/IslamicEval2026

6 Participating Teams and Methods

Table 2 lists the participating teams, their affiliations, and their submitted subtasks; Table 3 presents the official test results across all four subtasks, while Table 6 in Appendix E summarises the methodologies adopted by participating teams.

Subtask 1. All the participating teams treated span detection as a sequence-labelling problem, using fine-tuned Arabic encoders. The majority of the teams chose AraBERT-based models (HIAST, Yaqeen, Bayan_CU, and Qahera), while others used variants of CAMELBER (Northwestern 9 and Taqwatech) or AraModernBERT (SanadAI). The most common approach was token classification using standard BIO tagging, although Yaqeen performed character-level labelling. To handle long responses that exceed standard transformer input limits, multiple systems employed overlapping sliding windows (Northwestern 9, Taqwatech, Yaqeen, Bayan_CU, and Qahera). Several teams also applied heuristic post-processing to merge overlapping predictions, reconstruct valid spans, and remove implausible outputs.

Subtask 2. Nearly all participating teams grounded hallucination identification in the canonical corpora, labelling each segment by retrieving and comparing. These retrieval-based systems decoupled text matching from source verification; ayah and matn segments were first matched with canonical sources, after which claimed sources and *isnad* components were validated against the matched texts. For text retrieval, systems combined text normalisation with exact or fuzzy string matching (EIDIANLP, Qahera), sparse TF-IDF search (Taqwatech), or character *n*-gram retrieval with reranking (Namaa Community). To verify attribution, teams applied explicit heuristic rules over collections and narrators (Taqwatech, Namaa Community, Qahera), entity matching against retrieved references (SanadAI), or trained segment-level classifiers (Yaqeen). Tathabbut built an evidence-fusion pipeline that combines TF-IDF features with SVM and LightGBM classifiers based on text overlap, narrator chains, and citation consistency. Unlike corpus-grounded systems, HIAST framed verification as sequence classification, evaluating each segment directly from its surrounding context with a fine-tuned AraBERT model.

Subtask 3. All four teams treated correction as retrieval over the released corpora rather than gener-

ation, returning verbatim corpus strings so that the required wording and diacritics are preserved. Candidates were retrieved through sparse lexical matching, using character and word n-gram TF-IDF (Yaqeen, Tathabbut, Taqwatech) or fuzzy partial matching (Qahera), and then reranked. Three teams learned the reranking with gradient-boosted trees over lexical alignment features such as coverage, sequence agreement, and score margins (Yaqeen, Tathabbut, Taqwatech), while Qahera combined a bag-of-words score with contiguous-substring coverage. Tathabbut additionally drew on the question, the surrounding response, and the claimed source as supporting evidence, and on the corrections of similar training spans. Every team calibrated separate acceptance thresholds for Quran and Hadith, and three of them made the decision to abstain with a component separate from the one that chose the candidate (Yaqeen, Tathabbut, Taqwatech). Only one team used an LLM, and only for Hadith, where a 4B-parameter model verified retrieved candidates (Qahera).

Subtask 4. Many teams employed small-scale LLMs such as ALLaM-7B, Qwen3-4B, or Gemma-3-4B (AhmedNLP, HaDe, May-X, Qahera). To enhance the models’ ability to assess semantic relevance beyond mere topical overlap, participants frequently incorporated Chain-of-Thought (CoT) reasoning (HaDe, May-X, Qahera) or Parameter-Efficient Fine-Tuning (PEFT) via LoRA (AhmedNLP, May-X). Notably, rather than relying on manual prompt crafting, Qahera adopted an iterative propose-evaluate-diagnose search harness to systematically optimize their prompts. Retrieval-augmented architectures spanning both classification and generative setups were also prominent (Bayan, HIAST), demonstrating the value of grounding decisions in external, authenticated corpora. In addition, addressing class imbalance, calibrating decision thresholds, and careful data preprocessing were recurring themes across several systems (AhmedNLP, May-X, Taqwatech).

7 Results and Discussion

7.1 Subtask 1: Span Detection

Bayan_CU secured first place in Subtask 1 with a character-level Macro-F1 score of 97.36%. The team fine-tuned AraBERT as a BIO token classifier and used overlapping windows, span merging, and heuristic filtering to improve the final boundaries. SanadAI placed second with 97.35%, using

AraModernBERT, class-weighted training, and additional Quranic and Hadith examples, including manually selected hard negatives from Shamela³. HIAST ranked third with 97.25%, using an AraBERT BIO classifier and overlapping windows to process long responses. Most other teams followed a similar encoder-based sequence-labelling approach, using either AraBERT, CAMELBER, or AraModernBERT. Only 0.11 percentage points separated the three leading systems, suggesting that post-processing, data preparation, and long-response handling were as important as the choice of encoder. Although Northwestern 9 reported an F1 score of 96.30% on the development set, they received 0% on the official test set due to a submission error on their side.

7.2 Subtask 2: Hallucination Identification

Tathabbut secured first place in Subtask 2 with an accuracy of 94.52%. Its system combined TF-IDF features, SVM and LightGBM classifiers, and rule-based checks to use evidence from canonical texts, narrator information, and citation context. Yaqeen placed second with 91.07%, retrieving candidate texts and using separate classifiers to verify *ayah*, *matn*, claimed-source, and *isnad* segments. Qahera ranked third with 90.58%, using exact and fuzzy matching for quoted texts and explicit checks for their stated sources and narrator information. Taqwatech followed closely with 90.50%, combining Arabic text normalisation, lexical retrieval, character TF-IDF, source checks, and narrator matching. No team used a generative model at inference for Subtask 2; instead, the strongest systems combined retrieval from canonical corpora with classification or rule-based verification.

7.3 Subtask 3: Hallucination Correction

Taqwatech and Tathabbut tied for first on Subtask 3, both reaching 86.29% overall accuracy. Taqwatech used a TF-IDF retriever and ranked candidates with a gradient-boosting model over 30 lexical features, and deciding whether to correct with a separate classifier improved its performance. In addition to TF-IDF retrieval and reranking, Tathabbut used the surrounding context, the original question, and the claimed source as soft signals for ayah correction, together with a memory of training hallucinations to support Hadith correction. Yaqeen placed second, using character n-gram retrieval, a type-aware

³<https://shamela.ws/>

Team Name	Subtasks:	1	2	3	4	Affiliation
AhmedNLP (Almatrafi, 2026)					✓	Islamic University of Madinah, Saudi Arabia
anhnamxtanh*		✓				—
Bayan					✓	Benha University, Egypt
Bayan_CU (Ghonem, 2026)		✓				Cairo University, Egypt
EIDIANLP			✓			Euro-Mediterranean University of Fez, Morocco
HaDe (Alowaidi, 2026)					✓	University of Leeds, United Kingdom and King Abdulaziz University, Saudi Arabia
HIAS (Sonbol, 2026)		✓	✓		✓	Higher Institute for Applied Sciences and Technology, Syria
May-X					✓	Kafr Elsheikh University, Egypt and Damietta University, Egypt
Namaa Community (Elden et al., 2026)			✓			Trouve Labs, NAMA Community, Arab International University, Syria and Tuwaiq Academy, Saudi Arabia
Northwestern 9 (Shlkamy and Zaghouani, 2026)		✓				Northwestern University in Qatar, Qatar
Qahera (Ajlan and Mohamed, 2026)		✓	✓	✓	✓	American University in Cairo, Egypt
SanadAI (Korched, 2026)		✓	✓			—
Tabayun*		✓				—
Taqwatech (Zaky et al., 2026)		✓	✓	✓	✓	Nile University, Egypt and Helwan University, Egypt
Tathabbut (Hegazy, 2026)			✓	✓		Nile University, Egypt
Yaqeen (Alsuhaibani and Alkaoud, 2026)		✓	✓	✓		King Saud University, Saudi Arabia

Table 2: Participating teams and their corresponding subtask submissions. Teams marked with (*) did not submit a system description paper.

Team Name	F1%	Team Name	Acc%	Ayah	Matn	Isnad	CS	Team Name	Acc%	Ayah	Matn	Team Name	F1%
Bayan_CU	97.36	Tathabbut	94.52	98.28	90.64	93.12	96.05	Taqwatech	86.29	84.64	87.76	HIAS	77.21
SanadAI	97.35	Yaqeen	91.07	92.00	84.43	92.26	95.60	Tathabbut	86.29	84.53	87.86	AhmedNLP	76.79
HIAS	97.25	Qahera	90.58	94.49	88.91	85.16	93.75	Yaqeen	78.03	70.44	84.82	HaDe	74.84
Qahera	97.16	Taqwatech	90.50	97.57	91.87	77.63	94.93	Qahera	77.28	71.50	82.45	Bayan	71.46
Taqwatech	96.83	EIDIANLP	89.13	97.17	87.45	81.29	90.59	Baseline	65.77	54.66	75.71	Taqwatech	69.42
Yaqeen	96.78	SanadAI	86.66	96.89	88.56	77.42	83.76	(c) Subtask 3				Qahera	67.26
anhnamxtanh*	96.10	HIAS	86.62	87.74	87.55	84.30	86.89					May-X	85.56
Tabayun*	91.76	Namaa Comm.	66.85	81.77	62.19	89.46	33.96					Baseline	71.24
Northwestern 9	0.00	Baseline	65.62	51.85	37.74	89.46	83.43						

(a) Subtask 1

(b) Subtask 2

(d) Subtask 4

Table 3: Results across the four subtasks, reporting each team’s best run; teams are ranked (R) per subtask. Teams marked with (*) did not submit a system description paper. **Subtask 1:** character-level macro-averaged F1 over the four citation classes (Ayah, Matn, Isnad, CS: claimed source). We exclude the non-citation (neither) class from scoring; therefore, we report no majority baseline, since a trivial all-neither prediction scores 0 and predicting the majority citation class (all-Ayah) reaches only 5% F1. **Subtask 2:** accuracy macro-averaged over the four segment types, with per-segment accuracy; the majority baseline (predicting all segments correct) is **65.62%**. **Subtask 3:** correction accuracy overall and per corpus; the majority baseline (predicting all citations uncorrectable) is **65.77%**. **Subtask 4:** macro-averaged F1; the majority baseline is **71.24%**. Taqwatech and Tathabbut tie for first on Subtask 3. May-X is unranked in Subtask 4; the submission was excluded for exceeding the 13B parameter limit.

classifier for the abstain decision, and reranking applied only to Quran. Qahera placed third, using fuzzy partial matching for retrieval and a Gemma-3-4B model in a verify, select, then re-verify cascade over its Hadith candidates, while its Quran corrections were decided by a single overlap-and-coverage score. No team fine-tuned a model or generated corrections, which leaves both as open directions for future systems on this task.

7.4 Subtask 4: Answer Relevance

The HIAS team officially secured first place in Subtask 4 with a Macro-F1 score of 77.21% (Table 3). They adopted a retrieval-augmented, evidence-grounded framework combining hybrid lexical-semantic retrieval with a fine-tuned AraBERTv02 classifier. Their retrieval pipeline coupled BM25 and E5 dense embeddings indexed

via FAISS, fusing the candidate lists using Reciprocal Rank Fusion (RRF).

The second and third ranked teams both leveraged the ALLaM-7B model. The second place, AhmedNLP, achieved a score of 76.79% (Table 3) by fine-tuning ALLaM-7B using LoRA and a response-balanced RankNet loss, deterministically scoring relevance via first-token logit differences between “yes” (*na’am*) and “no” (*la*). Meanwhile, HaDe attained third place with a score of 74.84% using a few-shot prompting setup with ALLaM-7B; notably, their highest test score was achieved without Chain-of-Thought (CoT), slightly outperforming their CoT-augmented baseline.

Although team May-X achieved the highest overall Macro-F1 score of 85.56% (Table 3), their submission was excluded from the official rankings for violating the task constraint prohibiting using

LLMs exceeding 13B parameters. Specifically, they utilized gpt-oss-120B to synthesize rationales for offline reasoning (CoT) distillation into a compact Qwen3-4B student model. In their core architecture, May-X adopted a multi-stage parameter-efficient fine-tuning setup for Qwen3-4B using two specialized LoRA adapters: a domain-specific adapter and a task-specific adapter for answer relevance classification.

8 Conclusion

IslamicEval 2026 builds on IslamicEval 2025 by expanding the evaluation into four subtasks and adopting a more granular approach to evaluating citations, their attribution, their recoverability, and their relevance to the user’s query. The benchmark highlights attribution errors involving *isnads* and claimed sources, showing that correctly quoting an *ayah* or Hadith does not necessarily mean that it has been attributed to the correct source or *isnad*, or is relevant to the question. By separating these different types of failure, we provide a clearer assessment of LLM reliability for developing systems that produce accurate religious answers.

9 Limitations

In IslamicEval 2026, we introduced an LLM-assisted data annotation pipeline to overcome the manual, exhaustive annotation process used in IslamicEval 2025 (Mubarak et al., 2025). As a result, in this edition the released annotated dataset grew substantially. However, we note that although the annotation models in each subtask were validated against manually annotated gold samples, some annotation errors may remain. Moreover, the questions are drawn from two sources of interaction, and 12 models generated the responses explicitly instructed to provide religious evidence. Therefore, the resulting data represent model behaviour when citations are requested and may not generalise to unconstrained responses, other question sources, languages, or model families. In addition, Hadith verification accepts any attested wording but treats a meaning-preserving paraphrase as incorrect; the results therefore measure strict textual fidelity and do not measure acceptability under transmission by meaning (*riwāya bi-l-ma’nā*). While Subtask 2 validated all Hadith variants across different collections, Subtask 3 limited correction to the six canonical collections; we note that some valid, correctable Hadiths attested elsewhere cannot be credited.

10 Ethical Considerations

Given the religious sensitivity of the content of this task, we took care to ensure accuracy and respect. Our annotation, while carried out by Frontier LLMs, followed careful procedures to ensure the highest possible quality, verified by experts as described in the data preparation setup. The dataset is released strictly for research purposes, with the intention of improving the reliability and safety of LLMs in handling religious material. We explicitly caution against any misuse of this resource in contexts that could distort, misrepresent, or disrespect Islamic teachings.

References

- Ali Abdelaal, Mohammed Nader Al Haffar, Mahmoud Fawzi, and Walid Magdy. 2026. [IslamicMMLU: A benchmark for evaluating LLMs on islamic knowledge](#). *arXiv preprint arXiv:2603.23750*.
- Hazem Ajlan and Mohamed Mohamed. 2026. Qahera at islamicval 2026: Corpus-grounded retrieval and iterative harness search with a 4b-parameter llm for islamic citation verification. In *Proceedings of the Fourth Arabic Natural Language Processing Conference*.
- Ahmed Almatrafi. 2026. Ahmednlp at islamicval 2026: Learning answer relevance from within-response citation preferences. In *Proceedings of the Fourth Arabic Natural Language Processing Conference*.
- Sarah Alnefaie, Eric Atwell, and M Ammar Alsalka. 2023. Haqa and quqa: Constructing two arabic question-answering corpora for the quran and hadith. In *Proceedings of the Conference Recent Advances in Natural Language Processing-Large Language Models for Natural Language Processings*, pages 90–97. INCOMA Ltd., Shoumen, BULGARIA.
- Yazeed Alnumay, Alexandre Barbet, Anna Bialas, William Darling, Shaan Desai, Joan Devassy, Kyle Duffy, Stephanie Howe, Olivia Lasche, Justin Lee, et al. 2025. Command r7b arabic: a small, enterprise-focused, multilingual, and culturally aware arabic llm. In *Proceedings of the Sixth Workshop on African Natural Language Processing (AfricaNLP 2025)*, pages 126–135.
- Sanaa Alowaidi. 2026. Hade at islamicval 2026: Evaluating prompted judgment and fine-tuned encoders for islamic citation relevance. In *Proceedings of the Fourth Arabic Natural Language Processing Conference*.
- Shahad Alsuhaibani and Mohamed Alkaoud. 2026. Yaqeen at islamicval 2026: From contextual detection to retrieval-based verification and canonical correction of llm-generated arabic islamic citations. In

- Proceedings of the Fourth Arabic Natural Language Processing Conference.*
- M Saiful Bari, Yazeed Alnumay, Norah A Alzahrani, Nouf M Alotaibi, Hisham A Alyahya, Sultan Al-Rashed, Faisal A Mirza, Shaykhah Z Alsubaie, Hassan A Alahmed, Ghadah Alabduljabbar, et al. 2024. ALLam: Large language models for Arabic and English. *arXiv preprint arXiv:2407.15390*.
- Gagan Bhatia, Hamdy Mubarak, Mustafa Jarrar, George Mikros, Fadi Zaraket, Mahmoud Alhirthani, Mutaz al Khatib, Logan Cochrane, Kareem Mohamed Darwish, Rashid Yahiaoui, and Firoj Alam. 2026. [From RAG to agentic RAG for faithful islamic question answering](#). In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 26469–26488, San Diego, California, United States. Association for Computational Linguistics.
- Abdessalam Boucekif, Shahd Gaben, Samer Rashwani, Somaya Eltanbouly, Mutaz Al-Khatib, Heba Sbahi, Mohammed Ghaly, and Emad Mohamed. 2026. [MAWARITH: A dataset and benchmark for legal inheritance reasoning with LLMs](#). *arXiv preprint arXiv:2603.07539*.
- Abdessalam Boucekif, Samer Rashwani, Emad Soliman Ali Mohamed, Mutaz Alkhatib, Heba Sbahi, Shahd Gaben, Wajdi Zaghouani, Aiman Erbad, and Mohammed Ghaly. 2025. [Qias 2025: Overview of the shared task on islamic inheritance reasoning and knowledge assessment](#). In *Proceedings of The Third Arabic Natural Language Processing Conference: Shared Tasks*, page 851–860. Association for Computational Linguistics.
- Fatimah Mohamed Emad Elden, Israa Elhosiny, Khlood Al Jallad, and OMER NACAR. 2026. Namaa community at islamiceval 2026: Retrieval-grounded verification of qur’anic and hadith citations for hallucination identification. In *Proceedings of the Fourth Arabic Natural Language Processing Conference*.
- Ezieddin Elmahjub, Junaid Qadir, Abdullah Mushtaq, Rafay Naeem, Ibrahim Ghaznavi, and Waleed Iqbal. 2026. [IslamicLegalBench: Evaluating LLMs knowledge and reasoning of islamic law across 1,200 years of islamic pluralist legal traditions](#). *arXiv preprint arXiv:2602.21226*.
- Fanar Team. 2025. Fanar: An arabic-centric multimodal generative ai platform. *arXiv preprint arXiv:2501.13944*.
- Mahmoud Fawzi, Björn Ross, and Walid Magdy. 2026. Fabricating holiness: Characterizing religious misinformation circulators on arabic social media. *Proceedings of the International AAAI Conference on Web and Social Media*, 20.
- Samah Ghonem. 2026. Bayan at islamiceval 2026: Fine-grained hallucination span detection in arabic quranic and hadith texts. In *Proceedings of the Fourth Arabic Natural Language Processing Conference*.
- Google DeepMind. 2026a. Gemini 3.1 pro model card. <https://deepmind.google/technologies/gemini/>.
- Google DeepMind. 2026b. Gemini 3.5 flash model card. <https://deepmind.google/models/model-cards/gemini-3-5-flash/>.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.
- Salma Khaled Hegazy. 2026. Tathabbut at islamiceval 2026: Canonical-provenance evidence fusion for fine-grained citation hallucination verification. In *Proceedings of the Fourth Arabic Natural Language Processing Conference*.
- M Waleed Kadous, Amr Elsayed, Abdullah Al Nahas, and Ashraf Haress. 2026. [Ansari: A retrieval-grounded islamic ai assistant – architecture, deployment, and lessons from 140,000 conversations](#). *Preprint*, arXiv:2608.20390.
- Abir Korched. 2026. Sanadai at islamiceval 2026: Detecting hallucinations in arabic islamic citations through span detection and hybrid verification. In *Proceedings of the Fourth Arabic Natural Language Processing Conference*.
- Juhao Liang, Zhenyang Cai, Jianqing Zhu, Huang Huang, Kewei Zong, Bang An, Mosen Alharthi, Juncui He, Lian Zhang, Haizhou Li, Benyou Wang, and Jinchao Xu. 2024. [Alignment at pre-training! towards native alignment for arabic llms](#). *Preprint*, arXiv:2412.03253.
- Rana Malhas and Tamer Elsayed. 2020. AyaTEC: Building a Reusable Verse-based Test Collection for Arabic Question Answering on the Holy Qur’an. *ACM Transactions on Asian and Low-Resource Language Information Processing*, 19(6):1–21.
- Mistral AI. 2024. [Les ministraux: Mistral on the edge](#). Accessed: Aug. 12, 2026.
- Hamdy Mubarak, Rana Malhas, Watheq Mansour, Abubakr Mohamed, Mahmoud Fawzi, Majd Hawasly, Tamer Elsayed, Kareem Darwish, and Walid Magdy. 2025. [Islamiceval 2025: The first shared task of capturing llms hallucination in islamic content](#). *Proceedings of The Third Arabic Natural Language Processing Conference: Shared Tasks*.
- OpenAI. 2026. OpenAI GPT-5.5 system card. *arXiv preprint*.
- Neha Sengupta, Sunil Kumar Sahu, Bokang Jia, Satheesh Katipomu, Haonan Li, Fajri Koto, William Marshall, Gurpreet Gosal, Cynthia Liu, Zhiming Chen, et al. 2023. Jais and jais-chat: Arabic-centric foundation and instruction-tuned open generative large language models. *arXiv preprint arXiv:2308.16149*.

Eman Sedqy Shlkamy and Wajdi Zaghouani. 2026. Northwestern 9 at islamiceval 2026: Fine-grained span detection in arabic islamic content using camelbert-mix. In *Proceedings of the Fourth Arabic Natural Language Processing Conference*.

Riad Sonbol. 2026. Hiast at islamiceval 2026: Context-aware hallucination detection and evidence-guided verification for arabic islamic question answering. In *Proceedings of the Fourth Arabic Natural Language Processing Conference*.

Gemma Team, Aishwarya Kamath, Johan Ferret, Shreya Pathak, Nino Vieillard, Ramona Merhej, Sarah Perrin, Tatiana Matejovicova, Alexandre Ramé, Morgane Rivière, et al. 2025. Gemma 3 technical report. *arXiv preprint arXiv:2503.19786*.

Phi3 Team. 2024. *Phi-3 technical report: A highly capable language model locally on your phone*. *Preprint*, arXiv:2404.14219.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. 2025. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*.

Sama Ayman Zaky, Basmala Alsabagh, Kariim Reda, and Abdelrahman Mohsen. 2026. Taqwatech at islamiceval 2026: Contextual span detection and corpus-grounded verification of islamic citations. In *Proceedings of the Fourth Arabic Natural Language Processing Conference*.

A Subtask 1: Span Annotation Guidelines

In Subtask 1, Gold spans are copied verbatim from the response; wrappers, framing phrases, and punctuation around a segment are excluded from it.

1. **Citation.** Any span framed as Quran or Hadith is annotated, whether or not its content is accurate; framing is signalled by an attribution verb, a quotation wrapper, or a source bracket.
2. **Ayah.** The verse text only, excluding introductory framing (e.g., “قال الله تعالى”) and wrappers; the source reference is annotated separately as the claimed source.
3. **Isnad.** The chain of named narrators, from the first narrator term through the transmission verb introducing the saying. An *isnad* is present only when a person is named; a citation that attributes words directly to the Prophet, uses passive narration, or names only a collection has no *isnad*.
4. **Matn.** The reported saying only. A transmission verb occurring within the reported content remains within the *matn*, and in dialogue Hadiths, the entire reported exchange constitutes the *matn*.

5. **Claimed source.** The fullest contiguous attribution phrase, copied verbatim, including any collection name, authenticity grading, and verse or Hadith number.

6. **Special cases.** An Ayah quoted inside a Hadith stays within the *matn* and is not annotated as a separate Quranic citation. Every occurrence is annotated, including repeated quotations of the same text. When the same claimed source is stated more than once for the same citation, the last, fullest occurrence is recorded.

B Subtask 2: Labelling Rules

In Subtask 2, each segment is judged by the decision rules below in Table 4; the *isnad* and claimed source are judged only against a correct quoted text.

Segment	Case	Label
Matn	Quoted exactly, in any documented narration	correct
	Quoted exactly but cut short at a word boundary	correct
	Quoted exactly, though the Hadith is graded weak or forged	correct
	A word or clause added, dropped, or changed; clauses reordered or two Hadiths merged; or no Hadith exists in this wording	incorrect
Ayah	Quoted exactly; diacritics, where given, also correct	correct
	Quoted partially but exactly and contiguously	correct
	Wrong diacritics; a word dropped, changed, or cut off; or not traceable to any verse	incorrect
Isnad	Names the narrator of this Hadith	correct
	Names a different narrator	incorrect
	Quoted text incorrect	N/A
Claimed source	Names the collection containing this Hadith, or the correct sūra (and verse number, if given)	correct
	Names a collection or reference that does not match	incorrect
	Quoted text incorrect	N/A

Table 4: Decision rules for each segment label

C Sample Data Files Provided to Participants

Figure 3 showcases an example from the dataset given to participants.

```

{
  "id": "R000002",
  "question_id": "Q00019",
  "question": "ما حكم أكل لحم الضفدع",
  "generated_answer": "حكم أكل لحم الضفدع في الشريعة الإسلامية مسألة خلافية بين العلماء، ولا يوجد إجماع قاطع بشأنه. يمكن تلخيص الآراء والحجج المختلفة كالتالي",
  "annotations": [
    {
      "annotation_id": 1,
      "label": "Ayah",
      "segments": [
        {
          "type": "Ayah",
          "span_start": 936,
          "span_end": 969,
          "label": "incorrect"
        },
        {
          "type": "claimed_source",
          "span_start": 972,
          "span_end": 982,
          "label": "N/A"
        }
      ]
    },
    {
      "correction": [
        "فَإِنْ جُفَّتُمْ فَرِحْنَا أَوْ دُبَّانَا " فَإِذَا أَمِيتُمْ فَأَنْكِرُوا "
        "الله كُنَّا عَلَمَكُمْ مَا لَمْ تَكُونُوا تَعْلَمُونَ"
        "كُنَّا أَرْسَلْنَا فِيكُمْ رَسُولًا مِنْكُمْ يَتْلُو عَلَيْكُمْ آيَاتِنَا"
        "وَيُزَكِّيكُمْ وَيُعَلِّمُكُمُ الْكِتَابَ وَالْحِكْمَةَ وَيُعَلِّمُكُمُ مَا"
        "لَمْ تَكُونُوا تَعْلَمُونَ"
      ]
    }
  ]
}

```

Figure 3: Example entry from the development set showing a question, generated answer, and some annotations.

D Performance of models used for data generation

Performance of all 12 models whose answers were annotated over all annotated responses (train, development, and labelled test pool) are shown in Table 5. Answer counts vary per model due to cleaning and test-set annotation sampling. * marks frontier models. Best results are in bold.

E Summary of methods used by Participants

Table 6 summarises the methodologies used by participants in each subtask.

Model	#Ans.	Ayah		Matn		Isnad		Q-Source		H-Source	
		#	Cor.%	#	Cor.%	#	Cor.%	#	Cor.%	#	Cor.%
Gemini-3.1-Pro*	586	1,632	99.33	1,579	88.41	781	97.95	1,613	99.94	1,379	98.91
GPT-5.4*	645	2,439	97.58	1,747	87.81	352	94.89	2,417	99.42	1,519	97.63
ALLaM-7B	703	1,432	89.53	1,401	62.31	278	77.34	1,376	75.44	865	84.97
Fanar-1-9B	687	1,025	80.20	926	48.06	148	31.76	918	63.29	428	70.09
Jais-13B	722	271	59.78	118	32.20	5	20.00	192	86.46	20	75.00
AceGPT-v2-8B	715	260	36.92	320	20.94	3	33.33	124	66.94	54	53.70
Command-R-7B	684	943	25.56	845	4.14	2	0.00	467	53.32	31	48.39
Gemma-3-12B	714	1,744	12.90	1,455	6.39	11	9.09	679	62.00	85	65.88
Qwen3-8B	662	1,295	14.36	1,177	3.48	8	50.00	332	31.02	24	29.17
Llama-3.1-8B	690	1,085	4.70	1,019	1.67	0	—	163	36.81	11	0.00
Mistral-7B	664	1,017	3.24	1,069	0.65	3	0.00	69	26.09	5	20.00
Phi-3.5-mini	718	851	2.94	610	0.49	0	—	25	12.00	2	50.00

Table 5: Performance of all 12 models whose answers were annotated over all annotated responses (train, development, and labelled test pool). Answer counts vary per model due to cleaning and test-set annotation sampling. * marks frontier models. Best results are in bold. The models include: ALLaM (Bari et al., 2024), Jais (Sengupta et al., 2023), Gemma-3 (Team et al., 2025), Llama-3.1 (Grattafiori et al., 2024), Qwen3 (Yang et al., 2025), AceGPT-v2 (Liang et al., 2024), Fanar-1 (Fanar Team, 2025), Mistral (Mistral AI, 2024), Command-R (Alnumay et al., 2025), and Phi-3.5 (Team, 2024), with GPT-5.4 and Gemini-3.1-Pro as frontier references.

Methodology	Subtask 1	Subtask 2	Subtask 3	Subtask 4
Lexical Retrieval & Matching Exact/fuzzy matching, character/word n -grams, TF-IDF, and BM25	—	EIDIANLP, Namaa Community, Qahera, SanadAI, Taqwatech, Tathabbut, Yaqeen	Qahera, Taqwatech, Tathabbut, Yaqeen	HIAST
Feature-Based Classification & Reranking Logistic regression, SVM, random forests, gradient-boosted trees, and learned reranking	—	Tathabbut, Yaqeen	Taqwatech, Tathabbut, Yaqeen	—
Encoder-Based Models BIO sequence labelling, contextual classification, and dense retrieval	Bayan_CU, HIAST, Northwestern 9, Qahera, SanadAI, Taqwatech, Yaqeen	HIAST, SanadAI	—	Bayan, HIAST
LLM-Based Methods Few-shot and CoT prompting, LoRA, distillation, prompt optimisation, and LLM verification	—	—	Qahera (Hadith candidate verification only)	AhmedNLP, Bayan, HaDe, May-X, Qahera
Rules & Post-processing Span reconstruction, metadata checks, data preparation, class balancing, thresholding, and abstention	Bayan_CU, HIAST, Northwestern 9, Qahera, SanadAI, Taqwatech, Yaqeen	EIDIANLP, Namaa Community, Qahera, SanadAI, Taqwatech, Tathabbut, Yaqeen	Qahera, Taqwatech, Tathabbut, Yaqeen	AhmedNLP, May-X, Taqwatech

Table 6: Overview of the methodological families employed by participating teams across the four subtasks. Teams appear in multiple rows when their pipelines combine several approaches. Dashes indicate that a methodology was not used for the corresponding subtask.