

AKARROUCH at DialectSentEval 2026: A Generate-and-Rerank Approach for Arabic Dialect Sentiment Swapping

Mohamed Akarrouch

Faculty of Sciences and Technologies of Tangier

Abdelmalek Essaâdi University

Tetouan 93000, Morocco

mohamed.akarrouch@etu.uae.ac.ma

Abstract

We describe the AKARROUCH submission to DialectSentEval 2026 Subtask 2, which requires reversing the sentiment of an Arabic dialect sentence while preserving its content. We fine-tune AraT5v2 for direction-controlled generation and use it to produce multiple candidates. We then rerank these candidates using CAMELBERT target-sentiment probabilities and source-to-candidate chrF. The system ranked second in the official evaluation, achieving 78.51% sentiment accuracy, 49.82 BLEU, and 70.12 chrF.

1 Introduction

DialectSentEval 2026 Subtask 2 asks systems to reverse the sentiment of an Arabic dialect sentence while retaining the information expressed by the source (Ezzini et al., 2026a). The task extends AraSentEval’s work on Arabic sentiment analysis and swapping (Ezzini et al., 2026b). Our work uses a direction-controlled AraT5v2 generator followed by inference-time reranking. CAMELBERT is used only during inference and is not jointly trained with the generator. Code and the executable notebook are available at <https://github.com/MedAkarrouch/DialectSentEval2026>.

2 Background

Sentiment swapping is a text style transfer problem in which a target attribute is changed while other content is retained as far as possible. The shared-task training and validation records contain a source sentence, its polarity, and an opposite-sentiment target, while the test records provide sources for generation. Parallel resources remain limited for Arabic. Ma’aks provides manually curated Arabic sentence pairs expressing opposite sentiments about the same subject (Mughaus




Figure 1: Overview of the generate-and-rerank inference workflow.

et al., 2026). Prior work shows that limited parallel data can support supervised sentiment transfer (Mukherjee and Dusek, 2023).

Generate-and-rerank methods separate candidate production from final selection. In text style transfer, Suzgun et al. (2022) generate several target-style candidates and rank them using style strength, textual similarity, and fluency. Our system follows this general structure with models and ranking signals chosen for the shared task.

3 System Description

Figure 1 gives an overview. One input produces 20 candidates, which are reranked before the highest-ranked candidate is selected.

3.1 Direction-Controlled Generation

We initialize the generator from UBC-NLP/AraT5v2-base-1024, an Arabic text-to-text model associated with the Octopus toolkit (Elmadany et al., 2023; UBC Deep Learning and NLP Lab, 2023). Fine-tuning treats sentiment swapping as conditional sequence generation: the source sentence and the requested direction form the model input, and the paired sentiment-swapped sentence is the target. All AraT5v2 parameters are updated.

The two transformation directions are represented by special tokens. Table 1 shows how the

Source polarity	Target polarity	Model input
Negative	Positive	<negative_to_positive> {source}
Positive	Negative	<positive_to_negative> {source}

Table 1: Direction-controlled input templates used for fine-tuning and generation.

corresponding token is prepended to the source sentence for each direction.

The complete validation set is evaluated once per epoch for model selection. Training stops after three consecutive epochs without an improvement in validation loss, and the checkpoint with the lowest loss is used for inference. At inference, CAMELBERT compares its positive and negative probabilities for the source. The larger value determines the inferred source polarity, so neutral does not participate in this binary decision. The opposite polarity selects the direction token.

3.2 Candidate Generation and Reranking

The selected AraT5v2 checkpoint uses deterministic beam search with 20 beams to return 20 candidates. No sampling is applied, and the method does not assume that the candidates are unique. For each candidate y_i , CAMELBERT supplies the probability q_i of the requested positive or negative label (Inoue et al., 2021; CAMEL Lab, 2021).

The second score, h_i , is sentence-level chrF between the source and candidate (Popović, 2015). It measures character-level overlap and is used as a content-preservation proxy, not as direct evidence of semantic equivalence. During inference-time evaluation, the gold validation targets are used to calculate BLEU and chrF only after selection. They do not rank the generated candidates.

Algorithm 1 presents the complete decision procedure. Candidates are indexed in beam order. When the preceding scores are identical, the $-i$ term favors the candidate with the lower beam index. Tuple maximization is lexicographic from left to right.

The threshold separates the two ranking priorities. Candidates meeting it are ranked first by source-to-candidate chrF, then target-sentiment probability, and finally beam position. If none reaches the threshold, target-sentiment probability becomes the primary criterion, followed by chrF and beam position.

Algorithm 1 Generate-and-rerank inference

Require: source x , generator G , classifier C

```

1:  $s \leftarrow \arg \max_{\ell \in \{\text{pos}, \text{neg}\}} P_C(\ell \mid x)$ 
2:  $t \leftarrow \text{opposite}(s)$ 
3:  $p \leftarrow \text{DIRECTIONPROMPT}(x, t)$ 
4:  $Y \leftarrow \text{BEAMSEARCH}(G, p, K = 20)$ 
5: for  $i = 1, \dots, 20$  do
6:    $q_i \leftarrow P_C(t \mid y_i)$ 
7:    $h_i \leftarrow \text{chrF}(x, y_i)$ 
8: end for
9:  $S \leftarrow \{i : q_i \geq \tau\}$ , where  $\tau = 0.85$ 
10: if  $S \neq \emptyset$  then
11:    $j \leftarrow \arg \max_{i \in S} (h_i, q_i, -i)$ 
12: else
13:    $j \leftarrow \arg \max_{1 \leq i \leq 20} (q_i, h_i, -i)$ 
14: end if
15: return  $y_j$ 
```

4 Experimental Setup

We use the official training set for supervised fine-tuning and the complete validation set for epoch-level validation loss, early stopping, and checkpoint selection. Fine-tuning is performed in FP16 on two NVIDIA Tesla T4 GPUs. AdamW uses a learning rate of 5×10^{-5} and a linear schedule with 5% warmup. Appendix A reports the remaining reproducibility settings. Core software comprises PyTorch 2.10.0+cu128, Transformers 4.49.0, and SacreBLEU 2.5.1 (Post, 2018).

At inference, the selected checkpoint generates 20 candidates per input and Algorithm 1 selects one prediction. CAMELBERT is used only for source-direction inference and candidate scoring. It contributes no AraT5v2 training loss and receives no updates. The official evaluation reports sentiment accuracy, BLEU (Papineni et al., 2002), and chrF. Sentiment accuracy evaluates the requested polarity, while BLEU and chrF measure overlap with reference outputs.

5 Results and Discussion

Table 2 presents the official test results.

Sentiment accuracy reflects how often an output received the requested polarity under the official evaluation. BLEU and chrF measure overlap with the reference outputs, with chrF providing character-level sensitivity to morphological and orthographic variation. Together, these values report the end-to-end performance of the submitted system.

Metric	Official test result
Sentiment accuracy	78.51%
BLEU	49.82
chrF	70.12

Table 2: Official DialectSentEval 2026 Subtask 2 results for Team AKARROUCH.

Transformation	Model input	Model output
Positive → Negative	هذا الموقع جدا جميل 🥰	هذا الموقع جدا ممل 😞
Negative → Positive	تطبيق متجه للفشل 🤡 جدا فاشل	تطبيق متجه للنجاح 🤖 جدا ناجح

Table 3: Examples of sentiment swapping produced by the generate-and-rerank system.

Table 3 shows one positive-to-negative example and one negative-to-positive example. For each input, the model generates 20 candidates, reranks them using Algorithm 1, and selects the highest-ranked candidate as the output.

6 Conclusion

We presented direction-controlled AraT5v2 generation with CAMELBERT- and chrF-based reranking. Future work could either modify fine-tuning with explicit task-specific penalty terms for sentiment reversal and content preservation, or apply preference optimization such as DPO over chosen and rejected generated candidates (Rafailov et al., 2023). Controlled inference ablations should also examine candidate count and beam width, deterministic beam search versus sampling, sentiment threshold, and length penalty.

Limitations

Because CAMELBERT determines source polarity and scores candidates, its errors could affect both transformation-direction selection and reranking. Reliance on the same classifier used by the official sentiment metric may align the system with its decision boundary rather than broader human sentiment judgments. For the submitted model, we used a fixed inference configuration of 20 deterministic candidates, no sampling, and a 0.85 threshold. We did not conduct controlled ablations of these settings, and the selected configuration may therefore be suboptimal.

Acknowledgments

We thank the DialectSentEval organizers, contributors, and reviewers. We used ChatGPT for coding assistance, including the generation of individual functions. We used AI assistants for rephrasing and grammatical error correction in our writing.

References

- CAMEL Lab. 2021. [CAMELBERT-DA sentiment analysis model card](#). Hugging Face model repository. Accessed 1 August 2026.
- AbdelRahim Elmadany, El Moatez Billah Nagoudi, and Muhammad Abdul-Mageed. 2023. [Octopus: A multitask model and toolkit for Arabic natural language generation](#). In *Proceedings of ArabicNLP 2023*, pages 232--243, Singapore (Hybrid). Association for Computational Linguistics.
- Saad Ezzini, Shadi Abudalfa, Maram Alharbi, Salmane Chafik, Hind Alatawi, Mo El-Haj, Ahmed Abdelali, Osamah Alnahari, and Salima Lamsiyah. 2026a. Dialectsenteval 2026: Arabic dialect sentiment analysis and swapping shared task. In *Proceedings of the Fourth Arabic Natural Language Processing Conference*, Budapest, Hungary.
- Saad Ezzini, Shadi Abudalfa, Maram Alharbi, Salmane Chafik, Hamzah Luqman, Mo El-Haj, Paul Rayson, Samaher Alghamdi, and Reem Alotaibi. 2026b. AraSentEval: A shared task on sentiment analysis and swapping in Arabic. In *Proceedings of the 7th Workshop on Open-Source Arabic Corpora and Processing Tools (OSACT7), co-located with the Language Resources and Evaluation Conference (LREC 2026)*, Palma, Mallorca, Spain.
- Go Inoue, Bashar Alhafni, Nurpeiis Baimukan, Houda Bouamor, and Nizar Habash. 2021. [The interplay of variant, size, and task type in Arabic pre-trained language models](#). In *Proceedings of the Sixth Arabic Natural Language Processing Workshop*, pages 92--104, Kyiv, Ukraine (Virtual). Association for Computational Linguistics.
- Raed Mughaus, Shadi Abudalfa, Hamzah Luqman, Fahad Abdu, Mohammed AlAli, Nawaf Al-Dowayan, and Ahmed Abdelali. 2026. [Ma'aks: manually-curated parallel dataset for](#)

Arabic text sentiment swap. *Language Resources and Evaluation*, 60(1):1.

Sourabrata Mukherjee and Ondrej Dusek. 2023. [Leveraging low-resource parallel data for text style transfer](#). In *Proceedings of the 16th International Natural Language Generation Conference*, pages 388--395, Prague, Czechia. Association for Computational Linguistics.

Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. [BLEU: a method for automatic evaluation of machine translation](#). In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311--318, Philadelphia, Pennsylvania, USA. Association for Computational Linguistics.

Maja Popović. 2015. [chrF: character n-gram F-score for automatic MT evaluation](#). In *Proceedings of the Tenth Workshop on Statistical Machine Translation*, pages 392--395, Lisbon, Portugal. Association for Computational Linguistics.

Matt Post. 2018. [A call for clarity in reporting BLEU scores](#). In *Proceedings of the Third Conference on Machine Translation: Research Papers*, pages 186--191, Brussels, Belgium. Association for Computational Linguistics.

Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D. Manning, Stefano Ermon, and Chelsea Finn. 2023. [Direct preference optimization: Your language model is secretly a reward model](#). In *Advances in Neural Information Processing Systems*, volume 36, pages 53728--53741.

Mirac Suzgun, Luke Melas-Kyriazi, and Dan Jurafsky. 2022. [Prompt-and-rerank: A method for zero-shot and few-shot arbitrary textual style transfer with small language models](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 2195--2222, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.

UBC Deep Learning and NLP Lab. 2023. [AraT5v2-base-1024 model card](#). Hugging Face model repository. Accessed 1 August 2026.

A Implementation Details

Tables 4 and 5 list the remaining training and inference settings.

Setting	Value
Training examples	6,752
Validation examples	964
Test examples	1,931
Per-device training batch	4
Gradient accumulation	2
Weight decay	0.01
Gradient clipping	1.0
Maximum source length	256 tokens
Maximum target length	256 tokens
Maximum epochs	20
Early-stopping patience	3 epochs
Random seed	42

Table 4: Data and supervised fine-tuning configuration.

Setting	Value
Decoding method	Deterministic beam search
Number of beams	20
Returned candidates	20
Beam-search early stopping	Enabled
Length penalty	1.0
Minimum new tokens	3
Maximum new tokens	128
Target-sentiment threshold	0.85

Table 5: Inference and reranking configuration.