

NAMAA at StanceEval-2026: From Encoder and Decoder Model Fine-Tuning to Data-Centric Retrieval-Augmented In-Context Learning for Arabic Stance Detection

Abdelbasset Djamai^{1,6} Hassan Barmandah^{2,6} Nour Qandos³
Samar Ahmed^{4,6} Omer Nacar^{5,6}

¹Datategy ²Umm Al-Qura University ³Independent Researcher

⁴King Salman Global Academy for the Arabic Language

⁵Tuwaiq Academy ⁶NAMAA Community

¹contact@adjamai.com ²hassanhbarmandah@gmail.com ³nooramerg0@gmail.com

⁴samar.sass6@gmail.com ⁵o.najar@tuwaiq.edu.sa

Abstract

Stance detection identifies an author’s position toward a specified target, yet Arabic social-media text remains challenging because stance is often implicit and shaped by dialect, sarcasm, informal expression, and culturally specific references. StanceEval-2026 evaluates this problem in two transfer settings: a related target within a seen thematic domain (Track 1) and entirely unseen topics (Track 2). We compare four modeling strategies: single-task Arabic encoder fine-tuning, multi-task learning with sentiment and sarcasm supervision, 4-bit QLoRA adaptation of the ALLaM-7B model, and Retrieval-Augmented In-Context Learning (RAG-ICL) with frozen large language models (LLMs). Our strongest single-backend configuration is a data-centric RAG-ICL pipeline combining label-balanced retrieval, hashtag and negation adaptation, text cleaning, similarity filtering, and target descriptions. A majority-vote ensemble of this configuration across three LLM backends obtains official F_{avg2} scores of 0.8719 on Track 1 and 0.9016 on Track 2, ranking 3rd and 7th, respectively. Under the shared-task test conditions, these findings suggest that training-free RAG-ICL is particularly effective for unseen-topic transfer, outperforming the supervised systems reported here.

1 Introduction

Stance detection determines whether an author supports, opposes, or expresses no clear position toward a specified target. Arabic social-media stance is often implicit and complicated by dialectal variation, informal orthography, sarcasm, and culturally specific references. The Mawqif dataset established a benchmark for target-specific Arabic stance, sentiment, and sarcasm annotation (Alturayef et al., 2022), while Mawqif-XT extends this setting to

cross-target evaluation (Albalawi et al., 2026b). StanceEval introduced a shared-task setting in 2024 (Alturayef et al., 2024b), followed by a second edition in 2026 (Albalawi et al., 2026a).

This paper presents our systems for the StanceEval-2026 shared task and examines how different adaptation strategies support cross-target generalization. We compare four approaches spanning single-task encoder fine-tuning, multi-task learning (MTL) with sentiment and sarcasm supervision, QLoRA-based adaptation of ALLaM-7B, and retrieval-augmented in-context learning (RAG-ICL) with frozen large language models (LLMs).

Our key contributions are:

- A comparison of supervised fine-tuning and frozen-LLM prompting across related-target and unseen-topic transfer settings;
- A staged RAG-ICL evaluation covering label-balanced retrieval, task-specific prompt adaptation, data-centric cleaning and target descriptions, similarity filtering, and a separate rationale-augmented variant; and
- A majority-vote ensemble of three LLM backbones, based on our best method, that obtains official F_{avg2} scores of 0.8719 on Track 1 and 0.9016 on Track 2, ranking 3rd and 7th, respectively.

2 Related Work

SemEval-2016 Task 6 established the tweet-level *Favor/Against/None* formulation (Mohammad et al., 2016), followed by target-conditioned encoders that modeled post-target interactions explicitly (Augenstein et al., 2016). Subsequent work addressed zero-shot and cross-target transfer through generalized

topic representations and transferable semantic-emotion knowledge (Allaway and McKeown, 2020; Zhang et al., 2020).

Arabic resources range from claim-document fact-checking in AraStance (Alhindi et al., 2021) to target-specific social-media annotation in Mawqif (Alturayef et al., 2022), which comprises 4,121 tweets labeled for stance, sentiment, and sarcasm toward three targets. Mawqif supported StanceEval 2024, while Mawqif-XT extends evaluation to unseen targets (Alturayef et al., 2024b; Albalawi et al., 2026b). Closely related work on Arabic online polarization likewise compared encoder fine-tuning, zero-shot prompting, and RAG-ICL, finding retrieval-augmented prompting competitive with discriminative models (Djamaï et al., 2026).

Arabic pretrained encoders have played a central role in tackling stance detection and related tasks. Prominent model families include AraBERT, ARBERT/MARBERT, and the more recent AraModernBERT (Antoun et al., 2020; Abdul-Mageed et al., 2021; Elshehy et al., 2026).

StanceEval-2024 highlighted the strength of encoder-based approaches for Arabic stance detection. The winning system combined per-target AraBERT models, multi-task supervision for stance, sentiment, and sarcasm, and contrastive-loss ensembling, while other competitive entries relied on single-task fine-tuning or MTL ensembles (Badran et al., 2024; Melhem et al., 2024; Li and Caragea, 2019; Alturayef et al., 2024a; Shukla et al., 2024). LLM-based systems produced mixed results: fine-tuned GPT-3.5-Turbo was highly competitive, whereas LoRA-tuned Llama-3 and sentence-transformer approaches were less effective (Alghaslan and Almutairy, 2024; Al Hariri and Abu Farha, 2024). In this work, we compare encoder fine-tuning and multi-task learning with retrieval-augmented in-context learning, using query-specific demonstrations and ensembles of frozen LLMs.

3 Task and Dataset

StanceEval-2026 formulates target-conditioned Arabic stance detection: given an Arabic tweet and a specified target, a system predicts *Favor*, *Against*, or *None*. Track 1 (T1) evaluates transfer to a related target within a thematic domain represented during training, whereas Track 2 (T2) evaluates generalization to entirely unseen topics

(Albalawi et al., 2026a).

The task data build on Mawqif, a corpus of 4,121 Arabic tweets covering COVID-19 Vaccine (1,373), Digital Transformation (1,348), and Women Empowerment (1,400). Each tweet is annotated independently for stance, sentiment, and sarcasm. The Mawqif-XT cross-target extension adds 996 manually annotated tweets from Women Driving (352), E-Cars (332), and Trimester System (312), while retaining the same three annotation tasks (Albalawi et al., 2026b).

The official metric averages the class-specific F1 scores of the two polar stance classes:

$$F_{\text{avg2}} = \frac{F_1^{\text{Favor}} + F_1^{\text{Against}}}{2}. \quad (1)$$

4 Methods

Our experiments span single-task encoders, multi-task encoders, adaptation of a decoder model, and retrieval-augmented prompting with frozen LLMs. Figure 1 provides a unified overview of these model families.

4.1 Single-Task Encoder Fine-Tuning

We fine-tuned multiple Arabic-centric and multilingual encoders, retaining MARBERTv2 for separate experiments with supervised contrastive learning (SupCon), focal loss, and R-Drop (Khosla et al., 2020; Lin et al., 2017; Liang et al., 2021). On AraBERT-base-Twitter, we jointly combined cross-entropy (CE), supervised contrastive (SupConF), and center loss using learned uncertainty-based weighting (Wen et al., 2016; Kendall et al., 2018). Loss definitions and explanations appear in Section A.2; model and training details appear in Sections A.1, A.3 and A.4.

4.2 Multi-Task Encoder Learning

We evaluate two formulations that jointly learn stance, sentiment, and sarcasm. Parallel MARBERTv2 uses separate, directly supervised heads on a shared representation, with a weighted sum of cross-entropy losses; its contrastive variant additionally applies stance SupCon.

The stacked AlexUNLP-BH architecture (Badran et al., 2024), reproduced with AraBERT-base-Twitter, concatenates the logits of the three task heads for final stance prediction. Sentiment and sarcasm heads receive direct supervision; stance supervision is applied only to the stacked head, training the parallel stance head indirectly. Track 2




Figure 1: Conceptual overview of the explored method families.

additionally uses supervised NTXentLoss. We compare tweet-only and target-tweet inputs.

Parallel MARBERTv2 coefficients were set heuristically. For stacked AraBERT, we retained AlexUNLP-BH’s effective task weights and heuristically reduced the Track 2 contrastive coefficient. Neither formulation underwent systematic coefficient tuning. Full objectives and training settings appear in Sections A.4 and A.5; individual loss definitions appear in Section A.2.

4.3 LLM Fine-Tuning with QLoRA

We adapt ALLaM-7B-Instruct-preview as a three-class classifier using 4-bit NF4 QLoRA (Dettmers et al., 2023). Each instance uses the Arabic tweet-target classification template shown in Section A.6. The trainable adapters and score head minimize the unweighted three-way stance cross-entropy, while the quantized base-model weights remain frozen. The complete adapter, quantization, input, and optimization settings are provided in Section A.6.

4.4 Retrieval-Augmented In-Context Learning

These ICL methods use frozen LLMs without any training. All prompt templates used in our experiments are provided in Section D.

Baseline and retrieval. We first compare four LLMs, in a zero-shot setting, on the dev set and select the best-performing model (which was found to be Gemma 4 31B) as the anchor for the subsequent controlled RAG-ICL experiments. The direct-label prompt distinguishes stance from general sentiment and requires exactly one of the three labels. *Vanilla RAG-ICL* embeds the concatenated target and post, retrieves demonstrations by cosine similarity, balances them across labels, and inserts them before

the test instance.

The staged *Vanilla–Adapted–Data-centric* comparison was conducted on Track 1. For Track 2, we evaluated the zero-shot baseline and the final data-centric configuration; intermediate RAG variants were not evaluated on that track.

Domain and data-centric adaptation. The *adapted* variant adds neutral hashtag glosses, includes extracted hashtags in the embedding query, and explicitly highlights Arabic negation markers and their scope.

The final *data-centric* variant additionally removes URLs and mentions before embedding and injects concise Arabic target definitions to disambiguate short English labels. For demonstration retrieval only, Track 1 maps *Women Driving* to the related training target *Women Empowerment*; Track 2 retrieves from the provided training demonstrations without explicit cross-target mapping. The original target and post remain unchanged in the LLM prompt. Retrieval mappings, hashtag glosses, and target definitions were manually validated and fixed without access to gold test labels.

Retrieval retains at most $K_{\max} = 2$ examples per label with cosine similarity at least $\tau = 0.55$, a global fixed threshold chosen heuristically rather than systematically tuned. Resource construction, examples, and detailed retrieval settings appear in Appendix A.7.

Rationale variant and ensemble. We also test RAG-ICL with *rationale-augmented chain-of-thought (CoT)* in which three independently retrieved prompts request an explicit rationale followed by a label; branch predictions are majority-voted. This variant is evaluated as a separate branch and is not incorporated into the final data-centric

Method	LLM backend	Track 1	Track 2
MULTI-TASK ENCODERS			
Stacked AraBERT MTL, tweet only	—	0.6707	0.6308
Stacked AraBERT MTL, target + tweet	—	0.6643	0.6717
LLM FINE-TUNING			
ALLaM-7B w/ QLoRA	—	0.6794	—
DIRECT IN-CONTEXT PROMPTING			
Zero-shot baseline	Gemma 4 31B	0.8104	0.8691
RAG-ICL VARIANTS			
↪ Vanilla	Gemma 4 31B	0.8229	—
↪ Adapted	Gemma 4 31B	0.8307	—
↪ Data-centric	GPT-4o	0.8048	0.8830
	Gemma 4 31B	0.8570	0.8931
	Gemini 3.1 Pro	<u>0.8590</u>	<u>0.8954</u>
↪ Rationale CoT	Gemma 4 31B	0.8131	—
ENSEMBLE			
Heterogeneous ensemble	Three LLMs	0.8719	0.9016

Table 1: Official test results under $F_{\text{avg}2}$. Bold and underlining mark the best and second-best scores.

pipeline. Our final offline ensemble instead applies the data-centric configuration independently with Gemini 3.1 Pro, Gemma 4 31B, and GPT-4o, then takes a tweet-level majority vote. A three-way tie is resolved using Gemini 3.1 Pro.

5 Results and Discussion

Appendix B presents and analyzes the development results for the supervised systems and the preliminary zero-shot LLM comparison.

Table 1 reports the official test results. Zero-shot Gemma remained relatively stable on Track 2, decreasing from 0.8773 on development to 0.8691 on test, a difference of only 0.0082. This consistency is encouraging for cross-target transfer, although the development and test splits differ in target composition and label distribution, so the gap should be interpreted descriptively rather than as a controlled estimate of generalization.

The stacked AraBERT MTL systems were considerably weaker on test. The tweet-only variant obtained 0.6707 on Track 1 and 0.6308 on Track 2, while adding the target produced 0.6643 and 0.6717, respectively. Target conditioning therefore had a track-dependent effect in this architecture, slightly reducing Track 1 performance while improving Track 2 by 0.0409.

In contrast, the staged RAG-ICL experiments showed consistent gains for Gemma on Track 1: performance increased from 0.8104 in the zero-shot setting to 0.8229 with Vanilla RAG-ICL, 0.8307 with the adapted variant, and 0.8570 with the final data-centric configuration. The same configuration also improved Track 2 from 0.8691 to 0.8931.

These results indicate that retrieval and task-specific data adaptation provided substantial gains over direct prompting, although the joint contribution of the individual data-centric components cannot be isolated from this comparison alone.

Backend choice also affected performance. Gemini 3.1 Pro was the strongest individual data-centric backend, scoring 0.8590 on Track 1 and 0.8954 on Track 2. Majority voting across Gemini, Gemma, and GPT-4o improved these scores to 0.8719 and 0.9016. By contrast, rationale augmentation reduced Track 1 performance to 0.8131, below Vanilla RAG-ICL at 0.8229, suggesting no benefit from generating rationales before classification.

6 Error Analysis

We examined 30 target-balanced Track 1 development disagreements, with 10 instances per target. Twenty involved confusion between *None* and a polar label; in 17 cases, all three LLM backends agreed on the same non-gold label, and in 21 none predicted the gold label. The disagreements involved implicit stance, sarcasm, target–actor distinctions, and mixed positions. The textual evidence also suggests a likely annotation inconsistency in one displayed case, distinguishing benchmark disagreement from interpretive error. Because the sample includes only ensemble–gold disagreements and was not independently re-adjudicated, these findings illustrate difficult decision boundaries rather than corpus-level error rates. A detailed analysis with examples appear in Section C.

7 Conclusion

We compared supervised fine-tuning, multi-task learning, QLoRA adaptation, and retrieval-augmented ICL for Arabic stance detection. Data-centric RAG-ICL performed best among single backends, while majority voting across three LLMs further improved performance to 0.8719 on Track 1 and 0.9016 on Track 2, suggesting that retrieval-based inference can be highly effective for cross-target Arabic stance detection without additional model training. Future work should analyze the supervised dev-test gap and isolate the contribution of individual retrieval components.

Limitations

The data-centric RAG-ICL stage changes several components jointly. Consequently, the effect of

each retrieval intervention cannot be identified independently, and one-factor ablations are needed to separate their effects. We also do not compare RAG-ICL with a static few-shot baseline using a fixed demonstration set. Thus, our results do not isolate the benefit of query-specific retrieval from that of including labeled demonstrations. A controlled comparison with static few-shot prompting is left for future work. Hosted LLM revisions may also affect reproducibility, while manually curated cross-target retrieval mappings, hashtag glosses, and target definitions introduce domain knowledge that may not transfer directly to other datasets or unseen targets without a closely related training target. The error analysis is restricted to 30 selected Track 1 development-set disagreements and does not include independent re-adjudication; hence, it cannot estimate corpus-level phenomenon frequencies or annotation-error rates.

Our ALLaM-7B experiments are limited to 4-bit QLoRA; time constraints precluded a matched comparison using non-quantized base-model weights. We therefore cannot isolate the effect of quantization on stance detection performance.

Ethical Considerations

Stance predictions concern sensitive opinions and may be especially error-prone for dialectal, sarcastic, or culturally situated speech. The systems should not be used to infer personal beliefs, moderate users automatically, or support consequential decisions without human review.

Acknowledgments

We thank Tuwaiq Academy for its continued support of this work. We also thank the StanceEval-2026 organizers for coordinating the shared task and providing the data and evaluation framework.

References

- Ahmed Abdelali, Sabit Hassan, Hamdy Mubarak, Kareem Darwish, and Younes Samih. 2021. [Pre-training BERT on Arabic tweets: Practical considerations](#). *arXiv preprint arXiv:2102.10684*.
- Muhammad Abdul-Mageed, AbdelRahim Elmadany, and El Moatez Billah Nagoudi. 2021. [ARBERT & MARBERT: Deep bidirectional transformers for Arabic](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7088–7105, Online. Association for Computational Linguistics.
- Youssef Al Hariri and Ibrahim Abu Farha. 2024. [SMASH at StanceEval 2024: Prompt engineering LLMs for Arabic stance detection](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 800–806, Bangkok, Thailand. Association for Computational Linguistics.
- Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Saad Ezzini, Maram Kurdi, Asma Yamani, Ahmed Ashraf, Maged Al-Shaibani, and Nora Alturayef. 2026a. StanceEval-2026: The second stance detection shared task. In *Proceedings of the Fourth Arabic Natural Language Processing Conference (ArabicNLP 2026)*, Budapest, Hungary. Association for Computational Linguistics.
- Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Maram Kurdi, Saad Ezzini, Asma Yamani, and Ahmed Ashraf. 2026b. [Mawqif-XT: An Arabic benchmark dataset for cross-target stance detection](#). *Preprint*, arXiv:2608.09539.
- Mamoun Alghaslan and Khaled Almutairy. 2024. [MGKM at StanceEval2024 fine-tuning large language models for Arabic stance detection](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 816–822, Bangkok, Thailand. Association for Computational Linguistics.
- Tariq Alhindi, Amal Alabdulkarim, Ali Alshehri, Muhammad Abdul-Mageed, and Preslav Nakov. 2021. [AraStance: A multi-country and multi-domain dataset of Arabic stance detection for fact checking](#). In *Proceedings of the Fourth Workshop on NLP for Internet Freedom: Censorship, Disinformation, and Propaganda*, pages 57–65, Online. Association for Computational Linguistics.
- Emily Allaway and Kathleen McKeown. 2020. [Zero-shot stance detection: A dataset and model using generalized topic representations](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 8913–8931, Online. Association for Computational Linguistics.
- Nora Alturayef, Hamzah Luqman, and Moataz Ahmed. 2024a. [Enhancing stance detection through sequential weighted multi-task learning](#). *Social Network Analysis and Mining*, 14(1):7.
- Nora Alturayef, Hamzah Luqman, Zaid Alyafeai, and Asma Yamani. 2024b. StanceEval 2024: The first Arabic stance detection shared task. In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 774–782, Bangkok, Thailand. Association for Computational Linguistics.
- Nora Saleh Alturayef, Hamzah Abdullah Luqman, and Moataz Aly Kamaleldin Ahmed. 2022. [Mawqif: A multi-label Arabic dataset for target-specific stance detection](#). In *Proceedings of the Seventh Arabic Natural Language Processing Workshop (WANLP)*, pages 174–184, Abu Dhabi, United Arab Emirates (Hybrid). Association for Computational Linguistics.

- Wissam Antoun, Fady Baly, and Hazem Hajj. 2020. [AraBERT: Transformer-based model for Arabic language understanding](#). In *Proceedings of the 4th Workshop on Open-Source Arabic Corpora and Processing Tools, with a Shared Task on Offensive Language Detection*, pages 9–15, Marseille, France. European Language Resources Association.
- Isabelle Augenstein, Tim Rocktäschel, Andreas Vlachos, and Kalina Bontcheva. 2016. [Stance detection with bidirectional conditional encoding](#). In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 876–885, Austin, Texas. Association for Computational Linguistics.
- Mohamed Badran, Mo'men Hamdy, Marwan Torki, and Nagwa El-Makky. 2024. [AlexUNLP-BH at StanceEval2024: Multiple contrastive losses ensemble strategy with multi-task learning for stance detection in Arabic](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 823–827, Bangkok, Thailand. Association for Computational Linguistics.
- Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. 2020. A simple framework for contrastive learning of visual representations. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 1597–1607. PMLR.
- Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. 2020. [Unsupervised cross-lingual representation learning at scale](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 8440–8451, Online. Association for Computational Linguistics.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. [QLoRA: Efficient fine-tuning of quantized LLMs](#). In *Advances in Neural Information Processing Systems*, volume 36, pages 10088–10115.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. [BERT: Pre-training of deep bidirectional transformers for language understanding](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.
- Abdelbasset Djamai, Sahara Al-Madi, Norah Al-Zaid, Khlood Al Jallad, and Mona Azim. 2026. [NAMAA at SemEval-2026 task 9: Comparing generative, retrieval-augmented, and discriminative methods for Arabic online polarization detection and type classification](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 3562–3590, San Diego, California, USA. Association for Computational Linguistics.
- AbdelRahim Elmadany, ElMoatez Billah Nagoudi, and Muhammad Abdul-Mageed. 2023. [ORCA: A challenging benchmark for Arabic language understanding](#). In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 9559–9586, Toronto, Canada. Association for Computational Linguistics.
- Omar Elshehy, Omer Nacar, Abdelbasset Djamai, Muhammed Ragab, Khlood Al Jallad, and Mona Abdelazim. 2026. [AraModernBERT: Transtokenized initialization and long-context encoder modeling for Arabic](#). In *Proceedings of the 2nd Workshop on NLP for Languages Using Arabic Script*, pages 313–321, Rabat, Morocco. Association for Computational Linguistics.
- Go Inoue, Bashar Alhafni, Nurpeiis Baimukan, Houda Bouamor, and Nizar Habash. 2021. [The interplay of variant, size, and task type in Arabic pre-trained language models](#). In *Proceedings of the Sixth Arabic Natural Language Processing Workshop*, pages 92–104, Kyiv, Ukraine (Virtual). Association for Computational Linguistics.
- Alex Kendall, Yarin Gal, and Roberto Cipolla. 2018. [Multi-task learning using uncertainty to weigh losses for scene geometry and semantics](#). In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 7482–7491.
- Prannay Khosla, Piotr Teterwak, Chen Wang, Aaron Sarna, Yonglong Tian, Phillip Isola, Aaron Maschinot, Ce Liu, and Dilip Krishnan. 2020. Supervised contrastive learning. In *Advances in Neural Information Processing Systems*, volume 33, pages 18661–18673.
- Yingjie Li and Cornelia Caragea. 2019. [Multi-task stance detection with sentiment and stance lexicons](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 6299–6305, Hong Kong, China. Association for Computational Linguistics.
- Xiaobo Liang, Lijun Wu, Juntao Li, Yue Wang, Qi Meng, Tao Qin, Wei Chen, Min Zhang, and Tie-Yan Liu. 2021. R-Drop: Regularized dropout for neural networks. In *Advances in Neural Information Processing Systems*, volume 34, pages 10890–10905.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. 2017. [Focal loss for dense object detection](#). In *Proceedings of the IEEE International Conference on Computer Vision*, pages 2980–2988.
- Anas Melhem, Osama Hamed, and Thaer Sammar. 2024. [TAO at StanceEval2024 shared task: Arabic stance detection using AraBERT](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 842–846, Bangkok, Thailand. Association for Computational Linguistics.
- Saif Mohammad, Svetlana Kiritchenko, Parinaz Sobhani, Xiaodan Zhu, and Colin Cherry. 2016. [SemEval-2016 task 6: Detecting stance in tweets](#). In *Proceedings of*

the 10th International Workshop on Semantic Evaluation (SemEval-2016), pages 31–41, San Diego, California. Association for Computational Linguistics.

Kevin Musgrave, Serge Belongie, and Ser-Nam Lim. 2020. PyTorch Metric Learning. *arXiv preprint arXiv:2008.09164*.

Ishaan Shukla, Ankit Vaidya, and Geetanjali Kale. 2024. PICT at StanceEval2024: Stance detection in Arabic using ensemble of large language models. In *Proceedings of the Second Arabic Natural Language Processing Conference*, pages 837–841, Bangkok, Thailand. Association for Computational Linguistics.

Yandong Wen, Kaipeng Zhang, Zhifeng Li, and Yu Qiao. 2016. A discriminative feature learning approach for deep face recognition. In *Computer Vision—ECCV 2016*, pages 499–515. Springer.

Puxuan Yu, Luke Merrick, Gaurav Nuti, and Daniel F. Campos. 2025. Arctic-Embed 2.0: Multilingual retrieval without compromise. In *Proceedings of the Second Conference on Language Modeling*.

Bowen Zhang, Min Yang, Xutao Li, Yunming Ye, Xiaofei Xu, and Kuai Dai. 2020. Enhancing cross-target stance detection with transferable semantic-emotion knowledge. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 3188–3197, Online. Association for Computational Linguistics.

A Implementation Details, Experimental Settings, and Reproducibility

A.1 Encoder Checkpoints and References

Table 2 lists every distinct pretrained encoder used in our experiments. The stacked AraBERT multi-task system reused the AraBERT-base-Twitter checkpoint.

A.2 Training Objectives

This section defines the individual loss terms referenced in Section 4. Let a minibatch contain B instances, with gold class y_i , logits $\mathbf{z}_i$, and $p_{i,c} = \text{softmax}(\mathbf{z}_i)_c$. For the binary sarcasm head, let $s_i \in \{0, 1\}$ and $q_i = \text{sigmoid}(a_i)$.

Classification losses. With $W_B = \sum_i w_{y_i}$, the weighted multi-class and binary losses are

$$\mathcal{L}_{\text{CE}}^{(w)} = -\frac{1}{W_B} \sum_{i=1}^B w_{y_i} \log p_{i,y_i}, \quad (2)$$

$$\mathcal{L}_{\text{BCE}}^{(v)} = -\frac{1}{B} \sum_{i=1}^B \left[v_1 s_i \log q_i + v_0 (1 - s_i) \log (1 - q_i) \right]. \quad (3)$$

Unweighted losses set all class weights to one. The Track 1 stacked model uses $w_c \propto 1/n_c$ for stance and sentiment, where n_c is the training count of class c , and class-weighted BCE for sarcasm.

Supervised contrastive and focal losses. Supervised contrastive learning (SupCon) encourages representations with the same stance label to cluster together and separates those with different labels. Its loss is denoted by $\mathcal{L}_{\text{SC}}$. Focal loss downweights confidently classified examples, increasing the relative contribution of difficult examples.

For normalized pooled representations $\mathbf{u}_i$, define $A(i) = \{1, \dots, B\} \setminus \{i\}$, $P(i) = \{p \in A(i) : y_p = y_i\}$, and $I_+ = \{i : |P(i)| > 0\}$. Let $\tau_{\text{con}} > 0$ denote the contrastive temperature, distinct from the retrieval similarity threshold τ . The supervised contrastive objective of Khosla et al. (2020) is

$$\begin{aligned} \ell_i^{\text{SC}} &= -\frac{1}{|P(i)|} \sum_{p \in P(i)} \exp(\mathbf{u}_i^\top \mathbf{u}_p / \tau_{\text{con}}) \\ &\quad \times \log \frac{\exp(\mathbf{u}_i^\top \mathbf{u}_p / \tau_{\text{con}})}{\sum_{a \in A(i)} \exp(\mathbf{u}_i^\top \mathbf{u}_a / \tau_{\text{con}})}, \\ \mathcal{L}_{\text{SC}} &= \frac{1}{|I_+|} \sum_{i \in I_+} \ell_i^{\text{SC}}. \end{aligned} \quad (4)$$

The MARBERTv2 SupCon variant uses $\mathcal{L}_{\text{CE}} + 0.1\mathcal{L}_{\text{SC}}$ with $\tau_{\text{con}} = 0.07$. Its focal-loss variant follows Lin et al. (2017) and replaces cross-entropy with

$$\mathcal{L}_{\text{focal}} = -\frac{1}{B} \sum_{i=1}^B (1 - p_{i,y_i})^\gamma \log p_{i,y_i}, \quad \gamma = 2. \quad (5)$$

R-Drop. Two independently dropped-out passes produce distributions $\mathbf{p}_i^{(1)}$ and $\mathbf{p}_i^{(2)}$. Following Liang et al. (2021), we use the symmetric KL penalty

$$\begin{aligned} D_{\text{SKL}}(\mathbf{p}, \mathbf{q}) &= \frac{1}{2} D_{\text{KL}}(\mathbf{p} \parallel \mathbf{q}) + \frac{1}{2} D_{\text{KL}}(\mathbf{q} \parallel \mathbf{p}), \\ \mathcal{L}_{\text{RD}} &= \frac{1}{2} \left(\mathcal{L}_{\text{CE}}^{(1)} + \mathcal{L}_{\text{CE}}^{(2)} \right) \\ &\quad + \frac{\alpha}{B} \sum_{i=1}^B D_{\text{SKL}}(\mathbf{p}_i^{(1)}, \mathbf{p}_i^{(2)}), \\ \alpha &= 4. \end{aligned} \quad (6)$$

Center loss and automatic weighting. Center loss encourages each representation to remain close to a learned center for its stance class. Uncertainty-based weighting learns the relative weights of the combined loss terms during training.

Model used	Exact Hugging Face checkpoint	Primary reference(s)
AraBERT-base-Twitter	aubmindlab/bert-base-arabertv02-twitter	(Antoun et al., 2020)
AraBERT-large-Twitter	aubmindlab/bert-large-arabertv02-twitter	(Antoun et al., 2020)
AraBERT-large-MSA	aubmindlab/bert-large-arabertv02	(Antoun et al., 2020)
ARBERTv2	UBC-NLP/ARBERTv2	(Abdul-Mageed et al., 2021; Elmadany et al., 2023)
MARBERTv2	UBC-NLP/MARBERTv2	(Abdul-Mageed et al., 2021)
CAMeLBERT-DA	CAMEL-Lab/bert-base-arabic-camelbert-da	(Inoue et al., 2021)
QARiB	ahmedabdelali/bert-base-qarib	(Abdelali et al., 2021)
mBERT	google-bert/bert-base-multilingual-cased	(Devlin et al., 2019)
XLM-R-large	FacebookAI/xlm-roberta-large	(Conneau et al., 2020)

Table 2: Exact pretrained encoder checkpoints used in our experiments, with official Hugging Face repositories and primary references.

For pooled representation $\mathbf{h}_i$ and learned stance center $\mathbf{c}_{y_i}$, we use a batch-averaged formulation of the center loss of Wen et al. (2016):

$$\mathcal{L}_{\text{ctr}} = \frac{1}{2B} \sum_{i=1}^B \|\mathbf{h}_i - \mathbf{c}_{y_i}\|_2^2. \quad (7)$$

The AraBERT joint-loss experiment learns an unconstrained log-scale s_m for each component using Kendall-style uncertainty-based weighting (Kendall et al., 2018):

$$\mathcal{L}_{\text{auto}} = \sum_{m \in \{\text{CE}, \text{SC}, \text{ctr}\}} (e^{-s_m} \mathcal{L}_m + s_m). \quad (8)$$

Supervised NTXentLoss. The stacked Track 2 experiment uses supervised normalized temperature-scaled cross-entropy ($\mathcal{L}_{\text{NTX}}$), implemented as NTXentLoss in PyTorch Metric Learning (Chen et al., 2020; Musgrave et al., 2020).

Unlike the multi-positive denominator in Equation (4), the library implementation evaluates each ordered same-label pair separately. Let $\Pi = \{(i, p) : i \neq p, y_i = y_p\}$, $N(i) = \{a : y_a \neq y_i\}$, and let d_{ij} be cosine similarity. With $\tau_{\text{con}} = 0.07$,

$$Z_{ip} = e^{d_{ip}/\tau_{\text{con}}} + \sum_{a \in N(i)} e^{d_{ia}/\tau_{\text{con}}},$$

$$\mathcal{L}_{\text{NTX}} = -\frac{1}{|\Pi|} \sum_{(i,p) \in \Pi} \log \frac{e^{d_{ip}/\tau_{\text{con}}}}{Z_{ip}}. \quad (9)$$

A.3 Single-Task Encoder Fine-Tuning

A.4 MARBERTv2 Regularization and Parallel Multi-Task Settings

We evaluated parallel multi-task supervision and three separate loss/regularization variants on the MARBERTv2 baseline: SupCon, focal loss, and R-Drop. The losses were not all applied simultaneously.

Item	Value
Input format	[CLS] tweet [SEP] target [SEP]
Classifier head	Dropout(0.1) $\rightarrow$ Linear(hidden, 3) on pooled [CLS]
Loss	Cross-entropy
Max sequence length	128
Batch size	32
Optimizer	AdamW
Learning rate	2×10^{-5}
Warmup ratio	0.1
Weight decay	0.01
Epochs / early stop	5 / patience 3 on dev F_{avg2}
Seed	42

Table 3: Single-task encoder fine-tuning settings.

The parallel MARBERTv2 formulation attaches a separate prediction head to the shared representation for each task and minimizes

$$\mathcal{L}_{\text{parallel}} = \mathcal{L}_{\text{stance}}^{\text{CE}} + 0.3 \mathcal{L}_{\text{sentiment}}^{\text{CE}} + 0.15 \mathcal{L}_{\text{sarcasm}}^{\text{CE}}. \quad (10)$$

Its contrastive variant adds $0.1 \mathcal{L}_{\text{SC}}$, where $\mathcal{L}_{\text{SC}}$ is the stance SupCon loss. The task coefficients (1.0, 0.3, 0.15) and the SupCon coefficient 0.1 were set heuristically rather than through systematic tuning.

These experiments were exploratory proof-of-concept evaluations, with loss coefficients set heuristically rather than through systematic tuning. We prioritized LLM-based approaches after they showed stronger development-set performance, leaving further hyperparameter tuning outside the scope of this study.

A.5 Stacked AraBERT Multi-Task System

The stacked system is derived from the publicly released AlexUNLP-BH architecture. An Arabic BERT encoder reads the tweet or the target–tweet pair and produces one pooled vector. Three classifier heads sit on that vector—stance (3), sentiment (3), and sarcasm (1, sigmoid). A fourth stacked head concatenates the three heads’ raw logits (3+3+1=7)

Component	Value
MTL heads	Sentiment (3-way, CE $\times 0.3$) and sarcasm (2-way, CE $\times 0.15$), each reading the same pooled representation as the stance head and supervised directly.
Contrastive loss	SupConLoss, temperature 0.07, weight $\times 0.1$; Equation (4).
Focal loss	Unweighted multi-class focal loss with $\gamma = 2.0$, replacing cross-entropy; Equation (5).
R-Drop	Two forward passes, their averaged CE, and $\alpha = 4.0$ times symmetric KL divergence; Equation (6).
Everything else	Same as Table 3.

Table 4: MARBERTv2 auxiliary-task and regularization settings.

and emits the final three-way stance logits.

The stance loss is applied to the stacked head only: the parallel stance head receives no direct supervision and is trained through the stacked head, while the sentiment and sarcasm heads are supervised directly.

For Track $r \in \{1, 2\}$, the complete objective is

$$\begin{aligned} \mathcal{L}_{\text{stack}}^{(r)} = & 0.7\mathcal{L}_{\text{stance}}^{(r)} \\ & + 0.2\mathcal{L}_{\text{sentiment}}^{(r)} \\ & + 0.3\mathcal{L}_{\text{sarcasm}}^{(r)} + 0.5\mathbf{1}\{r = 2\}\mathcal{L}_{\text{NTX}}. \end{aligned} \quad (11)$$

Track 1 uses inverse-frequency weighted CE for stance and sentiment and weighted binary cross-entropy (BCE) for sarcasm; Track 2 uses their unweighted forms. The Track 2 implementation uses supervised NTXentLoss (Musgrave et al., 2020); its implementation-specific denominator is defined in Section A.2.

For the stacked model, we retained the released AlexUNLP-BH implementation’s effective task coefficients (0.7, 0.2, 0.3) (Badran et al., 2024): the configured sarcasm coefficient 0.1 is multiplied by an additional factor of three in the code. For Track 2, we heuristically reduced the contrastive coefficient from 1.0 to 0.5. These inherited and heuristic settings do not establish optimal loss coefficients for either formulation.

A.6 ALLaM-7B Fine-Tuning Details

Item	Value
Base model	humain-ai/ALLaM-7B-Instruct-preview
Quantization	4-bit NF4, double quantization, bf16 compute
Adapter	LoRA $r = 16$, $\alpha = 32$, dropout 0.05; target modules q/k/v/o_proj, gate/up/down_proj; score head also trained
Objective	Unweighted three-way stance cross-entropy
Input	Arabic tweet–target classification template
Max sequence length	192
Batch size	$8 \times$ gradient accumulation 4 (effective 32)
Optimizer	paged AdamW (8-bit)
Learning rate	10^{-4}
Warmup ratio	0.06
Epochs / early stop	6 / patience 3 on dev $F_{\text{avg}2}$

Table 6: ALLaM-7B fine-tuning settings.

Input prompt. The ALLaM-based classifier used the following Arabic input template.

ALLaM-7B input prompt

ما هو موقف كاتب هذه التغريدة تجاه
 ؟<[TARGET]>
 غرد أحد المستخدمين على منصة X (تويتر)
 بالتغريدة التالية:
 .<[TWEET]>

A.7 RAG-ICL Models, Hyperparameters, and Environment

Models and generation. The heterogeneous ensemble used three hosted backends: OpenRouter’s google/gemini-3.1-pro-preview, NVIDIA NIM’s google/gemma-4-31b-it, and the OpenAI API’s gpt-4o. We used temperature 0.0 for all experiments to reduce sampling variation. The maximum output length was 16 tokens for direct-label prompts and 300 tokens for the RAG-ICL with Rationale-Augmented CoT experiment. The embedding model was Snowflake/snowflake-arctic-embed-1-v2.0 (Yu et al., 2025), loaded through Hugging Face sentence-transformers.

Evaluation coverage. Table 7 summarizes the combinations of method, track, and LLM backend evaluated after the preliminary zero-shot model selection reported in Table 10. Gemma 4 31B was held fixed for the controlled zero-shot, Vanilla, Adapted, and rationale-augmented comparisons;

Item	Track 1 (related target / seen domain)	Track 2 (unseen targets)
Encoder	aubmindlab/bert-base-arabertv02-twitter (BERT-base, ~135M)	same
Dropout	0.1	same
Model input	Exp 1: tweet only; Exp 2: target + tweet	same
Heads (parallel)	stance (3), sentiment (3), sarcasm (1, sigmoid)	same
Stacked stance head	concat of three heads' raw logits (7) $\rightarrow$ linear $\rightarrow$ 3	same
Task weights	[0.7, 0.2, 0.1], with sarcasm additionally $\times 3$ in code; effective [0.7, 0.2, 0.3]	same
Stance / sentiment loss	CE with inverse-frequency class weights	unweighted CE (None is near-absent)
Sarcasm loss	weighted BCE	unweighted BCE
Contrastive loss	none	supervised NTXentLoss on stance labels, $\lambda = 0.5$, temperature 0.07; training only; Equation (9)
Max sequence length	128 (longest-first truncation for pairs)	same
Batch size	32 train; eval/test loaders use 1	same
Optimizer	AdamW, lr 2×10^{-5} , linear warmup (one-third epoch) then linear decay	same
Max epochs / early stop	60 / patience 7 on stance validation loss	same
Checkpoint selection	minimum stance validation loss	same
Seed	42	42
Label map	None=0, Favor=1, Against=2	same

Table 5: Stacked AraBERT multi-task settings.

the final data-centric configuration was then evaluated with all three ensemble backends on both tracks.

For standard RAG-ICL, the retriever first obtains $K_{\max} \times 10$ candidates before applying label balancing and the similarity threshold. The RAG-ICL with Rationale-Augmented CoT experiment uses five demonstrations per independently constructed branch.

The data-centric retriever applies a single global fixed cosine-similarity threshold of $\tau = 0.55$ to all posts, retaining at most $K_{\max} = 2$ examples per label. This threshold was chosen heuristically to filter out low-similarity candidates and reduce the risk of irrelevant demonstrations entering the prompt; systematic threshold tuning is left for future work.

Software and execution environment. The RAG experiments were run locally on macOS, with inference performed through hosted APIs. The implementation used Python 3, pandas for data manipulation and ensemble voting, sentence-transformers for dense embeddings, and scikit-learn for exact cosine-similarity retrieval. The official openai Python client provided a common interface for OpenAI-compatible endpoints by changing the base URL.

Hashtag glosses and negation guidance. The adapted and data-centric prompts contained neutral glosses for two task-relevant hashtags:

#نن_تقودي (“you (*female*) will not drive”) and #المراه_السعوديه_تسوق (“Saudi woman drives”).

The adapted variant also includes extracted hashtags in the embedding query and explicitly highlights Arabic negation markers such as لا، ما، نن، مو and their scope.

Target mappings and definitions. Table 9 lists the manually curated cross-target retrieval mapping and the Arabic target definitions used by the data-centric pipeline. These resources were derived through semi-automatic inspection of the target inventory and manually validated using domain knowledge. The retrieval mapping affects demonstration selection only, whereas the definitions are injected into the LLM prompt. Track 1 maps *Women Driving* to the semantically related training target *Women Empowerment*, retaining the original target in the prompt. Track 2 uses no explicit cross-target mapping and retrieves from the provided training demonstrations; target definitions disambiguate short English labels such as *E-Cars* and *Trimester*. The mappings, hashtag glosses, and target definitions were fixed without access to gold test labels.

B Development Results

Table 10 shows development-set results for the supervised systems and the preliminary zero-shot LLM comparison. Gemma 4 31B is the strongest

Method	Track(s)	Gemma 4 31B	Gemini 3.1 Pro	GPT-4o
Zero-shot baseline	T1, T2	✓	—	—
<i>RAG-ICL variants</i>				
↪ Vanilla	T1	✓	—	—
↪ Adapted	T1	✓	—	—
↪ Data-centric	T1, T2	✓	✓	✓
↪ Rationale-augmented CoT	T1	✓	—	—
Heterogeneous ensemble	T1, T2	✓	✓	✓

Table 7: RAG-ICL evaluation matrix by method, track, and LLM backend. A check mark denotes an evaluated combination; — denotes a combination that was not evaluated.

Hyperparameter	Value
Per-label retrieval cap, $K_{\max}$	2
Maximum demonstrations per prompt	6
RAG-ICL+CoT demonstrations per branch	5
Cosine threshold, τ	0.55
Candidate multiplier	10

Table 8: RAG-ICL retrieval hyperparameters.

zero-shot backbone, while ALLaM-7B is the strongest supervised model on both tracks.

Regularization and auxiliary-task supervision produced modest, model-dependent changes across the encoder experiments, with no single encoder configuration leading on both tracks.

On Track 1, AraBERT-large-Twitter scores 0.8139, slightly below AraBERT-base-Twitter at 0.8145, although both outperform XLM-R-large at 0.7790 and mBERT at 0.6907. Track 2 shows the same advantage for AraBERT-base-Twitter at 0.8091 over XLM-R-large at 0.7739 and mBERT at 0.5541. Within this comparison, dialect-relevant pretraining data appears more important than model size.

Focal loss on MARBERTv2 reaches 0.8263 on Track 1 and 0.8208 on Track 2, compared with 0.8198 and 0.8150 for parallel MTL with sentiment and sarcasm heads.

Stacked AraBERT is the strongest encoder on the Track 1 development set at 0.8286, but falls to 0.6707 on the test set. This is the largest Track 1 development-to-test decrease among the encoder systems, although it is smaller than the 0.1812 decrease observed for ALLaM-7B.

C Detailed Error Analysis

This analysis concerns the Track 1 development set only. We reviewed 30 instances for which the heterogeneous ensemble (our best method) disagreed with the gold label, comprising 10 examples from each Track 1 target: Digital Transformation, COVID-19

Track	Target	Retrieval mapping / definition
<i>Cross-target retrieval mapping</i>		
1	Women Driving	→ Women Empowerment
<i>Injected target definitions</i>		
1	Women Driving	قيادة المرأة في السعودية
1	Women Empowerment	تمكين المرأة في السعودية
2	Trimester	نظام الفصول الدراسية الثلاثة
2	E-Cars	السيارات الكهربائية
2	COVID-19 Vaccine	لقاح كورونا
2	Digital Transformation	التحول الرقمي

Table 9: Cross-target retrieval mapping and injected target definitions used in the data-centric RAG-ICL pipeline. The Track 1 mapping associates the test target *Women Driving* with the semantically related training target *Women Empowerment* for demonstration retrieval only. Target definitions are injected into the LLM prompt to disambiguate short English target labels. No gold test labels were used in constructing either resource.

Vaccine, and Women Empowerment. The counts characterize this selected sample rather than error rates on the full development set.

Twenty of the 30 cases involve a boundary between *None* and either *Favor* or *Against*; the remaining 10 exchange the two polar labels. The three backends produce the same non-gold label in 17 cases, a two-to-one majority in 10, and three different labels in three. In 21 cases no backend produces the gold label, whereas in the other nine exactly one backend does. Of these nine, six have a gold-matching minority prediction and three are three-way ties. The remaining four two-to-one splits contain no gold-matching prediction. This separates disagreements that vote selection could resolve from those in which the gold label is absent from all backend outputs.

Table 11 presents seven representative cases se-

System	T1 (dev)	T2 (dev)
SINGLE-TASK ENCODERS		
MARBERTv2	0.8052	0.8079
AraBERT-base-Twitter	0.8145	0.8091
AraBERT-large-Twitter	0.8139	—
QARiB	0.8099	0.7777
ARBERTv2	0.7893	0.7880
XLM-R-large	0.7790	0.7739
AraBERT-large-MSA	0.7653	—
CAMeLBERT-DA	0.7365	0.7461
mBERT	0.6907	0.5541
MARBERTv2 REGULARIZATION		
↪ Contrastive loss	0.8226	0.8103
↪ R-Drop	0.8242	0.8167
↪ Focal loss	0.8263	0.8208
ARA-BERT-BASE-TWITTER JOINT-LOSS VARIANT		
↪ SupCon + Center (auto-weighted)	0.8055	—
MULTI-TASK ENCODERS		
MARBERTv2, parallel MTL	0.8198	0.8150
↪ Contrastive stance loss	0.8213	0.8162
Stacked AraBERT MTL, tweet only	0.8286	0.7874
↪ Target + tweet input	0.8128	0.8048
LLM FINE-TUNING		
ALLaM-7B w/ QLoRA	<u>0.8606</u>	<u>0.8533</u>
ZERO-SHOT LLMs		
Fanar	0.7827	0.7134
DeepSeek V4 Flash	0.8438	0.6224
Mistral Small 4	0.8108	0.8439
Gemma 4 31B	0.8911	0.8773

Table 10: Development set results under F_{avg2} for each method family. Bold and underlining mark the best and second-best scores.

lected to cover all three Track 1 targets, the recurring interpretive phenomena, and the three observed vote structures. The gold label remains the benchmark reference throughout. We distinguish the observed predictions from our assessment of the textual evidence, including ambiguity and a likely annotation inconsistency in Case F. URL and mention placeholders are omitted from the displayed text, and emoji or omitted hashtag metadata are described in brackets when relevant. An English translation follows each Arabic post.

Implicit and reported stance. Cases A and E both receive an ensemble prediction of *None* against a gold *Favor* label, but differ in their evidence for author stance. Case A reports digital-transformation initiatives without explicit evaluation. Reporting the event alone does not establish endorsement, making the unanimous *None* prediction plausible from the displayed wording. In Case E, للأسف (“unfortunately”) expresses regret about the lack of vaccination centers, and the writer is concerned that

others cannot obtain a dose. This provides stronger evidence for *Favor*, despite the negative sentiment toward access conditions. Gemini matches that label, while Gemma and GPT-4o choose *None*. The distinction is between merely mentioning a target and expressing a desire for access to it, not simply between explicit and implicit evaluation.

Sarcasm and pragmatic intent. Case D turns on whether the microchip claim is endorsed or ridiculed. An endorsed claim would support an anti-vaccine reading, whereas mockery of the conspiracy would not itself establish opposition to vaccination. The smiling emoji is compatible with a joking reading but does not settle the direction of the ridicule. Gemma matches the gold *Against* label, while Gemini and GPT-4o choose *None*. In contrast, Case C anchors its sarcastic praise in a concrete deterioration from a fifteen-minute transaction to a week-long delay. Its criticism of the service is therefore clearer. These cases distinguish identifying sarcasm from determining what it evaluates.

Scope of criticism. Case C criticizes the “Sahel” platform and explicitly links its delays to electronic transformation, supporting the unanimous *Against* prediction. The gold *Favor* label might reflect a distinction between the underlying goal and its implementation, but dissatisfaction with a service does not by itself establish support for that goal. Case G instead criticizes women said to have exploited empowerment and calls for stopping their conduct. Extending this criticism to empowerment itself supports an *Against* reading, whereas restricting it to those beneficiaries leaves the broader stance less explicit. Gemini, Gemma, and GPT-4o predict *None*, *Favor*, and *Against*, respectively. Both cases concern how far local criticism extends to the evaluated target, rather than whether the language is negative.

Target versus opponent and annotation review. Case F provides stronger evidence for the ensemble label than for the gold label. The phrase **عاجبك تمكين المرأة** (“you dislike women’s empowerment”) attributes opposition to the addressee, not the writer. The writer presents empowerment as enabling women to reject mistreatment and mocks the addressee’s loss of privileges. These cues support *Favor*, the unanimous prediction, rather than the gold *Against*. We therefore regard this case as a likely annotation inconsistency and a candidate for

independent re-adjudication. The judgment rests on the text, not model agreement. The prediction is also consistent with the target-versus-actor instruction in our data-centric prompt (Section D). Unlike Case G, the post explicitly identifies an opponent of empowerment, making the relation between the criticized actor and the target clearer.

Mixed or concessive positions. Case B begins with explicit praise but uses **ولكن** (“but”) to introduce a concern about catastrophic data loss. The unanimous *Favor* prediction is consistent with treating that concern as a qualification of support. The gold *Against* label is consistent with treating it as overriding the initial praise. The decision boundary is whether the concern qualifies or reverses the positive position, rather than whether both polarities are present. Our data-centric prompt treats conditional support as generally *Favor*, making this boundary relevant to the scope of that instruction without establishing that it caused the prediction.

Implications for ensembling. Cases D and E contain a gold-matching minority prediction that majority voting suppresses. In Case G, the tie-breaking rule selects Gemini’s *None* despite GPT-4o predicting the gold *Against*. Across the sample, an oracle selecting among existing backend predictions could resolve at most nine of the 30 disagreements. This is an upper bound within the selected sample, not an expected gain on the full development set. In the other 21 cases, matching the gold label requires a changed backend prediction rather than a different selection among existing outputs. Case F also shows why unanimous disagreement should not automatically be treated as shared interpretive failure. Review based only on low consensus would miss this likely annotation inconsistency. The examples therefore motivate separate attention to vote selection, stance interpretation, and annotation review rather than a single remedy for all benchmark disagreements.

D Prompt Templates for ICL-based Methods

This appendix gives the prompt content used by the frozen-LLM in-context-learning methods and their zero-shot reference condition. It covers the zero-shot baseline, Vanilla RAG-ICL, Adapted RAG-ICL, data-centric RAG-ICL, and the rationale-augmented variant. Runtime substitutions are shown as colored bracketed placeholders.

D.1 Zero-Shot Baseline

System prompt

You are an expert Arabic NLP annotator specializing in stance detection.

Your task: Given an Arabic social media post and a specific target topic, classify the writer’s stance toward that target.

Classification rules:

- "Favor": The writer expresses support, agreement, or a positive position toward the target.
- "Against": The writer expresses opposition, disagreement, or a negative position toward the target.
- "None": The writer’s stance toward the target is unclear, neutral, or the post is irrelevant to the target.

Important notes:

- Stance is about the writer’s position toward the TARGET, not general sentiment.
- A post can have negative sentiment but still favor the target (or vice versa).
- Sarcastic posts may use positive words to express opposition-consider the true intent.
- The target may not be explicitly mentioned in the text.

You must respond with EXACTLY one word: Favor, Against, or None.

User prompt

Target: [TEST_TARGET]

Post: [TEST_TEXT]

What is the writer’s stance toward the target? Respond with exactly one word: Favor, Against, or None.

D.2 Vanilla RAG-ICL

The system prompt is identical to the zero-shot baseline.

User prompt

Here are some labelled examples to help you:

Target: [DEMO_TARGET]

Post: [DEMO_TEXT]

What is the writer’s stance toward the target? Respond with exactly one word: Favor, Against, or None.

Stance: [DEMO_LABEL]

[REPEATED FOR K EXAMPLES PER LABEL]

Target: [TEST_TARGET]

Post: [TEST_TEXT]

What is the writer's stance toward the target? Respond with exactly one word: Favor, Against, or None.

D.3 Adapted RAG-ICL

The adapted system prompt appends the following item to the baseline's *Important notes* section.

Additional system instruction

- Pay close attention to negation words (e.g., لا، ما، مو، لن). Check carefully what part of the sentence the negation applies to: it may negate the whole stance, only part of a clause, or be reversed later in the same post (e.g., a concessive "except when..." that flips the practical meaning).

User prompt

Here are some labelled examples to help you:

Target: [DEMO_TARGET]

Post: [DEMO_TEXT]

What is the writer's stance toward the target? Respond with exactly one word: Favor, Against, or None.

Stance: [DEMO_LABEL]

[REPEATED FOR K EXAMPLES PER LABEL]

Target: [TEST_TARGET]

Post: [TEST_TEXT]

Stance hashtags: [HASHTAG_GLOSS]

What is the writer's stance toward the target? Respond with exactly one word: Favor, Against, or None.

D.4 Data-Centric RAG-ICL

System prompt

You are an expert Arabic NLP annotator specializing in stance detection.

Your task: Given an Arabic social media post and a specific target topic, classify the writer's stance toward that target.

Classification rules:

- "Favor": The writer expresses support, agreement, or a positive position toward the target.

- "Against": The writer expresses opposition, disagreement, or a negative position toward the target.

- "None": The writer's stance toward the target is unclear, neutral, or the post is irrelevant to the target.

Important notes:

- Stance is about the writer's position toward the TARGET, not general sentiment.

- Target vs. Actor: If the writer attacks the opponents of the target, their stance toward the target is "Favor".

- Conditional Support: If the writer supports the target but adds a condition (e.g., "I agree, provided that..."), it is generally still "Favor".

- Religious/Cultural Invocations: Pay attention to implicit prayers. الله يستر (God protect us) usually implies fear of change ("Against"). الله يعز الحكومة (God strengthen the government) usually supports official decrees ("Favor"). حسبنا الله (God is sufficient) usually expresses opposition. These are indications only; context is important and each text must be analyzed as a whole.

- Hashtag hijacking: Users may use an opposing hashtag to mock or critique the opposing side. Base the stance on the overarching meaning of the post, not only on the literal translation of the hashtag.

- The text may contain Saudi dialect or typographical errors. Focus on the underlying meaning.

- The target may not be explicitly mentioned in the text.

[TARGET_DEFINITION_INJECTED_HERE]

You must respond with EXACTLY one word (provide only one, verbatim): Favor, Against, or None.

The data-centric user prompt is identical to the adapted RAG-ICL format and retains the runtime [HASHTAG_GLOSS] field.

D.5 RAG-ICL with Rationale-Augmented CoT

System prompt

You are an expert Arabic NLP annotator specializing in stance detection.

Your task: Given an Arabic social media post and a specific target topic, classify the writer's stance toward that target.

Classification rules:

- "Favor": The writer expresses support, agreement, or a positive position toward the target.
- "Against": The writer expresses opposition, disagreement, or a negative position toward the target.
- "None": The writer's stance toward the target is unclear, neutral, or the post is irrelevant to the target.

Important notes:

- Stance is about the writer's position toward the TARGET, not general sentiment.
- Sarcastic posts may use positive words to express opposition-consider the true intent.
- The target may not be explicitly mentioned in the text.

CRITICAL INSTRUCTION: Before you output your final stance, you MUST explain your reasoning in <thought></thought> tags. Analyze the sentiment, implicit meaning, and relevance to the target.

After closing the thought tags, output the final stance label exactly as:

Stance: Favor

OR

Stance: Against

OR

Stance: None

User prompt

Here are some labelled examples to help you understand the task:

Target: [DEMO_TARGET]

Post: [DEMO_TEXT]

Stance: [DEMO_LABEL]

[REPEATED FOR K EXAMPLES PER LABEL]

Target: [TEST_TARGET]

Post: [TEST_TEXT]

Explain your reasoning in <thought> tags, then output the stance label.

Case and phenomenon	Target, Arabic post, and English translation	Gold	Ensemble	Backend votes
A. Reported or implicit stance	Target: Digital Transformation الحكومة المصرية تستعرض ٧ مبادرات للتحويل الرقمي [hashtags and URLs omitted] <i>The Egyptian government presents seven digital-transformation initiatives.</i>	Favor	None	None / None / None
B. Mixed or concessive stance	Target: Digital Transformation فعلاً والله التحويل الإلكتروني حلو ومريح ولكن ماذا سيحدث إذا فقدنا كل البيانات؟ <i>Indeed, digital transformation is nice and convenient, but what would happen if we lost all the data?</i>	Against	Favor	Favor / Favor / Favor
C. Concept versus implementation	Target: Digital Transformation انا اقول غيرو اسم سهل الى صعب وايد ابرك في السابق كنا نروح الوزارة ربع ساعة تكون مجدد رخصة الاستيراد مع التحويل الإلكتروني المعامله صار لها اسبوع عند الموظف خووووش تحول الكتروني ومسمى اسم منصته سهل انتوا وين والتطور وين <i>I suggest changing the name from “Sahel” (Easy) to “Difficult”; that would be much better. Previously, we could renew an import license at the ministry in fifteen minutes. With electronic transformation, the transaction has been with the employee for a week. What a wonderful electronic transformation, yet you call the platform “Easy.” Where are you, and where is progress?</i>	Favor	Against	Against / Against / Against
D. Sarcasm and conspiracy reference	Target: COVID-19 Vaccine هالبيانات سربتها الشريحة اللي غرسوها فيك مع [smiling emoji] <i>Those data were leaked by the microchip they implanted in you with the COVID-19 vaccine.</i>	Against	None	None / Against / None
E. Practical barrier versus target stance	Target: COVID-19 Vaccine [Hashtag: Have you taken a vaccine dose or not yet?] للأسف ماحولنا مراكز تطعيم اقرب شي بعد ٠٥١ كيلو وان اخذته انا غيري كثير من اهل الديره مايقدرن ياخذونه [shrugging emoji] <i>Unfortunately, there are no vaccination centers near us; the closest is more than 150 kilometers away. Even if I receive it, many other people in the village cannot.</i>	Favor	None	Favor / None / None
F. Target versus opponent	Target: Women Empowerment صح عليك المرأة الممكنه المستقبله الموظفة وش يخليها تصبر على العنف والضرب والتعدد والخيانة والولاية ووو ؟؟ مافيه سبب عشان كذا مو عاجبك تمكين المرأة لأنه يسحب منك امتيازتك ويرجعك لحجمك الحقيقي وصدقني راح تبكي اكثر واكثر [mention, personal name, and laughing emoji omitted] <i>Exactly. Why would an empowered, employed woman tolerate violence, beating, polygamy, betrayal, guardianship, and so on? There is no reason. That is why you dislike women’s empowerment: it takes away your privileges and returns you to your true size. Believe me, you will cry more and more in the future.</i>	Against	Favor	Favor / Favor / Favor
G. Criticism of beneficiaries versus target	Target: Women Empowerment للأسف استغل تمكين الدولة للمرأة من بعض ضعيفات النفوس والباحثات عن الشهرة والمال واصبحن يشوهن صورة المرأة السعودية التي كانت على طول الازمان [hashtag omitted] <i>Unfortunately, some ill-intentioned women and seekers of fame and money have exploited the state’s empowerment of women and have begun to tarnish the image of Saudi women, long regarded as an example of honor; they therefore must be stopped from this misconduct.</i>	Against	None	None / Favor / Against

Table 11: Representative Track 1 dev-set disagreements. Backend votes are ordered as Gemini / Gemma / GPT-4o.