

Saigum at StanceEval-2026: Evaluating LLM-ensemble models and Retrieval Augmented Classification for Arabic Stance Detection

Sairam Narendra Babu

IIIT Hyderabad

sairam.babu@research.iiit.ac.in

Abstract

Arabic stance depends on quotation, sarcasm, negation, hashtags, and target scope, which surface polarity may miss. For StanceEval-2026, we study Codex and Gemini prompting through two submission systems: a replicated contextual Codex predictor and a gated judge that retrieves labeled examples using seven AI-generated discourse cues. On 619 Dev examples, cue retrieval changes Fav2 from 0.8580 to 0.8613 for GPT-5.6-luna, 0.8767 to 0.8844 for Gemini 3.5 Flash, and 0.8646 to 0.8554 for Flash-Lite. Random and semantic controls show no consistent retrieval advantage. On the 352-row hidden set, the best retained Codex member scores 0.8685 Fav2 and currently ranks fourth, while gated retrieval scores 0.8604. Codex and OPUS-MT translation pipelines score 0.7848 and 0.7304, and a synthetic-data classifier scores 0.7334. These results support replication and retrieval as auditable system components, not as general performance gains.

1 Introduction

Stance detection assigns Favor, Against, or None toward a supplied target (Küçük and Fazli, 2020). Unlike sentiment detection, stance detection must determine whether an opinion concerns the target and whether the author endorses it, merely comments on it, or expresses it sarcastically. Arabic social media adds dialect, implicit arguments, and rhetorical questions. We follow the Mawqif task and data papers (Alturayef et al., 2022; Albalawi et al., 2026b) and the StanceEval-2026 overview (Albalawi et al., 2026a).

Our contributions are threefold. First, we construct ai-generated categorical descriptions of seven stance-relevant discourse phenomena and use weighted categorical agreement to retrieve tweet-stance pairs without lexical or target overlap. Second, we combine this retrieval mechanism with ensemble-variability and ensemble-disagreement

signals to filter and choose to revisit a subset of uncertain examples. Third, through repeated comparisons against zero-shot, random, and semantic-secondary selection, we show that cue retrieval-augmented few-shot classification produces small model-dependent changes rather than a consistent performance advantage, while providing auditable evidence for individual decisions.

2 Related work

Arabic stance systems commonly fine-tune AraBERT, ARBERT, or MARBERT encoders (Antoun et al., 2020; Abdul-Mageed et al., 2021), including in earlier StanceEval work (Alturayef et al., 2024). We instead study the LLM inference interface, drawing on in-context learning and RAG (Brown et al., 2020; Lewis et al., 2020). Our retriever uses generated categorical discourse cues as weak signals (Ratner et al., 2017), rather than dense topical similarity alone.

3 Task and data

Track 1 Train has 3,502 tweets (1,021 Against, 2,148 Favor, and 333 None), Dev has 619, and the hidden evaluator has 352. Train supplies labels and retrieval examples; Dev supports controlled comparison; hidden labels are used only by the submission evaluator. Fav2 averages Favor and Against F1; we also report Fav3, accuracy, and class F1 on Dev. Weak cues describe an existing tweet’s pragmatic structure without seeing its gold stance, whereas synthetic examples are newly generated tweets with requested labels. Neither is human annotation.

4 Systems

Figure 1 summarizes the two main submission paths.




Figure 1: The two main submission systems. System B starts from the ten-member majority record produced by System A and leaves non-candidate rows unchanged.

4.1 System A: contextual Codex family

The contextual predictor processes batches of ten target and Arabic-tweet pairs. Its Arabic prompt defines the labels and asks the model to resolve dialect, hashtags, negation, quotation, sarcasm, and implicit arguments before returning structured JSON with a label and confidence. We ran ten replicas with the same prompt and stable IDs. Majority ties are broken by mean reported confidence; 269 of 352 hidden rows are unanimous.

Our top submission, `finale.zip`, contains replica 3 rather than the ten-replica majority-vote output. We selected replica 3 using both an external LLM’s assessment of the replicas’ reported confidence records and the hidden-evaluator scores obtained after evaluating all ten replicas. Its 0.8685 Favg2 score is therefore a leaderboard-selected best-member result, rather than an unbiased estimate of the expected performance of an individual replica or the majority-vote ensemble. The performance of the majority-vote ensemble of the hidden-set replicas is mentioned in Table 2.

4.2 AI-generated tweet annotations

For each tweet, an annotator receives only its ID, target, and Arabic text. It produces seven categorical fields:

```

target_stance_presence,
opinion_scope, speech_act,
hashtag_direction,
sarcasm_direction,

```

```

author_attribution, and
contrast_direction.

```

We discard rationales that name a stance. Train, Dev, and evaluation cues are generated separately and joined by ID.

For query q and Train example x , weighted categorical agreement is

$$s(q, x) = \frac{\sum_j w_j \mathbf{1}[q_j = x_j]}{\sum_j w_j}. \quad (1)$$

Weights are 1.0 for target-stance presence; 0.8 for opinion scope and speech act; and 1.4 for hashtag direction, sarcasm direction, author attribution, and contrast direction. Target identity, raw-text overlap, and gold stance are excluded from this score. The gold label is attached only after retrieval so the selected tweet can be shown as an in-context example.

4.3 System B: gated retrieval judge

The retrieval system starts from the ten-member majority record and selects a row when its winning vote share is below 0.80 or it disagrees with a retained binary baseline. The gate selects 62 of 352 rows, each paired with its four highest-scoring Train examples. The judge first predicts from the query, then considers the seed label and demonstrations; it changes 26 candidate labels and leaves all other rows unchanged.

4.4 Controlled variants and additional pipelines

On Dev, cue retrieval supplies three examples to every query. We compare it with random examples and with a embedding-based tie-break using cosine similarity over L2-normalized, 3,072-dimensional `gemini-embedding-001` tweet vectors. Two translation pipelines classify English text produced by Codex or `Helsinki-NLP/opus-mt-tc-big-ar-en`. For synthetic supervision, we generated 1,500 ordinary and 1,500 difficult Women Driving tweets with Gemini 3.5 Flash; the mixed classifier uses Mawqif Train, 756 ArabicStanceX examples (Alkhathlan et al., 2025), and 1,500 difficult synthetic examples. Direct controls encode the seven cues in XGBoost, alone or with target, sentiment, and sarcasm metadata.

5 Experimental setup and reproducibility

The initial Dev comparison evaluates GPT-5.6-luna, Gemini 3.5 Flash, and Gemini 3.5 Flash-Lite

in zero-shot and cue-retrieved three-shot conditions. The stricter selection comparison evaluates GPT-5.6-sol at medium reasoning effort and both Gemini models with either random examples or cue-primary, semantic-secondary examples. Every LLM condition has three full-Dev runs; reported deviations are sample standard deviations.

For the repeated Gemini 3.5 classification calls, temperature was 0 and `maxOutputTokens` was not set, leaving the limit at the provider default. The GPT-5.6 experiments used `codex exec`; the scripts did not expose or pass either temperature or a maximum output-token limit, so the Codex defaults were used. Table 6 reports the corresponding settings for the auxiliary proprietary-model calls.

The code repository is available at https://github.com/Saigum/llm_arabic_tweet_eval_benchmark/tree/main.

6 Results

6.1 Controlled Dev results

Table 1 gives the main controlled results.¹ *Cue retrieval improves two of the three models, but the effect does not hold across all models.* This does not support a general claim that cue-based retrieval helps across models.

Random and semantic retrieval also produce mixed results, with neither consistently outperforming the other. The seven AI-generated cues place many tweets in the same archetype, so random and cue-based retrieval can supply similar kinds of examples.

The cues are weak direct predictors: cue-only XGBoost falls below the TF-IDF baseline, while adding cues to the provided metadata improves over metadata alone. Cue-based label voting without an LLM is also weak, particularly for None. Cue profiles can organize evidence, but they do not determine the stance direction.

6.2 Hidden submission results

Table 2 reports only scores retained in the submission ledger. *The best contextual Codex member currently ranks fourth; gated retrieval remains below it, and both English pipelines are worse than direct Arabic prompting.* Strong zero-shot Dev performance does not guarantee comparable performance on the evaluation split. Machine translation

can remove dialectal, sarcastic, or hashtag cues before the stance model sees the tweet, which may reduce performance.

6.3 Synthetic supervision

The hard synthetic set appears strong under random splitting: LinearSVC reaches 0.9631 Favg2 on synthetic-only data and 0.9050 after mixing synthetic and real Women Driving examples. Holding out complete generation families reduces the best score to 0.7660. *The corresponding hidden submission remains well below the contextual LLM systems.* Generated examples were useful for constructing stress cases, but they did not improve hidden transfer. For this task, synthetic tweet generation does not appear to be an effective training strategy because the generated data does not adequately reproduce dialectal variation, sarcasm, or natural Arabic tweet style.

7 Analysis

7.1 Concrete error cases

Table 3 shows three Dev examples using short Arabic extracts. The first illustrates a case where retrieval helps with a qualified position. The other two show anchoring and target-scope failures. IDs refer to the released Dev data so the full Arabic text can be inspected.

Across all unanimous Gemini Flash comparisons, retrieval corrects seven cases that zero-shot misses and harms four cases that zero-shot gets right. Both conditions make the same wrong prediction on 83 further examples. These counts do not by themselves measure Favg2 because class frequencies and error types differ, but they show that the mean gain is concentrated in a small set of decisions.

7.2 Why retrieval effects vary

The seven cues compress nuanced text into exact categories. They can preserve attribution, contrast, and sarcasm patterns while losing the content that determines direction. Equal-score pools are also large. Cue-only retrieval selects 492 distinct Train tweets and 190 distinct triples for 619 Dev queries; one triple is reused 69 times. Semantic tie-breaking raises this to 1,221 distinct tweets and 603 triples, yet the primary metric does not improve consistently.

¹Temperatures are shown in parentheses; “def.” denotes the default temperature used by the Codex runner. Exact max-output-token limits, top-*p*, reasoning effort, and unconfigured provider defaults are documented in Appendix Table 6.

Model (temperature)	Prompt condition	Favg2	Favg3	Accuracy
GPT-5.6-luna (def.)	Zero-shot	0.8580 $\pm$.0082	0.7276 $\pm$.0161	0.8218 $\pm$.0104
GPT-5.6-luna (def.)	Cue-retrieved three-shot	0.8613 $\pm$.0057	0.7373 $\pm$.0086	0.8244 $\pm$.0052
Gemini 3.5 Flash (0)	Zero-shot	0.8767 $\pm$.0031	0.7363 $\pm$.0036	0.8352 $\pm$.0028
Gemini 3.5 Flash (0)	Cue-retrieved three-shot	0.8844 $\pm$.0018	0.7403 $\pm$.0023	0.8417 $\pm$.0016
Gemini 3.5 Flash (0)	Random three-shot	0.8829 $\pm$.0008	0.7553 $\pm$.0035	0.8465 $\pm$.0016
Gemini 3.5 Flash (0)	Semantic tie-break three-shot	0.8789 $\pm$.0025	0.7613 $\pm$.0061	0.8471 $\pm$.0034
Gemini 3.5 Flash-Lite (0)	Zero-shot	0.8646 $\pm$.0047	0.7434 $\pm$.0038	0.8293 $\pm$.0041
Gemini 3.5 Flash-Lite (0)	Cue-retrieved three-shot	0.8554 $\pm$.0050	0.7232 $\pm$.0129	0.8180 $\pm$.0073
Gemini 3.5 Flash-Lite (0)	Random three-shot	0.8557 $\pm$.0096	0.7402 $\pm$.0051	0.8277 $\pm$.0065
Gemini 3.5 Flash-Lite (0)	Semantic tie-break three-shot	0.8657 $\pm$.0020	0.7374 $\pm$.0063	0.8314 $\pm$.0034
GPT-5.6-sol (def.)	Random three-shot	0.8631 $\pm$.0044	0.7234 $\pm$.0052	0.8234 $\pm$.0041
GPT-5.6-sol (def.)	Semantic tie-break three-shot	0.8585 $\pm$.0027	0.7297 $\pm$.0032	0.8244 $\pm$.0025

Table 1: Mean and sample standard deviation over three runs on all 619 Dev examples. Bold marks the best value within the Gemini 3.5 Flash block, not a claim of statistical significance.

Hidden-set submission or group	Favg2
Contextual Codex, 10 replicas (mean)	0.8525 $\pm$.0106
Contextual Codex, 10-replica majority	0.8538
Contextual Codex, best member	0.8685
Gated weak-cue retrieval	0.8604
Gemini cue retrieval, 3 runs (mean)	0.8418 $\pm$.0039
Codex-translated English	0.7848
OPUS-MT English	0.7304
Synthetic-data LinearSVC	0.7334

Table 2: Hidden evaluator results retained in the submission ledger.

7.3 None and target scope

None remains the hardest class. It includes neutral reports, quoted positions without author adoption, criticism of a related actor, and claims about an implementation rather than the target. The official Favg2 metric omits None, so we report Favg3 and class-level behavior when comparing prompting methods. For example, semantic tie-breaking raises None F1 for GPT-5.6-sol while lowering Against F1, which leaves Favg2 below the random control.

8 Limitations

Dev was used repeatedly during development, so Dev performance is not an unbiased estimate of future transfer. The hidden evaluator was queried for many submissions, and the best Codex member was selected from ten observed scores. The retrieval weights were chosen manually and were not learned on a separate validation set. Cue annotations have not been independently checked by human annotators, and the Train and hidden cue files were produced by different historical pipelines with partially different category vocabularies.

The semantic embeddings omit the supplied tar-

get and use the same retrieval task setting for Train and Dev. Random demonstrations vary across repeats, whereas retrieved neighbors are fixed, so their standard deviations capture different sources of variation.

9 Conclusion

We compared a replicated contextual Codex predictor with a gated pragmatic-cue retrieval system for StanceEval-2026. The best retained Codex member scores 0.8685 Favg2 on the hidden evaluator, currently ranking fourth in the StanceEval2026-SeenTargets Track; gated retrieval scores 0.8604. Cue-retrieved demonstrations improve two of three models in the initial Dev comparison, but random and semantic controls show no consistent selection advantage. Translating into English before classification is also worse than direct Arabic prompting in both retained runs. Synthetic data produces high random-split scores and poor hidden transfer. The useful systems lesson is narrower: repeated LLM predictions and logged retrieval evidence expose uncertainty and individual failure cases, but neither component should be described as an inherent performance gain.

Acknowledgments

We thank the StanceEval-2026 organizers for providing the task data and evaluation infrastructure.

References

Muhammad Abdul-Mageed, AbdelRahim Elmadany, and El Moatez Billah Nagoudi. 2021. [ARBERT & MARBERT: Deep bidirectional transformers for Arabic](#). In *Proceedings of the 59th Annual Meeting of the*

Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers), pages 7088–7105, Online. Association for Computational Linguistics.

Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Saad Ezzini, Maram Kurdi, Asma Yamani, Ahmed Ashraf, Maged Al-Shaibani, and Nora Alturayef. 2026a. Stanceeval-2026: The second stance detection shared task. In *Proceedings of the Fourth Arabic Natural Language Processing Conference (ArabicNLP 2026)*, Budapest, Hungary. Association for Computational Linguistics.

Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Maram Kurdi, Saad Ezzini, Asma Yamani, and Ahmed Ashraf. 2026b. [Mawqif-xt: An arabic benchmark dataset for cross-target stance detection](#). Preprint, arXiv:2608.09539.

Ali Alkhathlan, Faris Alahmadi, Faris Kateb, and Hend Al-Khalifa. 2025. [Constructing and evaluating arabicstancex: a social media dataset for arabic stance detection](#). *Frontiers in Artificial Intelligence*, 8:1615800.

Nora Alturayef, Hamzah Luqman, Zaid Alyafeai, and Asma Yamani. 2024. [Stanceeval 2024: The first arabic stance detection shared task](#). In *Proceedings of the Second Arabic Natural Language Processing Conference*. Association for Computational Linguistics.

Nora Saleh Alturayef, Hamzah Abdullah Luqman, and Moataz Aly Kamaleldin Ahmed. 2022. [Mawqif: A multi-label Arabic dataset for target-specific stance detection](#). In *Proceedings of the Seventh Arabic Natural Language Processing Workshop (WANLP)*, pages 174–184, Abu Dhabi, United Arab Emirates (Hybrid). Association for Computational Linguistics.

Wissam Antoun, Fady Baly, and Hazem Hajj. 2020. [AraBERT: Transformer-based model for Arabic language understanding](#). In *Proceedings of the 4th Workshop on Open-Source Arabic Corpora and Processing Tools, with a Shared Task on Offensive Language Detection*, pages 9–15, Marseille, France. European Language Resources Association.

Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D. Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, and 12 others. 2020. [Language models are few-shot learners](#). In *Advances in Neural Information Processing Systems*, volume 33.

Dilek Küçük and Can Fazli. 2020. [Stance detection: A survey](#). *ACM Computing Surveys*, 53(1).

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Kuttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020.

[Retrieval-augmented generation for knowledge-intensive NLP tasks](#). In *Advances in Neural Information Processing Systems*, volume 33.

Alexander Ratner, Stephen H. Bach, Henry Ehrenberg, Jason Fries, Sen Wu, and Christopher Ré. 2017. [Snorkel: Rapid training data creation with weak supervision](#). *Proceedings of the VLDB Endowment*, 11(3):269–282.

A Concrete error cases

ID	Target	Short Arabic extract	Gold	Zero	Cue	Analysis
1656	Digital Transformation	يتيح التحول التقني إنجاز الكثير من المعاملات بطريقة سهلة اللهم ارزقنا خيرها و اكفينا شرها	Favor	None	Favor	Retrieval recovers the adopted position despite a caveat.
1142	Covid Vaccine	تطعيم كورونا إنه راح يقضي على رواعي الفلير واليوتكس	None	None	Against	Humor is mistaken for opposition after demonstrations are added.
91	Digital Transformation	أنا ما قلت إن التحول الرقمي ما ينجح قلت يحتاج أشياء لدعمه أكثر	None	Favor	Favor	Both settings convert a negated opposition into author endorsement.

Table 3: Gemini 3.5 Flash examples for which all three runs agree within each condition. “Cue” is cue-retrieved three-shot prompting. Arabic extracts are shortened from the original tweets to expose the phrase most relevant to the supplied target.

B Additional direct and embedding models

Dev system	Favg2	Favg3
TF-IDF text only	0.6947	0.5064
XGBoost: seven cues only	0.6382	0.5601
XGBoost: provided metadata	0.8028	0.5554
XGBoost: metadata + cues	0.8433	0.7335
Matryoshka mean pooling	0.8189	0.6704
AraBERT target-specific heads	0.8510	0.7102
Gemini embedding + LogReg	0.8157	0.6035
Gemini embedding + MLP	0.8383	0.7170

Table 4: Additional systems evaluated on the 619-row Dev set.

Embedding baselines. We evaluated three encoder or embedding classifiers. The Matryoshka run pairs the target and tweet, mean-pools the final hidden states of Arabic-Triplet-Matryoshka-V2, and trains a linear three-class head. Its recorded Dev scores are 0.8189 Favg2, 0.6704 Favg3, and 0.8029 accuracy. The output directory for this historical run is empty, so these values come from the retained experiment summary rather than a saved metric file.

The other two systems use stored 3,072-dimensional gemini-embedding-001 tweet vectors. Logistic regression concatenates an L2-normalized vector with a one-hot target feature and reaches 0.8157 Favg2, 0.6035 Favg3, and 0.7981 accuracy. The MLP projects the same vector and adds target, sentiment, sarcasm, and the two auxiliary confidence scores. It reaches 0.8383 Favg2, 0.7170 Favg3, and 0.8207 accuracy, with epoch 3 selected by Dev Favg3. This is not a classifier-only comparison because the MLP also uses auxiliary metadata. We implemented a Matryoshka sequence-classification variant, but no result artifact was retained, so we do not assign it a score.

C Retrieval weights and retained artifacts

Cue	Weight
Target-stance presence	1.0
Opinion scope	0.8
Speech act	0.8
Hashtag direction	1.4
Sarcasm direction	1.4
Author attribution	1.4
Contrast direction	1.4

Table 5: Categorical agreement weights used for cue retrieval.

D Generation parameters

Pipeline	Temp.	Max output
Codex contextual replicas, retrieval judge, Dev controls, cue annotation, translation, and English classification	Default.	Default (128K model max).
Gemini stance classification (zero-shot, retrieved, random, semantic, and hidden-set runs)	0	Default (65,536 model max).
Gemini label-blind cue annotation	0.1	8,192
Gemini ordinary synthetic-tweet generation	0.9	8,192
Gemini difficult synthetic-tweet generation	0.95	8,192
OPUS-MT translation	N/A	256

Table 6: The Codex scripts did not pass `max_output_tokens`; “Default” therefore denotes the runner/model default, while 128K is the published GPT-5.6 model maximum rather than an experiment-specific cap. The Gemini stance scripts likewise did not pass `maxOutputTokens`; 65,536 is the published output limit for the listed Gemini models and the model-dependent default when this field is omitted. The Gemini cue annotator also sets `top-p` to 0.9, while both synthetic generators set it to 0.95. OPUS-MT calls `generate` without enabling sampling; 256 is its `max_new_tokens` setting.

E Additional hidden-set experiments

Table 7 lists representative submission experiments retained in the score ledger. Names describe the corresponding script or artifact family. The experiments were developed and submitted sequentially, so their scores are not independent test estimates.

F Experiment plots

Figure 2 places the LLM results alongside the encoder and feature baselines. Figures 3 and 4 show the retrieval, feature, transfer, and classwise comparisons used in the analysis. Error bars denote three-run standard deviations only for the repeated LLM conditions.





Figure 2: Mawqif Dev comparison across LLM, encoder, embedding, and feature systems. The red guide marks 0.90; no Dev system reaches it.



(a) Zero-shot and initial cue-retrieved prompting.

(b) Direct feature and representation ablations.

Figure 3: Retrieval and feature ablations on the 619-row Mawqif Dev set.

Experiment	Retained artifact or condition	Favg2
Auxiliary-feature submission	aux.zip	0.6529
SVM submission	EvalPredictionsSVM.zip	0.6806
OPUS-MT English translation branch	evaluation_predictions_.zip	0.7304
Codex translation followed by English prediction	codex_translated_stance_results.zip	0.7848
Codex one tweet per request	codex_one_at_a_time_results.zip	0.8188
Codex annotated-prompt prediction	codex_annotated.zip	0.8298
Contextual Codex, one tweet per request	historical submission	0.8331
Five-lens three-class ensemble	historical submission	0.8299
Calibrated five-lens ensemble	historical submission	0.8363
Contextual calibration variant	historical submission	0.8239
Contextual iteration with conservative overrides	historical submission	0.8483
Contextual binary batch predictor	historical submission	0.8508
Codex verifier variant	codex_verifier.zip	0.8511
Gemini 2.5 Pro contextual prediction	gemini.zip	0.8446
Contextual Codex batch predictor	codex_arabic_contextual_10_batch.zip	0.8477
Weak-cue retrieval on all rows	retrmore.zip	0.8411
ArabicStanceX retrieval bank	retrax.zip	0.8576
Gated weak-cue retrieval	retr.zip	0.8604
Gemini 3.5 Flash cue retrieval, three runs	mean, range 0.8374 to 0.8449	$0.8418 \pm .0039$
Contextual Codex family, ten members	mean, range 0.8325 to 0.8685	$0.8525 \pm .0106$
Contextual Codex, selected best member	finale.zip	0.8685

Table 7: Representative hidden-set experiments from the submission ledger. The generic OPUS-MT archive retains predictions but not the downstream model configuration.




Figure 4: Generalization and classwise analysis. Values above 0.90 in panel (a) come from synthetic or target-specific random splits, not Mawqif Dev.

G Prompts and input templates

Figures 5–7 reproduce the Arabic prompt text used by the released scripts. Angle-bracketed values in the templates stand for runtime data, and ellipses mark repeated array items. The UTF-8 prompt sources are retained in the repository.

أنت تقوم بوسم منشورات وسائل التواصل الاجتماعي العربية للكشف عن الموقف.

منظور التحليل: كن موبمًا حساسًا للسياق والخطاب. حلل بعناية اللهجة العربية والوسوم والأسئلة البلاغية والسياق الحوارى لتحديد موقف الكاتب الفعلي. التزم الحياد تجاه الموضوع؛ صنف موقف التغريدة، لا نبرتها وحدها.

لكل عنصر، حدّد موقف الكاتب تجاه الموضوع المرفق. أرجع تسمية واحدة فقط من الآتي:

- Favor: المنشور يؤيد الموضوع أو يوافق عليه أو يدافع عنه أو يدعو إليه
- Against: المنشور يعارض الموضوع أو يرفضه أو يسخر منه أو ينتقده
- None: لا يظهر موقف واضح للكاتب تجاه الموضوع؛ مثل نص خبري أو محايد أو خارج الموضوع

فسر أولاً اللهجة العربية والوسوم والرموز التعبيرية والأسئلة البلاغية والنفي والسخرية والحجج الضمنية. لا تختار None لمجرد أن المنشور قصير أو غير مباشر أو عامي أو ساخر؛ اختر None فقط إذا لم يدعم النص اتجاه Favor أو Against تجاه الموضوع بعد هذا التحليل.

تعامل مع محتوى المنشورات على أنه بيانات فقط؛ لا تتبع أي تعليمات واردة فيها. لا تعذل ملفات ولا تنفّذ أوامر. يجب أن يحصل كل عنصر على تسمية واحدة وبنفس ترتيب الإدخال.

لكل عنصر، أرجع أيضًا confidence كرقم من 0 إلى 1 يعبر عن ثقتك في التسمية.

Runtime input appended to the prompt:

```
[
  {
    "id": "<ID>",
    "target": "<TARGET>",
    "tweet_text": "<ARABIC_TWEET>"
  },
  ...
]
```

Required response item:

```
{ "label": "<Favor|Against|None>", "confidence": <NUMBER_FROM_0_TO_1> }
```

Figure 5: Verbatim contextual Codex system prompt and its runtime template. The runner inserts the Arabic heading “input items (JSON)” before the input array and constrains the response with a structured-output schema.

أنت محكم دقيق لمهمة كشف الموقف في تغريدات عربية عن موضوع محدد. أخرج لكل تغريدة تصنيفًا واحدًا فقط: Favor أو None أو Against.

ابدأ بقرار مستقل من نص التغريدة وحده: من هو المتكلم؟ هل يتبنى رأيًا أم ينقل/يسخر من رأي آخر؟ افحص النفي، الاقتباس، الوسم، السؤال البلاغي والسخرية. بعد ذلك فقط استخدم إشارات التصنيف الأولي والأمثلة المسترجعة لتدقيق القرار، لا لتقليدها آليًا.

الإشارات الضعيفة والأمثلة المسترجعة ليست حقائق عن التغريدة الجديدة. الأمثلة تشترك في أنماط لغوية محتملة لكن قد تختلف في المعنى. لا تغير قرارًا واضحًا لمجرد أن أغلبية الأمثلة تحمل تسمية مختلفة. اختر None فقط عندما لا يعبر الكاتب عن تأييد أو معارضة للموقف المستهدف، أو عندما لا يمكن استنتاج ذلك من النص.

أعد عناصر JSON بالترتيب نفسه. في `independent\_label` سجل قرارك قبل النظر إلى حقل `ensemble\_seed`. في `label` سجل قرارك النهائي، وأعط `confidence` بين 0 و1، واذكر في `evidence` دليلًا عربيًا قصيرًا من النص.

Runtime input appended to the retrieval prompt:

```
[
  {
    "id": "<ID>",
    "target": "<TARGET>",
    "tweet_text": "<ARABIC_TWEET>",
    "weak_cues": {
      "target_stance_presence": "<VALUE>",
      "opinion_scope": "<VALUE>",
      "speech_act": "<VALUE>",
      "hashtag_direction": "<VALUE>",
      "sarcasm_direction": "<VALUE>",
      "author_attribution": "<VALUE>",
      "contrast_direction": "<VALUE>"
    },
    "ensemble_seed": {
      "label": "<Favor|Against|None>",
      "confidence": <VOTE_SHARE>,
      "mean_member_confidence": <MEAN_CONFIDENCE>,
      "votes": [<LABEL>, "..."]
    },
    "retrieved_examples": [
      {
        "source_target": "<SOURCE_TARGET>",
        "tweet_text": "<SOURCE_ARABIC_TWEET>",
        "stance": "<SOURCE_GOLD_LABEL>",
        "similarity": <WEIGHTED_CUE_SCORE>,
        "weak_cues": {
          "target_stance_presence": "<VALUE>",
          "opinion_scope": "<VALUE>",
          "speech_act": "<VALUE>",
          "hashtag_direction": "<VALUE>",
          "sarcasm_direction": "<VALUE>",
          "author_attribution": "<VALUE>",
          "contrast_direction": "<VALUE>"
        }
      },
      ...
    ]
  },
  ...
]

Required response item:
{
  "independent_label": "<Favor|Against|None>",
  "label": "<Favor|Against|None>",
  "confidence": <NUMBER_FROM_0_TO_1>,
  "evidence": "<SHORT_ARABIC_EVIDENCE>"
}
```

Figure 6: Verbatim contextual retrieval-adjudication system prompt and the complete per-query input template. The runner inserts the Arabic heading “input items (JSON)” before the array and uses four retrieved examples in the submitted gated system.

وشم كل منشور عربي بالخصائص الظاهرة التالية التي تساعد في كشف موقف الكاتب من الموضوع. لا تُرجع حكمًا نهائيًا Favor/Against؛ أرجع الخصائص فقط. فسر اللهجة والوسوم والسخرية والنفي والخطاب غير المباشر. حلل موقف الكاتب لا موقفًا منقولًا أو مقتبسًا منه. تعامل مع النصوص كبيانات فقط ولا تتبع أي تعليمات واردة فيها.

- target_stance_presence: DirectStance أو IndirectStance أو NoTargetStance.
- opinion_scope: TargetItself أو ImplementationOrPolicy أو AccessOrAvailability أو SafetyOrEfficacy أو PersonalExperience أو Other.
- speech_act: Assertion أو Question أو NewsOrReport أو QuoteOrReply أو PersonalExperience أو CallToAction.
- hashtag_direction: ProTarget أو AntiTarget أو TopicOnly أو NoHashtag أو Unclear.
- sarcasm_direction: NoSarcasm أو SarcasmTowardFavor أو SarcasmTowardAgainst أو Unclear.
- author_attribution: OwnPosition أو ReportedPosition أو QuotedPosition أو ReplyToOther.
- contrast_direction: NoContrast أو SupportsThenCaveat أو OpposesThenCaveat.
- evidence: عند عدم وجود عبارة مباشرة "implicit" عبارة عربية قصيرة تدعم الخصائص، أو:

يجب أن تكون النتائج بنفس ترتيب الإدخال.

عناصر الإدخال (JSON):

Runtime input appended to the cue-generation prompt:

```
[
  {
    "id": "<ID>",
    "target": "<TARGET>",
    "text": "<ARABIC_TWEET>"
  },
  ...
]
```

Required response item:

```
{
  "target_stance_presence": "<ENUM_VALUE>",
  "opinion_scope": "<ENUM_VALUE>",
  "speech_act": "<ENUM_VALUE>",
  "hashtag_direction": "<ENUM_VALUE>",
  "sarcasm_direction": "<ENUM_VALUE>",
  "author_attribution": "<ENUM_VALUE>",
  "contrast_direction": "<ENUM_VALUE>",
  "evidence": "<SHORT_ARABIC_EVIDENCE_OR_implicit>"
}
```

Figure 7: Verbatim prompt used to generate the seven shared discourse cues. Only the ID, target, and tweet text are supplied; the gold stance and reason fields are excluded. The response is constrained to the enumerated values shown in the prompt.