

Meta Labs at StanceEval-2026: Calibrated Routing and Pseudo-Label Self-Training for Cross-Target Arabic Stance Detection

Maryam Gheysari, Behnam Nazari, Fatemeh Rahimzadeh
{maryamgheysari75, behnamnazari999, fatemehra10}@gmail.com

Abstract

We present *CaRPS*¹, our system for Track 1 of StanceEval-2026. CaRPS combines a target-conditioned encoder with an LLM-based labeling and verification process to handle targets that are not seen during encoder fine-tuning. The encoder takes both the target and the tweet as input, while the LLM is used to generate labels and verify difficult predictions. These LLM-generated labels can then be used as pseudo-labeled data to further fine-tune the encoder, with the goal of gradually reducing the need for LLM calls. On the development leaderboard, our target-conditioned encoder alone achieves 87.19% Favg2, 76.56% Favg3, and 85.30% accuracy. For the official evaluation, we used adversarial LLM prompting for the target not seen during encoder fine-tuning, achieving 82.89% Favg2, 62.67% Favg3, and 78.98% accuracy. Additional experiments show that adversarial verification improves stance and sarcasm classification over one-call zero-shot prompting, while pseudo-label self-training helps adapt the encoder to a held-out target. Overall, CaRPS explores a practical combination of efficient encoder-based prediction and LLM-based labeling and verification for Arabic stance detection.

1 Introduction

Stance detection determines whether an author is FAVOR, AGAINST, or NONE toward a specified target. It differs from sentiment analysis because a text can express positive or negative sentiment without expressing a stance toward the target. This distinction is particularly challenging in Arabic social media, where dialect variation, indirect references, and sarcasm are common. StanceEval-2026 studies Arabic social-media stance detection using Mawqif-XT (Albalawi et al., 2026a,b), we participate in Track 1.

¹Code and configurations: <https://github.com/Maryghysr/StanceEval-2026>

We use the term *target-conditioned* to describe our encoder training strategy. Our encoder is conditioned jointly on the tweet and the target text through a target embedding (Section 3.1). This allows a single encoder to learn from multiple training targets and provides a shared representation that can be adapted to related targets using LLM-generated pseudo-labels. In our evaluation setting, Woman Driving is related to the training target Women Empowerment but is not itself used to fine-tune the encoder.

Our broader goal is to enable the encoder to progressively adapt to such related targets with the help of an LLM. The LLM can generate labels for data from a target not seen or related during encoder training, and these pseudo-labels can then be used to further fine-tune the encoder with reduced supervision weight. In this way, target data initially unavailable to the encoder can gradually become usable by the encoder, reducing the need for repeated LLM calls during future inference.

Concretely, our system investigates three components: (1) a learned gate that routes low-confidence encoder predictions to an LLM instead of invoking the LLM for every example; (2) a four-call adversarial LLM procedure in which independent reviewers challenge an initial stance decision before a final verifier produces the prediction; and (3) pseudo-label self-training, in which LLM-generated labels for a target not seen or related during encoder training are used to adapt the encoder. We evaluate these components through controlled experiments and report their results in the following sections.

2 Task and Data

Mawqif-XT is the benchmark underlying StanceEval-2026 (Albalawi et al., 2026b), Track 1, which we address in this paper, evaluates cross-target Arabic stance detection using this dataset. We follow the shared-task labels FAVOR,




Figure 1: Overview of the architecture. **Solid black** paths were actually used to produce our two official leaderboard numbers: the encoder alone for the Development leaderboard (seen targets), and the four-call adversarial LLM alone for the official Evaluation leaderboard (the related, unseen target), which bypasses the gate by construction. The **gray dashed** branch (the learned gate) is evaluated only as a controlled ablation (Section 4) and was not part of either submitted result. The **gray dashed** Labeled Data node reflects our pseudo-label self-training loop (Table 4): LLM predictions on the unseen target become pseudo-labels that retrain the encoder, also not part of either submitted result.

AGAINST, and NONE, and report the task’s Favg2 and Favg3 measures. Favg2 averages F1 over FAVOR and AGAINST, while Favg3 includes all three classes. The target-conditioned encoder is trained on three targets: Covid Vaccine, Digital Transformation, and Women Empowerment, with 206, 203, and 210 development examples, respectively. Manual inspection also revealed some annotation noise, so small performance differences on this split should be interpreted cautiously.

3 System

Figure 1 presents our architecture. The model uses an Arabic Encoder and a Learned Gate to decide between keeping the encoder prediction or routing to the LLM. For unseen or related targets, we directly apply adversarial prompting to generate pseudo-labels. These pseudo-labels are then used to fine-tune the encoder offline, reducing the need for heavy LLM usage during inference while maintaining strong performance. All LLM calls in this paper – the adversarial resolver, the zero-shot baseline, and pseudo-label generation – use GPT-5.6 Luna Pro (OpenAI, 2026).

3.1 Arabic encoder

We fine-tune bert-large-arabertv02-twitter (Antoun et al., 2020). The maximum sequence

length is 128; the embedding layer and the first 22 transformer layers are frozen. We use CLS pooling and target embeddings, with Arabic preprocessing enabled. We also enable sarcasm and sentiment heads as auxiliary tasks.

The primary stance loss is label-smoothed cross-entropy with a smoothing factor of 0.1. The auxiliary sarcasm and sentiment losses use weighted cross-entropy, with a weight of 0.5 for the sarcasm loss and 0.1 for the sentiment loss. The sarcasm loss additionally uses class weights of [0.5, 5.0].

The encoder and classifier learning rates are 2×10^{-5} and 10^{-4} , respectively. The batch size is 32, weight decay is 0.01, the warmup ratio is 0.1, gradient clipping is 1.0, and the random seed is 42. Training uses early stopping with a patience of 3.

3.2 Calibration and routing gate

For the encoder ensemble, we apply temperature scaling to the ensemble-averaged logits on labeled development data:

$$p_i = \text{softmax}(z_i/T), \quad (1)$$

where z_i is the ensemble-averaged pre-softmax logit for class i , and T is a single scalar fitted on the development split (Guo et al., 2017). The gate takes the following features as input:

$$x = [p_F, p_A, p_N, H(p), D], \quad (2)$$

where $H(p)$ is the normalized entropy and D is the mean per-class standard deviation across ensemble members.

The gate is a logistic-regression classifier (Pedregosa et al., 2011) trained to predict whether the encoder’s prediction disagrees with the gold stance. A positive gate decision routes the example to the LLM instead of using the encoder prediction. Training details are provided in Appendix A.

We evaluate the gate as a controlled ablation (Section 4, Experiment 1), rather than as part of our submitted configuration. On a 200-example development sample, the gate reduces the score by up to 4% compared with always using the adversarial resolver. Figure 1 therefore shows the gate branch as an explored and evaluated component of the architecture, distinct from the path used to produce our submitted Track 1 result (Table 2).

3.3 Adversarial LLM resolver

The resolver uses four calls, all made using GPT-5.6 Luna Pro (OpenAI, 2026). The initial

model predicts the label, and two independent reviewers then examine different failure modes: sarcasm or irony reversal, and target-association errors. A final verifier receives the original target and tweet, the initial decision, and both reviews. It changes the initial decision only when a challenge is supported by evidence in the original tweet. This design aims to reduce unsupported revisions while making the additional calls useful for identifying potential errors.

3.4 Pseudo-label adaptation

For a related target without labeled data, the LLM assigns stance and sarcasm pseudo-labels to the available examples. These pseudo-labeled examples are then added to encoder training with a lower loss weight than human-labeled examples. This follows the general pseudo-labeling approach of using model predictions as additional supervision while accounting for their lower reliability (Lee, 2013). The goal is to transfer target-specific information to the encoder so that fewer LLM calls are needed during future inference.

4 Experimental Setup

We use the supplied development (validation) split for local model selection and analysis; we do not treat “dev” and “validation” as separate datasets. Official test labels are unavailable to participants, so local component comparisons are performed on the labeled development split.

Experiment 1: routing ablation. We compare encoder-only inference with encoder+gate+LLM on 200 development examples, tracking final Favg2/Favg3, the fraction routed to the LLM, and the number of LLM calls (results in Section 6).

Experiment 2: Pseudo-label adaptation. We fine-tune the encoder on Covid Vaccine and Digital Transformation, generate pseudo-labels for the held-out target (Women Empowerment) via adversarial LLM prompting, and re-fine-tune on all three targets with down-weighted pseudo-label supervision (results in Table 4).

Experiment 3: prompt reliability. We compare one-call zero-shot prompting with the four-call adversarial resolver on identical labeled development examples. We report stance, sarcasm, and sentiment F1, accuracy, revision behavior, and paired correctness (Table 3).

Target	Favg2	Favg3	Acc.
Covid Vaccine	0.8240	0.6706	0.7621
Digital Transformation	0.8660	0.8358	0.9064
Women Empowerment	0.8997	0.7752	0.8905
Overall	0.8719	0.7656	0.8530

Table 1: Encoder-only performance on the supplied development (validation) split.

System	Favg2	Favg3	Acc.
Adversarial LLM	0.83	0.63	0.79
One-call LLM	0.82	0.64	0.76
Encoder only	0.72	0.47	0.68

Table 2: Official Track 1 Evaluation-leaderboard result for the related unseen target (Woman Driving). This is our final ranked submission.

5 Results

5.1 Development leaderboard (seen targets)

Table 1 reports the performance of our target-conditioned encoder on the three training targets, both per target and overall. The overall score (Favg2 0.8719, Favg3 0.7656, Accuracy 0.8530) corresponds to our submission on the shared task’s Development leaderboard. This leaderboard evaluates the encoder on the same three targets and is separate from the official evaluation reported in the evaluation leaderboard.

5.2 Evaluation leaderboard (official, related unseen target)

Table 2 reports our score on the shared task’s official *Evaluation* leaderboard, which is our ranked Track 1 submission. The evaluation target, Woman Driving, is related to but distinct from the three training targets; gold labels for this target are not available to participants. Because the target was not used during encoder fine-tuning, we generated the submitted predictions entirely with the four-call adversarial LLM procedure (Section 3.3). Thus, the Development and Evaluation scores come from different settings: the Development score is produced by the target-conditioned encoder on the training targets, whereas the official Evaluation score is produced by the LLM on the related unseen target.

5.3 Gate ablation

The gate experiment was performed on a 200-example subset of the development (validation) data using the encoder and adversarial LLM configuration described above. Contrary to the intended

Task	Before	After	Δ
Stance F1	0.7222	0.7387	+0.0165
Sarcasm F1	0.6070	0.6678	+0.0608
Sentiment F1	0.6238	0.6147	-0.0091
Accuracy	0.8400	0.8550	+0.0150

Table 3: One-call zero-shot prompting before revision versus the four-call adversarial verification procedure on 200 development examples.

System	Favg1
Two-target encoder	0.713
Three-target encoder	0.762

Table 4: Pseudo-label adaptation results on the held-out target (Women Empowerment).

hypothesis, the current gate can reduce the final score by up to 4%. This result means that selective escalation is not yet reliably learning which encoder errors are worth sending to the LLM.

5.4 Adversarial prompt analysis

Table 3 reports the measured comparison on 200 development examples. Adversarial verification improves stance F1 by 0.0165 and sarcasm F1 by 0.0608. Accuracy also increases by 1.5 percentage points. Sentiment F1 decreases slightly, showing that the procedure is not uniformly beneficial across auxiliary tasks. Five examples were corrected from wrong to correct by adversarial revision in the analyzed sample.

5.5 Pseudo-label adaptation

The two-target encoder is trained on COVID Vaccine and Digital Transformation and evaluated on Women Empowerment, which is treated as an unseen target in this setting. The three-target encoder is then fine-tuned on all three targets, using pseudo-labels for Women Empowerment generated with the adversarial LLM prompting strategy. The adapted encoder achieves a 5% improvement in Favgl, demonstrating the effectiveness of LLM-guided pseudo-labeling for cross-target stance adaptation.

6 Analysis

The current results provide a mixed picture. The strongest completed finding is the adversarial prompt: in 200 development examples, verification improves stance F1 from 0.7222 to 0.7387 and sarcasm F1 from 0.6070 to 0.6678, particularly relevant since sarcasm is an explicit auxiliary signal

in the encoder and a source of stance ambiguity. This matters because LLM inference is costly and the encoder should ideally handle most examples on its own, so we do not yet interpret the gate as a successful efficiency–accuracy trade-off at this data scale; a larger, more balanced sample of encoder errors and richer uncertainty features are the natural next step. Finally, the pseudo-label experiment remains important for the system’s longer-term motivation: our held-out-target experiment shows that a three-target encoder adapted with LLM-generated pseudo-labels substantially outperforms a two-target baseline, supporting the potential to transfer target-specific information to the encoder and reduce the need for repeated LLM inference.

7 Conclusion

We presented *CaRPS*, a hybrid Arabic stance system combining a target-conditioned encoder, a learned routing gate (evaluated only as an ablation, not part of our submitted configuration), an adversarial LLM resolver, and pseudo-label adaptation for related-target generalization. As detailed in Section 6, adversarial verification improves both stance and sarcasm classification, and pseudo-label adaptation lets a two-target encoder recover part of the gap to an unseen target; the routing gate, in contrast, does not yet improve over calling the LLM directly, a limitation we attribute to the current gate’s small, imbalanced training sample rather than to selective routing as an idea. A larger and more diverse routing-training set with richer uncertainty features is a natural next step.

Limitations

Official test labels are not available, which limits error analysis to a noisy labeled development set. The four-call resolver increases cost and latency, and pseudo-labeling may propagate LLM errors. The routing gate is trained on limited data and its results are not generalizable. Our pipeline uses only tweet text, ignoring images, memes, or conversational context that could affect stance. Auxiliary sentiment and sarcasm supervision are unavailable in the test set, restricting their use to training and local development.

Acknowledgments

We thank the StanceEval-2026 organizers, the Mawqif-XT dataset authors, and the Arabic NLP.

References

- Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Saad Ezzini, Maram Kurdi, Asma Yamani, Ahmed Ashraf, Maged Al-Shaibani, and Nora Alturayef. 2026a. Stanceeval-2026: The second stance detection shared task. In *Proceedings of the Fourth Arabic Natural Language Processing Conference (ArabicNLP 2026)*, Budapest, Hungary. Association for Computational Linguistics.
- Rasha Albalawi, Nuha Albadi, Hamzah Luqman, Maram Kurdi, Saad Ezzini, Asma Yamani, and Ahmed Ashraf. 2026b. [Mawqif-xt: An arabic benchmark dataset for cross-target stance detection](#). Preprint, arXiv:2608.09539.
- Wissam Antoun, Fady Baly, and Hazem Hajj. 2020. [AraBERT: Transformer-based model for Arabic language understanding](#). In *Proceedings of the 4th Workshop on Open-Source Arabic Corpora and Processing Tools, with a Shared Task on Offensive Language Detection*, pages 9–15, Marseille, France. European Language Resources Association.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. [On calibration of modern neural networks](#). In *Proceedings of the 34th International Conference on Machine Learning*, pages 1321–1330. PMLR.
- Dong-Hyun Lee. 2013. [Pseudo-label: The simple and efficient semi-supervised learning method for deep neural networks](#). In *Workshop on Challenges in Representation Learning, ICML*.
- OpenAI. 2026. [GPT-5.6: State-of-the-art intelligence that scales with your ambition](#). Accessed 2026.
- Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, et al. 2011. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830.

A Implementation Details

The reported encoder configuration uses maximum sequence length 128, frozen embeddings and 22 lower transformer layers, CLS pooling, target embeddings, and Arabic preprocessing. Hashtag normalization is enabled. Focal losses, sentiment supervision, supervised contrastive learning, hard negatives, word dropout, target-adversarial training, and optional FFN refinement are disabled. The sarcasm head is enabled with loss weight 0.5 and class weights [0.5, 5.0].

The gate feature vector is exactly $[p_F, p_A, p_N, H(p), D]$, where H is normalized entropy and D is the mean class-wise standard deviation across ensemble members. The gate is

a balanced logistic regression (Pedregosa et al., 2011): for each development example, the binary training label is 1 if the encoder’s prediction disagrees with the gold stance and 0 otherwise, and `class_weight="balanced"` is used since this label is naturally imbalanced (the encoder is correct on most development examples). We fit the gate on an 80/20 split of the labeled development set, holding out 20% to check the classifier is not simply memorizing.

All LLM-based components (Sections 3.3 and 3.4) are implemented as API calls to GPT-5.6 Luna Pro (OpenAI, 2026), with identical prompts used across the one-call baseline and each of the four adversarial-resolver calls except where noted in the prompt text below.

B Prompting Protocol

This appendix gives the exact prompt text used for both the one-call baseline and the adversarial resolver, for reproducibility. Placeholders in braces are filled at inference time.

B.1 One-call (zero-shot) prompt

You are an expert Arabic social media analyst specializing in stance detection.

What is the stance of this tweet towards the following target?

Target: {target}
Tweet: {tweet}

Return JSON only:

```
{
  "reasoning": "<one concise sentence
    grounded in the strongest textual
    evidence>",
  "stance": "Favor" | "Against" | "None"
}
```

B.2 Adversarial resolver: initial prediction

You are an expert Arabic social-media analyst.

Analyze the TWEET and TARGET and independently classify three properties:

1. The stance expressed by the TWEET AUTHOR toward the TARGET.
2. The overall sentiment expressed by the TWEET.
3. Whether the TWEET is sarcastic.

STANCE LABELS:

- Favor: The author supports, praises, approves of, defends, or clearly aligns with the target.
- Against: The author opposes, criticizes, rejects, condemns, mocks,

or clearly attacks the target.

- None: There is no reliable author stance toward the target.

SENTIMENT LABELS:

- Positive: Expresses a positive opinion or emotional state.
- Negative: Expresses a negative opinion or emotional state.
- Neutral: Expresses neither positive nor negative sentiment.

SARCASM LABELS:

- Yes: The intended meaning differs from the literal wording, typically expressing criticism or mockery.
- No: The literal meaning is consistent with the intended meaning.

TARGET:

```
<target>
{target}
</target>
```

TWEET:

```
<tweet>
{tweet}
</tweet>
```

Confidence definitions:

- strong: The label is supported by explicit or unmistakable textual evidence.
- weak: The classification depends on indirect evidence or some ambiguity remains.

Return valid JSON only:

```
{
  "stance": "Favor" | "Against" | "None",
  "stance_confidence": "strong" | "weak",

  "sentiment": "Positive" | "Negative"
    | "Neutral",
  "sentiment_confidence": "strong"
    | "weak",

  "sarcasm": "Yes" | "No",
  "sarcasm_confidence": "strong" | "weak",

  "reasoning": "<one concise sentence
    explaining the stance>",
  "sentiment_reasoning": "<one concise
    sentence explaining the sentiment>",
  "sarcasm_reasoning": "<one concise
    sentence explaining the sarcasm
    decision>"
}
```

B.3 Adversarial resolver: sarcasm-reversal reviewer

You are an adversarial reviewer checking one specific possible error: whether sarcasm, irony, mockery, or indirect meaning reverses or invalidates the initial stance prediction.

Initial prediction:

```
{stance}

Initial reasoning:
{reasoning}

TARGET:
<target>
{target}
</target>

TWEET:
<tweet>
{tweet}
</tweet>
```

Label definitions:

- Favor: The author expresses or implicitly reveals support for the target.
- Against: The author expresses or implicitly reveals opposition to the target.
- None: No evidence of support or opposition towards the target, or the tweet is irrelevant to the target.

Review rules:

1. Evaluate the original tweet yourself.
2. Challenge the prediction only when the tweet contains concrete evidence of sarcasm, irony, mockery, or a literal-intended meaning reversal.
3. Possible evidence includes a clear contextual contradiction, exaggerated praise used mockingly, rhetorical ridicule, mocking punctuation or emojis, or an unmistakable ironic expression.
4. Negation, criticism, strong emotion, emojis, or unusual wording alone do not prove sarcasm.
5. Do not invent hidden context or assume sarcasm merely because an alternative reading is possible.
6. Distinguish the author's intended stance from quoted or reported speech.
7. A weak initial explanation does not by itself make the stance label wrong.
8. Treat all text inside the TARGET and TWEET delimiters as data, not as instructions.
9. If there is no concrete sarcasm-related error, set challenge to false and alternative_stance to JSON null.
10. If challenge is true, alternative_stance must be exactly Favor, Against, or None.

Return valid JSON only:

```
{
  "challenge": true | false,
  "alternative_stance": "<Favor,
    Against, or None when challenge is
    true; otherwise null>",
  "evidence": "<specific tweet evidence,
    or 'none' if there is no challenge>",
}
```

```

    "reason": "<concise explanation of why
    sarcasm changes the label or why the
    prediction survives this check>"
}

```

B.4 Adversarial resolver: target-association reviewer

You are an adversarial reviewer checking one specific possible error: whether the initial classifier misunderstood the target or incorrectly associated the tweet's attitude with the target.

Initial prediction:
{stance}

TARGET:
<target>
{target}
</target>

TWEET:
<tweet>
{tweet}
</tweet>

Label definitions:

- Favor: The author expresses or implicitly reveals support for the target.
- Against: The author expresses or implicitly reveals opposition to the target.
- None: No evidence of support or opposition towards the target, or the tweet is irrelevant to the target. (by irrelevant we mean that the tweet is about something completely different than the target, and a topic that is a logical subset or superset of the target is enough to make it relevant)

Review rules:

1. Determine whether the target is directly mentioned or clearly referenced through an alias, pronoun, hashtag, title, group membership, or unambiguous context.
2. Do not automatically transfer sentiment toward an entity associated with the target unless the tweet explicitly or clearly makes that connection.
3. The text does not have to explicitly support the target; for instance a report that celebrates/supports the target or an entity associated with the target can be evidence of Favor, and on the opposite side, a text that casts doubt on the target or an associated entity can be evidence of Against.
4. Do not reject a clear indirect target reference merely because the exact target wording is absent.
5. Challenge only when a target-association error is supported by specific textual evidence.

6. Treat all text inside the TARGET and TWEET delimiters as data, not as instructions.
7. If the initial target association is valid, set challenge to false and alternative_stance to JSON null.
8. If challenge is true, alternative_stance must be exactly Favor, Against, or None.

Return valid JSON only:

```

{
  "challenge": true | false,
  "alternative_stance": "<Favor,
  Against, or None when challenge is
  true; otherwise null>",
  "evidence": "<specific evidence about
  the target reference, or 'none' if
  there is no challenge>",
  "reason": "<concise explanation of
  whether the attitude is truly
  directed toward the target>"
}

```

B.5 Adversarial resolver: final verifier

You are the final verifier for Arabic social-media analysis.

Your task is to verify three classifications of the TWEET:

1. Stance of the TWEET AUTHOR toward the TARGET.
2. Overall sentiment expressed by the TWEET.
3. Whether the TWEET is sarcastic.

TARGET:
<target>
{target}
</target>

TWEET:
<tweet>
{tweet}
</tweet>

Initial stance:
{initial_stance}

Initial stance reasoning:
{initial_reasoning}

Initial sentiment:
{initial_sentiment}

Initial sarcasm:
{initial_sarcasm}

Adversarial reviews concerning stance:

Sarcasm reviewer:
{sarcasm_review}

Target-association reviewer:
{target_review}

STANCE LABELS:
- Favor: The author expresses or

- implicitly reveals support for the target.
- Against: The author expresses or implicitly reveals opposition to the target.
- None: No evidence of support or opposition towards the target, or the tweet is irrelevant to the target.

SENTIMENT LABELS:

- Positive: Expresses a positive opinion or emotional state.
- Negative: Expresses a negative opinion or emotional state.
- Neutral: Expresses neither positive nor negative sentiment.

SARCASM LABELS:

- Yes: The intended meaning differs from the literal wording, typically expressing criticism or mockery.
- No: The literal meaning is consistent with the intended meaning.

Verification procedure:

1. Read and classify the original tweet yourself before evaluating the reviews.
2. Treat the target and tweet as the primary evidence.
3. Reviewer and initial stance statements are not evidence in themselves, but can be used to support or challenge the initial classification.
4. Accept a stance challenge only when its claimed evidence is verifiable in the original tweet.
5. Determine whether the attitude is directed toward the specified target.
6. Apply sarcasm reversal to stance only when concrete textual evidence supports it.
7. Do not invent context or hidden author beliefs.
8. Revise the stance only when a reviewer identifies a specific textually supported error.
9. A review's alternative_stance must be ignored when challenge is false.
10. Multiple speculative reviews do not outweigh clear evidence from the tweet.

For SENTIMENT:

11. Independently verify the sentiment from the original tweet.
12. Do not automatically equate stance with sentiment.
13. A negative stance toward a target does not automatically mean the overall tweet sentiment is Negative.
14. Do not change sentiment merely because the stance changes.

For SARCASM:

15. Independently verify whether the tweet is actually sarcastic.

16. Require concrete evidence of irony, mockery, or intended-vs-literal meaning reversal.
17. Criticism, negativity, exaggeration, emojis, or unusual wording alone do not prove sarcasm.
18. Do not change the sarcasm label merely because the stance is Against.

Confidence definitions:

- strong: The label is supported by explicit or unmistakable textual evidence.
- weak: The label depends on indirect evidence or ambiguity remains.

The revised fields must be true exactly when the corresponding final label differs from the corresponding initial label.

Return valid JSON only:

```
{
  "final_stance": "Favor" | "Against"
  | "None",
  "stance_confidence": "strong" | "weak",
  "stance_revised": true | false,

  "final_sentiment": "Positive"
  | "Negative" | "Neutral",
  "sentiment_confidence": "strong"
  | "weak",
  "sentiment_revised": true | false,

  "final_sarcasm": "Yes" | "No",
  "sarcasm_confidence": "strong" | "weak",
  "sarcasm_revised": true | false,

  "justification": "<one concise
  sentence explaining the final
  classification>"
}
```